

Peer-to-Peer and Cloud Computing

MapReduce and Related Technologies

14th January 2019

Universität Augsburg
Institut für Informatik
Lehrstuhl für Organic Computing

MapReduce

The “next generation” of Big Data algorithms

MapReduce

- initially created by Google Inc.
- over the years: “hundreds of special-purpose computations that process large amounts of raw data”
 - crawled documents → inverted indices
 - web request logs → frequencies
 - etc.
- computations themselves are straightforward
- however: a *lot* of data
⇒ distributed (parallel) computation necessary

- simple programming model
- messy details are hidden, notably
 - parallelisation,
 - fault-tolerance,
 - data distribution and
 - load balancing

- inspired by functional programming (esp. Lisp)
- example usage of Lisp's **map**:

```
(map 'list #'(lambda (x) (+ x 1)) '(1 2 3 4 5))  
⇒ (2 3 4 5 6)
```

- takes a function and a sequence and
- applies that function to every element of the sequence
returning a new sequence

- **reduce** also comes from functional languages (esp. Lisp):
`(reduce #'(lambda (x y) (* x y)) '(1 2 3 4 5))`
 \Rightarrow (120)
- takes a function and a sequence
- the first element serves as the first *accumulator value*
- then, the next accumulator value is the value of the binary function applied to the *current accumulator value* and the *next element*

A MapReduce application has to implement two functions with the following signatures:

- `map :: (k1, v1) → [(k2, v2)]`
- `reduce :: (k2, [v2]) → [v2]`

Note: The **map**-function from MapReduce has a different signature than the *original* **map**-function from functional programming (e.g. from Lisp)! It more resembles a function that would be used *with* a **map**-function from functional programming.

- given a set of documents,
- calculate the number of occurrences for each word
- algorithmically simple task, however: assume a very large set of documents

```
public Map<Word, Integer>
count(Set<Document> documents) {
    Map<Word, Integer> counts = new Map<>();
    for (Document d : documents) {
        for (Word w: d.toWords()) {
            counts.merge(w, 1, Integer::add);
        }
    }
}
```

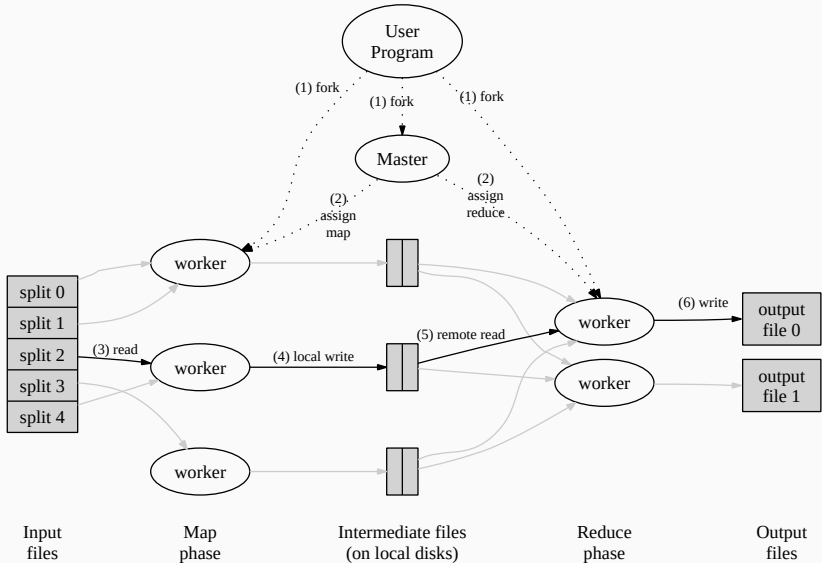
MapReduce separates processing a single document from “merging” that document’s word counts with the other documents’ word counts.

```
public void map(Key key, Document document) {  
    for (Word w : document.toWords()) {  
        emitIntermediate(w, 1);  
    }  
}
```

Note: While **key** is not used here, it may be required in another application.

Why use `emitIntermediate` instead of aggregating a list and `returning`?

```
public void reduce(Word w, Set<Integer> counts) {  
    Integer sum = 0;  
    for (Integer i : counts) {  
        sum += i;  
    }  
    emit(w, sum);  
}
```



- input is partitioned into pieces of a user-given size (e.g. 16MB)
⇒ distributed **map**'ing
- key space of intermediary results is partitioned using user-given partitioning function
⇒ distributed **reduce**'ing
- “emitting” (intermediary) results \approx writing them to a shared memory
- role of the master node:
 - picks idle workers for map or reduce tasks
 - gives them the memory address of their input

Why is *parallel execution* so easily possible?


Because there are no dependencies between **map** calls
(*partition* of input space)!

- on worker failure: simply reschedule all its tasks on other workers (no additional overhead necessary because no dependencies!)
- on master failure: either
 - use latest backup of its state and restart it, or
 - simply re-run the whole MapReduce task

- Hadoop project consists of
 - Hadoop Distributed File System (HDFS)** high-throughput access to application data
 - Hadoop YARN** framework for job scheduling and cluster resource management
 - Hadoop MapReduce** YARN-based implementation of MapReduce
- Open source!
- YARN and HDFS are general purpose
 - ⇒ they are relied on by other projects as well

The “next generation” of Big Data algorithms


Google trend graph for “MapReduce” keyword



- from a **developer's perspective**: MapReduce unnecessarily restricted!
 - Why only exactly support the MapReduce pipeline?
 - Why only support functional “primitives” **map** and **reduce**?
 - E.g., why not also
 $\text{filter} :: (a \rightarrow \text{Bool}) \rightarrow [a] \rightarrow [a]$
- from a **performance perspective**: iterative applications (such as machine learning) not properly supported!
 - don't behave well if too much I/O is involved
 - MapReduce always writes results to (distributed) disk
 \Rightarrow too much overhead through I/O

- more general than MapReduce: Directed Acyclic Graphs (DAGs)
- MapReduce is a DAG itself:
in \longrightarrow **map** \longrightarrow **reduce** \longrightarrow **out**
- DAGs are really composed (pure!) functions
 \Rightarrow *functional programming*

Example for DAG and the corresponding function



This DAG equals the following function:

```
out(in1, in2, in3) =  
  filter(  
    join(  
      reduce(in1),  
      union(  
        map(in2),  
        map(in3)))
```




- “A fast and general engine for large-scale data processing.”
- more general than “just” MapReduce
- many more parallelised functional operations
- builds on a distributed storage (e.g. HDFS or a simple shared file system)
- holds more stuff in memory than (original) MapReduce
⇒ greatly(!) increased performance




- Scala has proper type inference \Rightarrow only difference to “non-distributed” (functional) Scala code to initially create an implicit **SparkContext**

```
// sc: SparkContext
val lines = sc.textFile("data.txt")
val lineLengths = lines.map(s => s.length)
val totalLength =
    lineLengths.reduce((a, b) => a + b)
```

- As long as you keep programming *purely functional*, no need to think about parallelism at all!
- Need **lineLengths** somewhere else in your application?
 \rightarrow simply add **lineLengths.persist()** and that constant is written to the distributed storage

- increasingly less knowledge required for using parallelism
- *functional programming* is becoming more and more important

-  Jeffrey Dean and Sanjay Ghemawat. “MapReduce: Simplified Data Processing on Large Clusters”. In: *Communications of the ACM* 51.1 (2008), pp. 107–113.
-  Google. *MapReduce on Google Trends*. 2017. URL: <https://trends.google.com/trends/explore?date=all&q=MapReduce> (visited on 10/11/2017).
-  LispWorks Ltd. *Lisp Function MAP*. 2017. URL: http://clhs.lisp.se/Body/f_map.htm (visited on 10/11/2017).

-  LispWorks Ltd. *Lisp Function REDUCE*. 2017. URL: http://clhs.lisp.se/Body/f_reduce.htm (visited on 10/11/2017).
-  Apache Spark Team. *Apache Spark Programming Guide*. 2017. URL: <https://spark.apache.org/docs/latest/rdd-programming-guide.html> (visited on 10/10/2017).
-  Apache Spark Team. *Apache Spark Website*. 2017. URL: <https://spark.apache.org/> (visited on 10/10/2017).