

# A Unified Theory of Shared Memory Consistency

ROBERT C. STEINKE AND GARY J. NUTT

*University of Colorado at Boulder, Boulder, Colorado*

**Abstract.** The traditional assumption about memory is that a read returns the value written by the most recent write. However, in a shared memory multiprocessor several processes independently and simultaneously submit reads and writes resulting in a partial order of memory operations. In this partial order, the definition of most recent write may be ambiguous. Memory consistency models have been developed to specify what values may be returned by a read given that memory operations may only be partially ordered. Before this work, consistency models were defined independently. Each model followed a set of rules which was separate from the rules of every other model. In our work, we have defined a set of four consistency properties. Any subset of the four properties yields a set of rules which constitute a consistency model. Every consistency model previously described in the literature can be defined based on our four properties. Therefore, we present these properties as a unified theory of shared memory consistency.

Our unified theory provides several benefits. First, we claim that these four properties capture the underlying structure of memory consistency. That is, the goal of memory consistency is to ensure certain declarative properties which can be intuitively understood by a programmer, and hence allow him or her to write a correct program. Our unified theory provides a uniform, formal definition of all previously described consistency models, and in addition some combinations of properties produce new models that have not yet been described. We believe these new models will prove to be useful because they are based on declarative properties which programmers desire to be enforced. Finally, we introduce the idea of selecting a consistency model as an on-line activity. Before our work, a shared memory program would run start to finish under a single consistency model. Our unified theory allows the consistency model to change as the program runs while maintaining a consistent definition of what values may be returned by each read.

Categories and Subject Descriptors: F.1.2 [Computation by Abstract Devices]: Modes of Computation—*parallelism and concurrency*

General Terms: Theory

Additional Key Words and Phrases: Distributed shared memory systems, memory consistency models, memory consistency model lattice

## 1. Introduction

Shared memory is a powerful abstraction for interprocess communication. The concept of shared memory originated from multiprogramming on uniprocessors

---

Authors' address: University of Colorado at Boulder, Department of Computer Science, 430 UCB, Boulder, CO 80309-0430. Corresponding e-mail: rsteinke@altrionet.com.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or direct commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 1515 Broadway, New York, NY 10036 USA, fax: +1 (212) 869-0481, or [permissions@acm.org](mailto:permissions@acm.org).

© 2004 ACM 0004-5411/04/0900-0800 \$5.00

and bus-based multiprocessors. In these environments there is a simple model of the memory system enforced in hardware. The model can be stated as:

- There is a physical memory cell that represents each variable. The state of this memory cell is the state of the variable.
- Memory operations take place sequentially. They are atomic and there is a total order on all memory operations. Read operations return the current state of the physical memory cell. Write operations change the current state of the physical memory cell and the change becomes observable to all processes simultaneously.
- The operations of each process take place in the order specified by its program.

These conditions are enforced by the hardware architecture. In a multiprogrammed uniprocessor, there really is only one process submitting memory operations at a time. In a bus-based multiprocessor with no cache, the bus serves as a serialization mechanism that allows operations to reach memory sequentially.

For many years these assumptions were implicit and any computer scientist would tell you, “That’s just how memory works.” Then two things happened. The first is that memory systems in multiprocessors got more and more complicated [Dubois and Scheurich 1990; Dubois et al. 1986; Gharachorloo et al. 1990; Lenoski et al. 1990]. The second is the invention of distributed shared memory (DSM) for the message-passing multicomputer [Amza et al. 1996; Bennett et al. 1990, 1995; Bershad et al. 1993; Bershad and Zekauskas 1991; Li 1986; Li and Hudak 1989]. Caching and out-of-order instruction dispatching can pose a problem for multiprocessors. The hardware of each processor enforces the restriction that the processor sees its own memory operations in the order specified by its program, but this does not automatically protect processors from seeing each other’s operations out of order. DSM provides the illusion of a shared address space on top of hardware that only supports message passing. In DSM systems, asynchronous messages and replicated copies of data can cause the same problems.

These problems led to the concept of consistency models. A consistency model is a specification of the allowable behavior of memory. It can be seen as a contract between the memory implementation and the program utilizing memory [Tanenbaum 1995]. The input to memory is a set of memory operations (reads and writes) partially ordered by program order. The output of memory is the collection of values returned by all read operations. A consistency model is a function that maps each input to a set of allowable outputs. The memory implementation guarantees that for any input it will produce some output from the set of allowable outputs specified by the consistency model. The program must be written to work correctly for any output allowed by the consistency model. This idea was originally described by Lamport [1979] when he defined sequential consistency. A sequentially consistent multiprocessor allows conventional reasoning about the correctness of programs. Essentially, it allows the programmer to treat the machine as a multiprogrammed uniprocessor. Enforcing sequential consistency can be very costly. Soon weaker consistency models were discovered that were less expensive in terms of communication. Multiprocessors were generally used for large numerical programs that were already programmed with a constrained programming style to avoid data race conditions. With slight modifications to the programming style, algorithms could still be written to execute correctly for non-sequentially consistent memory systems.




FIG. 1. Shared memory as an API.

With consistency models, the concept of shared memory is no longer tied to the physical implementation of memory cells. A programmer can write a correct program using the abstractions of concurrent processes and shared memory with little knowledge about the underlying implementation that will eventually execute the program. All that the programmer needs to know is the consistency model enforced by memory. To give the memory implementor more flexibility for optimization, the memory might enforce fewer guarantees. Or to make the programmer's job easier the memory might enforce more guarantees. Many choices have been made along this ease of use to efficient implementation continuum. The results are the consistency models described in the literature [Ahmad et al. 1991; Bershad and Zekauskas 1991; Dubois et al. 1986; Gao and Sarkar 2000; Gharachorloo et al. 1990; Goodman 1989; Herlihy and Wing 1990; Hutto and Ahmad 1990; Iftode et al. 1996; Keleher et al. 1992; Lamport 1979; Lipton and Sandberg 1988]

This leads to the idea of shared memory as an application programming interface (API) as shown in Figure 1. The program and memory agree on a consistency model. Then the program executes using the shared memory API, and the program's processes share information in a common address space. No knowledge is needed of the memory implementation.

The hypothesis of our work is that every consistency model can be described as different combinations of a set orthogonal properties with each property defined as a constraint on per-process views of memory operations. If all of the constraints for a particular consistency model are satisfied in the views of all processes then that execution satisfies that consistency model. In fact, we have found just such a set of four orthogonal properties which describe all previously discovered consistency models. One of the major contributions of this work is that cache consistency and sequential consistency can now be described and checked against per-process views. This provides a uniform framework and notation for all consistency models. An additional contribution is that enumerating all combinations of these four properties produces a lattice of consistency models. With this, lattice relationships among different consistency models can easily be seen, and there are combinations in the lattice which correspond to consistency models which have not yet been discovered.

A final contribution is the description of synchronized consistency models as consistency transitions. Certain consistency models have special memory synchronization operations. We have discovered that these can be accurately modeled as transitioning to a stronger consistency model, issuing ordinary read and write operations that obey the stronger model, and then transitioning back to a weaker consistency model. Rather than having two independent classes of consistency models both synchronized and nonsynchronized models can now be described in the same framework with nonsynchronized models as merely a special case where

the program never initiates a consistency transition. This may open up avenues for new synchronized consistency models by examining different possible transition combinations in the consistency lattice. It also suggests the idea of consistency model selection as an on-line activity where a program might be written requiring different consistency guarantees at different times.

Section 2 discusses related work. Section 3 presents the formalism we use to describe consistency models. Section 4 presents the four consistency properties and the lattice of consistency models. Section 5 presents our model of consistency transitions and shows its equivalence to synchronized consistency models. Section 6 discusses our conclusions and future work.

## 2. Related Work

**2.1. SHARED MEMORY.** A common trend in the literature is the development of uniform frameworks and notation to represent several previously defined consistency models [Adve and Hill 1993; Adve and Gharachorloo 1996; Bataller and Bernabeu 1997; Mosberger 1993]. Our unified theory is an improvement over these methods because we expose the underlying structure of declarative properties enforced by various models, and we predict new models that have not yet been discovered. There are currently two common methods of characterizing consistency models. One method is to describe restrictions on the way in which processes are allowed to issue memory operations, which we will call the “issue” method (e.g., see Adve and Gharachorloo [1996].) Another method is to describe restrictions on the apparent order of events visible to processes which we will call the “view” method (e.g., see Bataller and Bernabeu [1997].) Adve and Gharachorloo [1996] use the “issue” method of defining consistency models. They identify two conditions that together will enforce sequential consistency. They call these the process order property, and the write atomicity property.

*Process order property.* Program order must be maintained among operations from individual processes.

*Write atomicity property.* In cache based systems with multiple copies of a memory location, writes must be atomic.

The first condition can be enforced by having a process not issue an operation until all previous operations are complete. Complete means that a read has returned its value, or a write has been applied and acknowledged. The second condition can be enforced by a cache coherence protocol. Adve and Gharachorloo [1996] use this implementation of sequential consistency as a basis for their definitions of other consistency models. Every other model is allowed to violate some of the restrictions required for sequential consistency. Violating a restriction allows for optimization in the implementation. They identify five optimizations that may be allowed.

- Allow a read to be issued before a previous write is complete.
- Allow a write to be issued before a previous write is complete.
- Allow a read or write to be issued before a previous read is complete.
- Allow a read to view another process’ write before the write is applied everywhere.
- Allow a read to view one’s own write before the write is applied everywhere.

The first optimization combined with the last two result in processor consistency as it was defined for the DASH multiprocessor [Lenoski et al. 1990]. All five optimizations combined result in slow consistency [Hutto and Ahamad 1990] which is used for nonsynchronizing operations in synchronized consistency models such as release and weak consistency. For each consistency model, Adve and Gharachorloo describe a “safety net” which would enforce sequential consistency on top of that model. A safety net consists of a programming paradigm where certain operations are replaced with special purpose synchronization operations such as test and set or acquire/release. They also describe the concept of a programmer centric framework. In this framework, for any consistency model, a programmer can determine what synchronizations must be performed for a program to simulate sequential consistency on top of that model.

The goal of consistency models in this view is to simulate sequential consistency with an efficient implementation. The trade off is speed versus the complexity exposed to the programmer. This work does not attempt to characterize the order of events as seen by any particular process in a non-sequential execution. Instead, they characterize what changes a programmer must make to a program to simulate sequential consistency. Other work taking this view has been done to present an efficient, sequentially consistent interface to the programmer through instruction level parallelism and speculative execution [Gniady et al. 1999; Ranganathan et al. 1997a, 1997b].

We believe that using weaker consistency models solely to simulate sequential consistency with an efficient implementation should not be the only goal of shared memory research. Our work is based on the idea of declarative consistency properties weaker than sequential consistency, but still intuitively useful to programmers. Therefore, we find the formalisms of the “issue” method less useful to us.

An alternative to the “issue” method is the “view” method where each process has a view of the order of events in the system. For example, PRAM consistency [Lipton and Sandberg 1988] states that each process must see all operations to occur in an order that respects program order, but different processes may see different orders. This essentially places restrictions on when operations may become visible to other processes, and not on when they may be issued. For our purposes, the view method of defining consistency models is most appropriate. What matters is the possible orders of events from the process’ (programmer’s) point of view. The programmer does not care how the shared memory is implemented. In this view, optimizations such as local bypassing, write atomicity, etc. are memory implementation issues, and not part of the consistency model. If two different implementations produce the same set of possible views they should be considered equivalent. For this reason, our work uses view based definitions of consistency models. We believe they are more independent of implementation details. Several surveys of view based definitions have been presented in the literature [Bataller and Bernabeu 1997; Mosberger 1993; Tanenbaum 1995]. We review these view based definitions in Sections 2.2 and 2.3.

The only prior comparison in the literature of the issue and view methods is by Ahamad et al. [1993]. In their paper, they compare Goodman’s definition of processor consistency (which is view based) to the DASH definition (which is issue based.) Their conclusion was that both definitions are weaker than sequential consistency, and stronger than both PRAM and cache consistency. This is the strength relationship commonly understood for processor consistency, and the two models have often been considered equivalent. However, Ahamad et al. showed that the

two definitions are not equivalent, and are in fact incomparable. This showed that it is not trivial to compare consistency models defined under the two formalisms. More work relating the two formalisms is needed. This paper concentrates on view based definitions. Generally, issue based definitions have a view based definition that is analogous.

The most closely related work to this article is the Mume project [Bataller and Bernabeu-Auban 1998] which specifies three consistency properties (orderings): total order, total order with mutual exclusion, and causal order. The Mume project showed that these orderings can be used to provide an alternative and equivalent specification of existing consistency models. However, unlike our work, there is no notion of combining properties in arbitrary ways to produce a lattice of consistency models, or of consistency transitions within that lattice.

**2.2. CONSISTENCY MODEL DEFINITIONS.** Leslie Lamport defined sequential consistency [Lamport 1979]:

*Definition 2.1.* A multiprocessor is *Sequentially Consistent* if the result of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by its program.

Lamport also gave two implementation requirements which, if met, would enforce sequential consistency.

*R1.* Each processor issues memory requests in the order specified by its program.

*R2.* Memory requests from all processors issued to an individual memory module are serviced from a single FIFO queue. Issuing a memory request consists of entering the request on this queue.

Linearizability [Herlihy and Wing 1990] also called atomic memory [Lamport 1986] is essentially sequential consistency with a real-time constraint. Each operation is given a begin time and end time in reference to a global Newtonian clock. For an execution to be linearizable, it must be sequentially consistent, and the sequential total order must correspond to an order realizable by placing each operation at a single point in time between its begin and end times. Essentially, if two operations' time spans do not overlap they cannot be re-ordered even in the absence of any other dependency. Even though linearizability is stronger, sequential consistency is the strongest consistency model used in practice [Adve and Gharachorloo 1996; Tanenbaum 1995]. Sequential consistency is considered strong enough for conventional reasoning about the correctness of shared memory programs.

Lipton and Sandberg [1988] defined PRAM (Pipelined RAM) consistency, and Goodman [1989] defined cache consistency:

*Definition 2.2.* A multiprocessor is *PRAM Consistent* if writes performed by a single process are seen by all other processes in the order in which they were issued, but writes from different processes may be seen in different orders by different processes.

*Definition 2.3.* A multiprocessor is *Cache Consistent* if all writes to the same memory location are performed in some sequential order.

In the same paper, Goodman defined processor consistency.




FIG. 2. An execution that is processor, but not sequentially consistent.

**Definition 2.4.** A multiprocessor is *Processor Consistent* if it is PRAM consistent and writes to the same memory location are seen in the same sequential order by all processes.

One consistency model is said to be stronger than another if every condition required by the weaker model is also required by the stronger one. Thus, a stronger consistency model has a more highly constrained behavior than a weaker one. Sequential consistency is strictly stronger than processor consistency which is strictly stronger than both PRAM and cache consistency. PRAM and cache consistency are incomparable. PRAM and cache consistency are very similar to Lamport's conditions R1 and R2, enforcing R1 and R2 enforces sequential consistency, processor consistency enforces PRAM consistency and cache consistency, but processor consistency is weaker than sequential consistency. How can this be?

Consider Figure 2. In this figure, time proceeds from left to right, and variables are assumed to have an initial value of  $\perp$ . Process  $p_1$  writes to  $x$ , and then reads from  $y$ . Likewise, process  $p_2$  writes to  $y$ , and then reads from  $x$ . Both processes read the initial value of the variable instead of each other's write. Both processes perceive that their write went first so the execution is not sequential. However, it is processor consistent. There is only one write by  $p_1$  and one by  $p_2$  so it is trivially PRAM consistent. There is only one write to  $x$  and one write to  $y$  so it is trivially cache consistent. This example demonstrates how processor consistency is weaker than sequential consistency. Writes by different processes to different variables may be seen to occur in different orders.

The question remains, does the execution in Figure 2 satisfy R1 and R2? The answer is no because R2 requires that read operations be placed in the queue along with write operations. Neither process can place its read operation in the queue until its write operation has been placed in the queue so at least one of the processes must read the other's write. On the other hand, processor consistency only requires that write operations become visible in the correct order. The write operations can be pending while each process does its read, and then the write operations are applied in the correct order.

Causal memory [Ahamad et al. 1991] is a consistency model drawn from Lamport's concept of potential causality [Lamport 1978]. Causal memory is weaker than sequential consistency, stronger than PRAM, and incomparable to processor and cache consistency. It was defined by Ahamad et al. [1991] as:

**Definition 2.5.** A multiprocessor is *Causally Consistent* if for each process the operations of that process plus all writes known to that process appear to that process to occur in a total order that respects potential causality. Potential causality is as defined by Lamport [1978] with writes interpreted as sends and reads interpreted as receives.

Slow consistency [Hutto and Ahamad 1990] is weaker than both PRAM and cache consistency.

*Definition 2.6.* A multiprocessor is *slow consistent* if reads must return some value that has been previously written to the location being read. Once a value has been read, no earlier writes to that location (by the processor that wrote the value read) can be returned. Writes by a process must be immediately visible to itself.

Local consistency [Bataller and Bernabeu 1997] refers to the weakest consistency model for shared memory.

*Definition 2.7.* A multiprocessor is *Locally Consistent* if each process' own operations appear to occur in the order specified by its program. There is no restriction on the order in which writes by other processes appear to occur, and different processes may see different orders.

It is important to note that every consistency model is stronger than local consistency and weaker than sequential consistency which is weaker than linearizability [Herlihy and Wing 1990]. This fact implies that consistency models could be placed in a lattice.

**2.3. SYNCHRONIZED CONSISTENCY MODELS.** Some consistency models include explicit synchronization actions which are treated differently than ordinary memory operations. Synchronization operations are processed at a high level of consistency, usually sequential consistency. Ordinary operations are processed at a low level of consistency, usually slow consistency, but the presence of synchronization operations places additional ordering restrictions on ordinary operations. Dubois et al. [1986] defined weak consistency.

*Definition 2.8.* A multiprocessor is *Weak Consistent* if:

- (1) Accesses to global synchronizing variables are strongly ordered [sequentially consistent].
- (2) No access to a synchronizing variable is issued in a processor before all previous global data accesses have been performed.
- (3) No access to global data is issued by a processor before a previous access to a synchronizing variable has been performed.

An ordinary operation is issued either before or after a synchronization operation. All processes must see the ordinary operation occur in this order with respect to the synchronization operation. This provides a sufficient programming environment for constructs such as critical sections and barriers. For example, a barrier is defined to be a synchronization operation, and all operations issued before the barrier must appear to occur before the barrier. However, this condition is sometimes stronger than necessary. Synchronizing operations can be used just to import information, as with the acquiring of a lock, or just to export information, as with the release of a lock. Taking advantage of this as an opportunity for optimization leads to a different consistency model called release consistency [Gharachorloo et al. 1990].

*Definition 2.9.* A multiprocessor is *Release Consistent* if:

- (1) Before an ordinary LOAD or STORE access is allowed to perform with respect to any other processor, all previous *acquire* accesses must be performed.
- (2) Before a *release* access is allowed to perform with respect to any other processor, all previous ordinary LOAD and STORE accesses must be performed.






FIG. 3. An execution that is location, but not entry consistent.

- (3) *Special accesses* [including acquire and release] are sequentially consistent with respect to one another.

A process performs an acquire to get up-to-date information. Only that process is guaranteed to be up-to-date, and then only up to the point of the latest release on every other process. A different implementation called lazy release consistency [Keleher et al. 1992] enforces the same consistency model, but sends updates as late as possible. The distinction between release and weak consistency is that release forces the program to give more detailed instructions on what must be up-to-date at a synchronization. This trend is continued with entry consistency [Bershad and Zekauskas 1991] and scope consistency [Iftode et al. 1996]. In entry consistency [Bershad and Zekauskas 1991] each synchronization variable is associated with one or more ordinary variables. Acquires and releases only bring up-to-date those ordinary variables associated with a particular synchronization variable. The Java Memory Model [Gosling et al. 1996; Roychoudhury and Mitra 2002] is equivalent to entry consistency because locks are associated with objects, and locally cached values are only required to be read from or flushed to a global main memory when a lock is acquired or released. In scope consistency [Iftode et al. 1996] this set of variables is not static, but rather any ordinary variables accessed between an acquire and release of a synchronization variable must be brought up-to-date to the point of the release on all subsequent acquires of the same synchronization variable.

A final synchronized model called location consistency [Gao and Sarkar 2000] is significantly different. location consistency is similar to entry consistency in that each ordinary variable is associated with a synchronization variable, and a release or acquire is ordered with an ordinary operation if their variables are associated. However, location consistency is different in that it allows the state of a variable to be a partial order, and not a total order.

For example, in Figure 3 two processes both write to the variable  $x$ . In entry consistency, the order of these two writes is undefined. They could be seen to occur in either order, and two different processes do not have to agree on the order. However, there is an implicit assumption that for a single process the two operations occur in some order, and the second one overwrites the first. So, when  $p_2$  reads 1 from  $x$  one can deduce that the order seen by  $p_2$  is:

$$(w, p_2, x, 2) < (w, p_1, x, 1) < (r, p_2, x, 1).$$

Therefore,  $p_2$  will never again read from  $(w, p_2, x, 2)$  because it has been overwritten. The operation  $(r, p_2, x, 2)$  violates entry consistency, but not location consistency. Location consistency assumes that each process sees a partial order of writes, and any read can return the value of any write that is not dominated by another write. Writes are only ordered when they are by the same process, or when they are separated by a release-acquire pair. Therefore, under location consistency  $p_2$  can continue forever alternately reading the values 1 and 2 from  $x$  (barring further write, acquire, or release operations). The purpose of location consistency is that

if a program separates every pair of competing writes with a release-acquire pair (called a data-race-free program), then it is equivalent to entry consistency, but still might be able to take advantage of the location model for efficiency optimizations.

Dag consistency [Blumofe et al. 1996] was developed to eliminate the unbounded nature of this non-determinism in location consistency. In his thesis, Frigo [1997] proved that this could not be done without strengthening the model all the way to cache consistency, and he argued for cache consistency as the weakest reasonable memory model. This is interesting because models like PRAM and slow consistency which are incomparable to and weaker than cache consistency respectively have proved very useful in practice. For example, under PRAM consistency if process 1 writes to  $x$ , and process 2 writes a different value to  $x$ , then processes might disagree on which went first, and thus have different final values of  $x$ . This situation where different processes disagree on the value of a variable may continue indefinitely. In this sense, PRAM does not have the “confines nondeterminism” property defined by Frigo.

However, a simple solution to this problem exists: don’t include these dangerous situations in your programs when using these memory models. When using PRAM consistency don’t have two processes write to the same variable without some other control flow dependence in place to determine the order. For example, the common scientific computing paradigm of compute-barrier-communicate can easily be implemented under PRAM consistency. To implement the barrier have one variable per process with an initial value of zero. When a process arrives at its barrier it sets its variable to one. When all variables are one everyone has reached the barrier. Then PRAM consistency enforces the requirement that all variables written to in the compute phase will be available in the communication phase. Likewise, slow consistency has mostly been useful for ordinary operations in synchronized consistency models where the synchronization operations provide the necessary control and data dependencies. So these consistency models are still useful if programmed correctly.

### 3. A Formalism for Shared Memory Consistency Models

This section presents formal, declarative definitions of the well known consistency models introduced in Section 2. When a shared memory implementation satisfies a particular consistency model it must produce only executions acceptable to that model. In this way, a consistency model can be thought of as a criteria to accept or reject program executions. Therefore, a model can be defined by specifying its set of accepted executions [Collier 1992]. This is the technique we will use in the rest of the article.

In “view” based definitions of consistency models, memory operations must *appear* to be processed in a certain order. For example, under sequential consistency, there must appear to be a single total order on all operations. Under Cache consistency, there must appear to be a total order on the operations to each variable. Each process sees, through its read operations, a particular order of events in the memory system. However, each process has limited information because it may not read every write. Therefore, there could be many orders of events that would be consistent with the values returned by a process’ reads. If any of these orders satisfies a consistency model, then the process cannot prove that the memory system violated

$$\begin{array}{c|c}
\begin{array}{l}
(w, p_1, x, 1) <_{PO} (r, p_1, x, 2) \\
(w, p_2, x, 2) <_{PO} (r, p_2, x, 1) \\
(w, p_1, x, 1) \mapsto (r, p_2, x, 1) \\
(w, p_2, x, 2) \mapsto (r, p_1, x, 2)
\end{array} &
\begin{array}{l}
(w, p_1, x, 1) <_{PO} (w, p_1, y, 2) \\
(r, p_2, y, 2) <_{PO} (r, p_2, x, \perp) \\
(w, p_1, y, 2) \mapsto (r, p_2, y, 2) \\
(w, \epsilon, x, \perp) \mapsto (r, p_2, x, \perp)
\end{array}
\end{array}$$

(a)
(b)

FIG. 4. Examples for PRAM and cache consistency.

that model. If some acceptable order exists for every process, then the execution must be accepted.

Our formalism is defined in Appendix A. It is taken from Ahamad et al. [1993] and Bataller and Bernabeu [1997], and is very similar to Collier [1992]. Basically, an execution consists of a set of read and write operations, and two orders on those operations. Process order,  $<_{PO}$ , is the order in which operations are submitted by each process. Writes-to order,  $\mapsto$ , defines which write is read by each read. Figure 4 defines two example executions by giving the operation tuples (read or write, process id, variable, value) and the process order and writes-to links between those operations.

A consistency model is defined based on serial views. A serial view is a total order on a subset of the operations in an execution that represents one process' view of the order in which those operations took place in memory. In a serial view, each read must read from the most recent write to the same variable. Each consistency model defines a view or set of views that must exist for the consistency model to be satisfied. If the given views can be constructed for an execution, then that execution is accepted by that consistency model, otherwise the execution is rejected. The set of executions accepted by this test provides the definition of the consistency model.

**THEOREM 3.1.** *An execution is Sequentially Consistent iff*

$$\exists \text{ SerialView}(<_{PO}).$$

**PROOF.** This is merely a restatement of the definition of sequential consistency using our formalism. It states that for an execution to be sequentially consistent there must exist a serial view (a total order in which each read reads from the most recent write) on all operations that respects process order ( $<_{PO}$ ).  $\square$

The actual execution may not have occurred in this order, but the values returned by the reads are exactly the same as the values that would have been returned had this been the execution order. Therefore, no process external to the memory system can prove that the execution did not actually happen in this order.

**THEOREM 3.2.** *An execution is PRAM Consistent iff*

$$\forall i \in P \exists \text{ SerialView}(<_{PO} \mid (*, i, *, *) \cup (w, *, *, *)).$$

**PROOF.** This is a restatement of the definition of PRAM consistency. PRAM consistency requires that each process see a view where the writes of all process occur in the order in which they were issued (process order), but writes by different processes may be seen in any order. The operations visible to each process are its own reads and all writes. For this reason, the view of process  $i$  is restricted to  $(*, i, *, *)$ , all of its own operations, and  $(w, *, *, *)$ , all writes. If a serial view respecting process order can be constructed for this subset of operations, then this process cannot argue that the memory system has violated PRAM. If such a view

can be constructed for every process then no external observer can argue that the memory system has violated PRAM.  $\square$

**THEOREM 3.3.** *An execution is Cache Consistent iff*

$$\forall_{x \in V} \exists \text{ SerialView}(<_{PO} | (*, *, x, *)).$$

**PROOF.** Once again, this is a restatement of cache consistency. Cache consistency requires that the operations on each variable,  $x$ , must be seen to occur in a total order that respects process order. The views that must be constructed to satisfy the above definition are exactly the total orders required for the original definition of cache consistency.  $\square$

Consider Figure 4. The sets  $P$ ,  $V$ , and  $O$  and the initial writes can usually be deduced from the descriptions of process order and writes-to order. For this reason they will be omitted in this and further examples unless required for clarity. In Figure 4(a), both processes write and then read  $x$ , and both read the other's write. This execution can be shown to be PRAM consistent by the following views.

$$\begin{aligned} p_1 : (w, p_1, x, 1) <_{p_1} (w, p_2, x, 2) <_{p_1} (r, p_1, x, 2) \\ p_2 : (w, p_2, x, 2) <_{p_2} (w, p_1, x, 1) <_{p_2} (r, p_2, x, 1). \end{aligned}$$

This execution is not sequential. One would have to add  $(r, p_2, x, 1)$  to  $<_{p_1}$ , or  $(r, p_1, x, 2)$  to  $<_{p_2}$ . In  $<_{p_1}$ ,  $(r, p_2, x, 1)$  cannot come before  $(w, p_2, x, 2)$  because that would violate process order. It also cannot come after  $(w, p_2, x, 2)$  because then it would be after, but not reading from,  $(w, p_2, x, 2)$  which would violate the serial property. A similar argument can be made for  $<_{p_2}$ . No single view can satisfy both processes so the execution is not sequentially consistent.

In Figure 4(b), process 1 writes to both  $x$  and  $y$  while process 2 reads both  $x$  and  $y$ . Process 2 reads process 1's second write to  $y$  and the initial value of  $x$ . This execution can be shown to be Cache consistent by the following views. Note, the initial writes must be accounted for in all views, but are omitted in examples where their placement is trivial.  $(w, \epsilon, x, \perp)$  is shown in  $<_x$  because its value is later read.

$$\begin{aligned} x : (w, \epsilon, x, \perp) <_x (r, p_2, x, \perp) <_x (w, p_1, x, 1) \\ y : (w, p_1, y, 2) <_y (r, p_2, y, 2) \end{aligned}$$

Figure 4(b) is not sequentially consistent. In a view with every operation,  $(w, p_1, x, 1)$  would have to come before  $(w, p_1, y, 2)$  by process order.  $(w, p_1, y, 2)$  would have to come before  $(r, p_2, y, 2)$  for the view to be serial.  $(r, p_2, y, 2)$  would have to come before  $(r, p_2, x, \perp)$  by process order. This implies  $(r, p_2, x, \perp)$  would come after  $(w, p_1, x, 1)$  but read from the initial write so the view could not be serial.

Also, Figure 4(a) is not Cache consistent, and Figure 4(b) is not PRAM consistent. In Figure 4(a) all operations are on the same variable so there would need to be a serial view on all operations. In disproving sequential consistency, we have already shown this is impossible. For Figure 4(b) to be PRAM consistent the view  $<_{p_2}$  would need to be constructed containing all of  $p_1$ 's writes, and all of  $p_2$ 's operations. This would include all of the operations which have likewise been shown to be impossible.

THEOREM 3.4. *An execution is Processor Consistent iff*

$$\begin{aligned} &\forall_{x \in V} \exists <_x = \text{SerialView}(<_{PO} \mid (*, *, x, *)), \text{ and} \\ &\forall_{i \in P} \exists \text{SerialView}((\cup_{x \in V} <_x) \cup <_{PO} \mid (*, i, *, *) \cup (w, *, *, *)). \end{aligned}$$

PROOF. This is a restatement of processor consistency which says that the execution must be cache consistent, PRAM consistent, and the PRAM serial views must respect the order of operations in the cache serial views. This is equivalent to Goodman's definition of processor consistency.  $\square$

Definition 3.5. *The Causal Relation,  $<_{CR}$ ,*

$$\begin{aligned} &\forall_{o_i, o_j \in O} o_i <_{CR} o_j \text{ iff} \\ &\quad o_i <_{PO} o_j, \text{ or} \\ &\quad o_i \mapsto o_j, \text{ or} \\ &\quad \exists o_k \in O \text{ such that } o_i <_{CR} o_k <_{CR} o_j. \end{aligned}$$

THEOREM 3.6. *An execution is Causally Consistent iff*

$$\forall_{i \in P} \exists \text{SerialView}(<_{CR} \mid (*, i, *, *) \cup (w, *, *, *))$$

PROOF. This is a restatement of causal consistency which says that each processor must see a serial view which respects the causal relation.  $\square$

THEOREM 3.7. *An execution is Slow Consistent iff*

$$\forall_{i \in P, x \in V} \exists \text{SerialView}(<_{PO} \mid (*, i, x, *) \cup (w, *, x, *))$$

PROOF. This is a restatement of slow consistency which says that the set of operations on a single variable must be PRAM consistent.  $\square$

THEOREM 3.8. *An execution is Locally Consistent iff*

$$\forall_{i \in P} \exists \text{SerialView}(<_i \mid (*, i, *, *) \cup (w, *, *, *))$$

PROOF. This is a restatement of local consistency which says that each processor must see a view which respects its own local order.  $\square$

3.1. SYNCHRONIZED CONSISTENCY MODELS. Synchronized consistency models require additional definitions. First of all, operations are divided into two types, ordinary and synchronization operations. In some models, such as weak consistency, reads and writes are designated as synchronization operations. In other models, such as release consistency, synchronization operations are new types of operations, acquire and release. In either case, the operation type  $s$  is used to designate synchronization operations. For example,  $(s, *, *, *)$  designates the set of all synchronization operations whether those are read, write, acquire, or release. Also, we need to explicitly state that the writes-to relation is defined on synchronization operations. For this purpose, acquires are treated as reads, and releases are treated as writes. Essentially, synchronization operations must be aware of which acquire corresponds to which release. Defining the writes-to relation in this way allows the existing definition of serial view to be used for this purpose. Finally, for each synchronized consistency model, certain ordinary operations must come before or after certain synchronization operations.

*Definition 3.9.*  $D_-(s)$  denotes the set of ordinary operations that must come before synchronization operation  $s$ .  $D_+(s)$  denotes the set of ordinary operations that must come after synchronization operation  $s$ .  $<_D$  denotes the relation:

$$\forall o \in D_-(s) o <_D s \cup \forall o \in D_+(s) s <_D o$$

Synchronized consistency models support different consistency for ordinary operations than synchronization operations. For some models, ordinary operations are processed under slow consistency, and for some models under cache consistency. Bataller and Bernabeu [1997] argue that this distinction is not a significant design feature, but rather was primarily an artifact of the implementation for which each model was originally defined. They present formal definitions of all models assuming that ordinary operations are processed under slow consistency. Synchronization operations are generally processed under sequential consistency, although a variation of release consistency was presented where synchronization operations were processed under processor consistency. Below, we assume synchronization operations obey sequential consistency and ordinary operations obey slow consistency. Variations will be dealt with in the section on consistency transitions (see Section 5.)

Every synchronized consistency model obeys the following condition. The only difference between models is in the definition of  $D_-(s)$  and  $D_+(s)$ .

*Definition 3.10.* For a given definition of  $<_D$ , an execution is *synchronized model consistent* iff

$$\begin{aligned} &\exists <_{seq} = \text{SerialView}(<_{PO} \mid (s, *, *, *)), \text{ and} \\ &<_S = \text{the transitive closure of } <_D \cup <_{seq}, \text{ and} \\ &\forall i \in P, x \in V \exists \text{SerialView}(<_S \cup <_{PO} \mid (*, i, x, *) \cup (w, *, x, *)) \end{aligned}$$

Definition 3.10 says that a sequential order exists on all synchronization operations. The per-process, per-variable views required by slow consistency exist. And the slow consistent views respect the transitive closure of the ordering  $<_D$  and the sequential order of synchronization operations. We will now discuss the differences between various synchronized consistency models.

In weak consistency [Dubois et al. 1986] there is only one synchronizing variable, and there is no distinction between acquire and release types of synchronizing operations.  $D_+(s)$  orders after  $s$  any operation ordered after it by process order.  $D_-(s)$  orders before  $s$  any operation ordered before it by process order.

For release consistency [Gharachorloo et al. 1990], there is only one synchronizing variable, but the distinction is made between acquire and release types of synchronizing operations.  $D_+(\text{acquire})$  orders after *acquire* any operation ordered after it by process order.  $D_-(\text{release})$  orders before *release* any operation ordered before it by process order.

Lazy release consistency [Keleher et al. 1992], does not force operations before a release to be ordered before that release, but they must be ordered before any subsequent acquire. There is only one synchronizing variable.  $D_+(\text{acquire})$  orders after *acquire* any operation ordered after it by process order.  $D_-(\text{acquire})$  orders before *acquire* any ordinary operation where there exists *release*  $<_S$  *acquire* such that the ordinary operation is ordered before *release* by process order. No ordinary operations are directly ordered with any release by  $<_D$ .

In entry consistency [Bershad and Zekauskas 1991], there can be more than one synchronization variable. Rather than synchronizing all of memory with each synchronization operation, each synchronization variable is associated with one or more ordinary variables, and only those variables are synchronized. An ordinary operation is ordered with a synchronization operation in the same way it would be by release consistency if and only if their variables are associated.

Scope consistency [Iftode et al. 1996] is similar to entry consistency in that not all of memory is synchronized by each synchronization operation. But there is no static mapping between ordinary and synchronization variables. Instead, whatever ordinary operations are performed between an acquire and a release are associated with those synchronization operations. An ordinary operation is ordered with a synchronization operation in the same way it would be by release consistency if and only if there is no other synchronization operation to the same variable ordered between them by process order. Essentially, ordinary operations are only ordered with respect to the most recent acquire and the next release to each synchronization variable.

Location consistency [Gao and Sarkar 2000] is different, but we will present it in a formalism as close as possible to that used for the other models. One important difference is that in location consistency synchronization operations are defined to provide a mutual exclusion function. If one process performs an acquire, then no other process may successfully perform an acquire until the first process performs a release. All subsequent acquires are ordered after that release. This provides control dependencies in addition to the data dependencies enforced by the consistency model. This exposes a fundamental difference of opinion about what the job of a consistency model should be. For example, under release consistency there is nothing to prevent two processes from both performing acquires and concurrently writing to the same variable. Release consistency specifies formally what data dependencies must be preserved by the memory system in that situation, that is, the writes are unordered and can be seen in different orders by different processes. If the program truly needs mutual exclusion it can be included in the program code as a locking algorithm that works correctly under release consistency [Gharachorloo et al. 1990].

Most synchronized consistency models were written in two parts, the consistency model itself, and a programming paradigm such as properly labeled [Gharachorloo et al. 1990] or data-race-free [Adve and Hill 1993] programs. The guarantee provided is that a program that obeys the programming paradigm executed on the consistency model will simulate sequential consistency. The authors of release consistency expected that it would be used in conjunction with control flow constructs in the program to simulate sequential consistency, but they did not directly embed the control flow into the consistency model. Instead they allowed the programmer to choose the appropriate control flow constructs. They also acknowledged that some programmers may not want to simulate sequential consistency, but rather deal directly with the semantics of release consistency. We take the approach that control dependencies should be dealt with in a programming paradigm, and not in the definition of the consistency model itself. Therefore, our definition of location consistency does not require mutual exclusion between acquires and releases. If a program obeys the mutual exclusion paradigm then our definition is equivalent to location consistency. Our definition is also valid in cases where the program does not obey mutual exclusion, and we believe it is a good extrapolation. But it is not equivalent because location consistency explicitly forbids these cases

Synchronization operations in location consistency are similar to entry consistency in that they are tagged with a variable, and only enforce dependencies with ordinary operations to that variable. The mutual exclusion assumption stated above requires that there is a total order on all synchronization operations to each variable so location consistency enforces at least cache consistency on synchronization operations. However, the description of location consistency [Gao and Sarkar 2000] does not specifically say that synchronization operations must obey sequential consistency. There is no example in the article with synchronization operations to more than one variable so it is difficult to say whether the authors intended synchronization operations to be sequentially consistent, or merely cache consistent. For similarity with previous models, sequential consistency is assumed.

We will now give the definition of the data dependencies implied by location consistency assuming synchronization operations are sequentially consistent without requiring the control dependencies of the mutual exclusion paradigm. In the original definition of location consistency, due to the mutual exclusion requirement there is an alternating order on the synchronization operations to each variable: acquire, release, acquire, release, etc. Each acquire is immediately after a release which is called its *most\_recent\_release*, and immediately before a release by the same process. The state of a variable,  $x$ , is defined to be a partial order,  $<$  which is the union of  $<_{PO} \mid (s, *, x, *) \cup (w, *, x, *)$  and the condition that all acquires to  $x$  are ordered after their *most\_recent\_release*.

Because  $<$  is a partial order there may be many writes that could be considered “most recent” in that there is no other write ordered after them. A read is allowed to return a value written by any one of these most recent writes. More formally,  $<$  is augmented with any process order edges between the read and any operation in  $(s, *, x, *) \cup (w, *, x, *)$  to produce  $<'$ . Then, the read,  $r$ , may return the value of any write,  $w$ , to the variable  $x$  such that  $\nexists w'$  where  $w <' w' <' x$ . To put this in a similar notation as the other synchronized models, the first requirement is the same that synchronization operations must be sequentially consistent.

$\exists <_{seq} = \text{SerialView}(<_{PO} \mid (s, *, *, *)), \text{ and}$   
 $<_S = \text{the transitive closure of } <_D \cup <_{seq} \text{ where } <_D \text{ is defined the same}$   
 as entry consistency

For programs that obey mutual exclusion there is already a total order,  $<_{seq}$ , on the synchronization operations to each variable. So  $<_S$  is merely the transitive closure of process order and *most\_recent\_release* order. Therefore,  $<_S$  is an equivalent definition of  $<$ . For programs that do not obey mutual exclusion This is a sensible extension of the definition of  $<$  that maintains similarity with other synchronized models. Now, location consistency defines the set of values that may be returned by any read. To capture this, we will add to the formalism the notion of a partial-ordered view.

**Definition 3.11.** There exists a *serial partial view* on a set of operations, *subset*, respecting a partial order,  $<$ , denoted  $\text{SerialPartialView}(< \mid \text{subset})$  iff

$$\forall_{w \mapsto r \in O} \nexists w' \text{ such that } w < w' < r.$$

A serial partial view is a minimal order, that is it doesn't add any edges to  $<$ , it just checks if each read reads from a non dominated write. This is unlike a serial view that must add edges to create a total order out of any partial order






FIG. 5. An initial consistency model lattice.

it respects. The order,  $<$ , must still be a partial order. For example, there cannot exist a serial partial view respecting a cyclic relation. Now, location consistency was defined where each read had its own serial partial view. However, if a serial partial view exists separately for two reads over the same set of writes and synchronization operations, then those two reads can be added to the same partial order, and it will still be a serial partial view. There is no interaction between the two reads. Therefore, the condition that all reads read a permissible value can be stated like this.

$$\forall_{i \in P, x \in V} \exists \text{ SerialPartialView}(<_S \cup <_{PO} \mid (*, i, x, *) \cup (w, *, x, *)).$$

Every process has multiple SerialPartialViews, one for each variable. Each of these views includes all of the operations by that process on that variable, and all of the writes by any process on that variable. These views must respect the synchronization order,  $<_S$ , and process order,  $<_{PO}$ , and be serial. That is, each read must read from a nondominated write. Therefore, the definition of location consistency is identical to the definition of entry consistency with SerialView replaced by SerialPartialView.

#### 4. Consistency Properties

Some existing consistency models have been viewed as a combination of other models. For example, processor consistency [Goodman 1989] is the combination of PRAM and Cache consistency. Causal ordering [Ahmad et al. 1991] is the transitive closure of process order and writes-to order. Lamport's [1979] original definition of sequential consistency included a pair of properties which, if independently enforced, would enforce sequential consistency. This suggests that perhaps many existing consistency models could be viewed as different combinations of a few primitive consistency properties. In this section we define four such properties. Global Process Order (GPO) is the condition that there is global agreement on the order of operations from each process. Global Data Order (GDO) is the condition that there is global agreement on the order of operations to each variable. Global Write-read-write Order (GWO) is the condition that there is global agreement on the order of potentially causally related writes. Global Anti Order (GAO) is the condition that there is global agreement on the order of any two writes when a process can prove it has read one before the other. Any combination of these properties results in a consistency model. Enumerating these models results in the lattice shown in Figure 13.

For pedagogical purposes, we will start with the lattice shown in Figure 5 and expand the lattice as properties are developed. The top of the lattice is defined to be sequential consistency, and the bottom is defined to be local consistency as these are the strongest and weakest properties in the literature (see Section 2.)

4.1. PROCESSOR CONSISTENCY AS A COMBINATION OF PROPERTIES. Processor Consistency is defined to be a combination of PRAM and cache consistency (see Theorem 3.4.) The given definition of processor consistency requires constructing per-variable views to satisfy cache consistency in addition to per-process views to satisfy PRAM consistency. To remove this inconvenience, we will define two properties, one equivalent to PRAM consistency, and one equivalent to cache consistency such that both properties can be combined in the same per-process views.

*Definition 4.1.* An execution is *Global Process Order (GPO)* iff

$$\forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{PO} | (*, i, *, *) \cup (w, *, *, *)).$$

**THEOREM 4.2.** *GPO is equivalent to PRAM consistency.*

**PROOF.** The definitions of GPO and PRAM are identical. The views for GPO are required to respect local order for similarity with the properties to follow. However, this requirement is redundant because process order is a superset of local order for any process.  $\square$

*Definition 4.3.* Two operations are ordered by *data order*,  $o_1 <_{DO} o_2$ , iff  $\text{variable}(o_1) = \text{variable}(o_2)$ , and either

- (1)  $o_1 <_{PO} o_2$ , or
- (2)  $o_1 \mapsto o_2$ , or
- (3)  $\exists_{r \in (r, *, *, *)} \text{variable}(r) = \text{variable}(o_1) \wedge \text{value}(r) \neq \text{value}(o_1) \wedge o_1 <_{PO} r \wedge o_2 \mapsto r$ , or
- (4) There exists an operation,  $o$ , such that  $o_1 <_{DO} o <_{DO} o_2$ .

Data order captures the restrictions involved in constructing the required views for cache consistency. The operations  $o_1$  and  $o_2$  can be either reads or writes, but must be to the same variable. The per-variable views for cache consistency must be serial and respect process order so data order contains writes-to order and process order restricted to pairs of operations to the same variable. For the third condition, a particular process reads or writes a value,  $o_1$ , and then at a later time reads a different value from the same variable,  $r$ . That process can deduce that a write,  $o_2$ , must have occurred between those two operations and so the restriction is included in data order. The fourth condition requires that data order is a transitive closure.

*Definition 4.4.* An execution is *Global Data Order (GDO)* iff

$$\forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{DO} | (*, i, *, *) \cup (w, *, *, *)).$$

GDO is equivalent to cache consistency, and an execution is cache consistent if and only if data order is acyclic. Proof of this is in Appendix B. This is an important result because it provides two new ways to define cache consistency. One can determine whether an execution is cache consistent by the original method of constructing per-variable serial views, or now by constructing per-process serial views, or even by testing the cyclicity of the data order relation. Now that cache consistency is defined over per-process views we can combine GPO and GDO more easily.

*Definition 4.5.* An execution is *GPO + GDO* iff

$$\forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{PO} \cup <_{DO} | (*, i, *, *) \cup (w, *, *, *)).$$

$$\begin{aligned}
(w, p_1, x, 1) &<_{PO} (w, p_1, z, 2) <_{PO} (r, p_1, y, \perp) \\
(w, p_2, y, 3) &<_{PO} (w, p_2, z, 4) <_{PO} (r, p_2, x, \perp) \\
(w, \epsilon, x, \perp) &\mapsto (r, p_2, x, \perp) \\
(w, \epsilon, y, \perp) &\mapsto (r, p_1, y, \perp)
\end{aligned}$$

FIG. 6. A GPO + GDO, but not processor consistent execution.

However, GPO + GDO is not quite equivalent to Goodman's definition of processor consistency. Processor consistency requires that all processes agree on a total order of operations to each variable. In Figure 6, the processes cannot agree on the order of the writes to  $z$ . If  $(w, p_1, z, 2)$  was first, then  $p_2$  should have read 1 from  $x$ . Likewise, if  $(w, p_2, z, 4)$  was first, then  $p_1$  should have read 3 from  $y$ . However, the two writes to  $z$  are not ordered by process order or data order. Under processor consistency they are allowed to occur in either order, but they are required to be seen in the same order by all processes. GPO + GDO does not enforce this. The problem can be solved by creating augmented data order,  $<_{DO'}$ . Augmented data order is any superset of data order that enforces a total order on all operations to each variable. Because GDO is equivalent to causal consistency, any GDO execution respects at least one augmented data order. The problem is that there may be more than one, and a single augmented data order may not be consistent with process order at all sites. GPO + GDO' is defined similarly to GPO + GDO.

**THEOREM 4.6.** *Goodman's definition of processor consistency (as given in Section 2.2) is equivalent to GPO + GDO'.*

**PROOF.** Augmented data order is equivalent to the per-variable cache consistency views required for processor consistency. The per-process views for GPO + GDO' respect process order and augmented data order. The per-process views for processor consistency respect process order, and the per-variable cache consistent views. Therefore, the two required sets of views are equivalent.  $\square$

Augmented data order solves the problem of equivalence to Goodman's definition of processor consistency. However, we feel that even without augmented data order GPO + GDO is in line with the intended purpose of process order. In Figure 6, the writes to  $z$  are unordered. Inserting reads to  $z$  to detect the order of the writes would create data order dependencies and eliminate the need for augmented data order.

Alternatively, GPO + GDO is equivalent to the following modified definition of processor consistency where the  $\forall_{i \in P}$  is moved outside of the  $\forall_{x \in V}$ , and each process respects a set of cache consistent views, but all processes do not have to respect the same set of views.

$$\begin{aligned}
&\forall_{i \in P} \forall_{x \in V} \exists <_x = \text{SerialView}(<_{PO} | (*, *, x, *)), \text{ and} \\
&\exists \text{SerialView}((\bigcup_{x \in V} <_x) \bigcup <_{PO} | (*, i, *, *) \cup (w, *, *, *))
\end{aligned}$$

Also, the execution in Figure 6 was taken from Ahamad et al. [1993] as an example of an execution accepted by the DASH definition of processor consistency, and rejected by Goodman's definition. We believe that GPO + GDO is a useful definition of processor consistency.

This issue is discussed in more detail in Section 5. The same issue comes up when defining synchronized consistency models as consistency transitions. The synchronization operations must be sequentially consistent, but there may be more than one

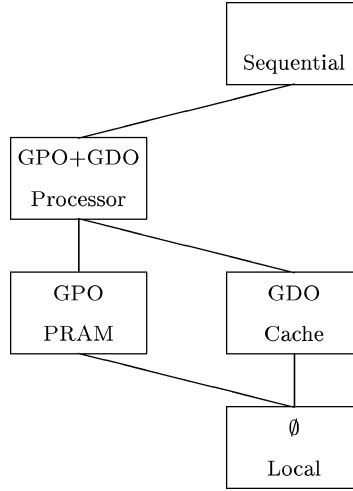


FIG. 7. A consistency model lattice including processor consistency.

$$\begin{aligned}
 (w, p_1, x, 1) &<_{PO} (w, p_1, y, 2) \\
 (r, p_2, y, 2) &<_{PO} (w, p_2, x, 3) \\
 (r, p_3, x, 3) &<_{PO} (r, p_3, x, 1) \\
 (w, p_1, x, 1) &\mapsto (r, p_3, x, 1) \\
 (w, p_1, y, 2) &\mapsto (r, p_2, y, 2) \\
 (w, p_2, x, 3) &\mapsto (r, p_3, x, 3)
 \end{aligned}$$

FIG. 8. A PRAM and cache, but not GPO + GDO consistent execution.

total order that would satisfy sequential consistency. The ordinary operations are not required to be sequentially consistent, and may demonstrate that different processes saw different sequential orders even though the synchronization operations in isolation are sequentially consistent.

GPO + GDO begins a framework for defining consistency properties (see Figure 7.) Any property that can be defined as a relation which must be respected by per-process views can be combined with process order and data order to create new consistency models.

There can also be executions that are GPO and GDO, but not GPO + GDO. Ahamad et al. [1993] provide the execution in Figure 8, which is PRAM and cache consistent, but not processor consistent. The execution is GPO because of the following views.

$$\begin{aligned}
 p_1 : (w, p_2, x, 3) &<_{p_1} (w, p_1, x, 1) <_{p_1} (w, p_1, y, 2) \\
 p_2 : (w, p_1, x, 1) &<_{p_2} (w, p_1, y, 2) <_{p_2} (r, p_2, y, 2) <_{p_2} (w, p_2, x, 3) \\
 p_3 : (w, p_2, x, 3) &<_{p_3} (r, p_3, x, 3) <_{p_3} (w, p_1, x, 1) <_{p_3} (r, p_3, x, 1) <_{p_3} \\
 &(w, p_1, y, 2)
 \end{aligned}$$

Data order is as follows:

$$\begin{aligned}
 (w, p_2, x, 3) &<_{DO} (r, p_3, x, 3) <_{DO} (w, p_1, x, 1) <_{DO} (r, p_3, x, 1) \\
 (w, p_1, y, 2) &<_{DO} (r, p_2, y, 2)
 \end{aligned}$$

$$\begin{aligned}
&(w, p_1, x, 1) <_{PO} (r, p_1, y, 3) <_{PO} (r, p_1, x, 1) \\
&(r, p_2, x, 1) <_{PO} (w, p_2, x, 2) <_{PO} (w, p_2, y, 3) \\
&(w, p_1, x, 1) \mapsto (r, p_1, x, 1) \\
&(w, p_1, x, 1) \mapsto (r, p_2, x, 1) \\
&(w, p_2, y, 3) \mapsto (r, p_1, y, 3)
\end{aligned}$$

FIG. 9. An execution that violates causal consistency.

The execution is GDO because of the following views:

$$\begin{aligned}
p_1 : & (w, p_2, x, 3) <_{p_1} (w, p_1, x, 1) <_{p_1} (w, p_1, y, 2) \\
p_2 : & (w, p_1, y, 2) <_{p_2} (r, p_2, y, 2) <_{p_2} (w, p_2, x, 3) <_{p_2} (w, p_1, x, 1) \\
p_3 : & (w, p_2, x, 3) <_{p_3} (r, p_3, x, 3) <_{p_3} (w, p_1, x, 1) <_{p_3} (r, p_3, x, 1) <_{p_3} \\
& (w, p_1, y, 2)
\end{aligned}$$

However, in  $<_{p_2}$ , the position of  $(w, p_1, x, 1)$  is different between the GPO and GDO views. There is no view  $<_{p_2}$  that conform to both  $<_{PO}$  and  $<_{DO}$ .

$$\begin{aligned}
&(w, p_1, x, 1) <_{PO} (w, p_1, y, 2) <_{DO} (r, p_2, y, 2) <_{PO} (w, p_2, x, 3) <_{DO} \\
&(w, p_1, x, 1)
\end{aligned}$$

so  $<_{DO} \cup <_{PO}$  has a cycle.  $<_{p_2}$  must contain all of these operations and thus cannot be constructed. This leads to the definition of another consistency model.

**Definition 4.7.** An execution is  $GPO \cap GDO$  iff

$$\begin{aligned}
&\forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{PO} | (*, i, *, *) \cup (w, *, *, *)) \bigwedge \\
&\forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{DO} | (*, i, *, *) \cup (w, *, *, *))
\end{aligned}$$

Any pair of properties can be combined in this way creating a new consistency model. The meaning of these models has not been explored previously in the literature, and we have not explored them in our work. They are mentioned here for completeness.

**4.2. CAUSAL CONSISTENCY AS A COMBINATION OF PROPERTIES.** Causal consistency is stronger than GPO, but incomparable to GPO + GDO. Therefore, there should be a property that enforces that part of causal not already covered by GPO. Causal consistency depends on the causal relation which is the transitive closure of process order and writes-to order. The causal relation is made up of three types of edges: edges in process order, edges in writes-to order, and edges not in either order, but in the transitive closure. Process order has already been identified as a primitive property, and any serial view respects writes-to order. Therefore, we now define another property which contains the edges in the transitive closure. This new property can be used with process order to define causal consistency.

For example, Figure 9 contains an execution that is not causally consistent even though the following serial views respect both process order and writes-to order.

$$\begin{aligned}
p_1 : & (w, p_2, x, 2) <_{p_1} (w, p_2, y, 3) <_{p_1} (w, p_1, x, 1) <_{p_1} (r, p_1, y, 3) <_{p_1} \\
& (r, p_1, x, 1) \\
p_2 : & (w, p_1, x, 1) <_{p_2} (r, p_2, x, 1) <_{p_2} (w, p_2, x, 2) <_{p_2} (w, p_2, y, 3)
\end{aligned}$$

There is a causal dependency from  $(w, p_1, x, 1)$  to  $(w, p_2, x, 2)$  because

$$(w, p_1, x, 1) \mapsto (r, p_2, x, 1) <_{PO} (w, p_2, x, 2).$$

1.	$o_1$	$\mapsto$	$r_1$	$<_{PO}$	$r_2$	$\mapsto$	$o_2$
2.	$o_1$	$\mapsto$	$r_1$			$\mapsto$	$o_2$
3.	$o_1$	$<_{PO}$	$r_1$	$<_{PO}$	$r_2$	$\mapsto$	$o_2$
4.	$o_1$	$<_{PO}$	$r_1$			$\mapsto$	$o_2$
5.	$o_1$	$<_{PO}$	$r_1$	$<_{PO}$	$r_2$	$<_{PO}$	$o_2$
6.	$o_1$	$<_{PO}$	$r_1$			$<_{PO}$	$o_2$
7.	$o_1$	$\mapsto$	$r_1$	$<_{PO}$	$r_2$	$<_{PO}$	$o_2$
8.	$o_1$	$\mapsto$	$r_1$			$<_{PO}$	$o_2$

FIG. 10. Enumerated possibilities for a causal transitive chain.

However,  $<_{p_1}$  places them in the opposite order because  $<_{p_1}$  does not contain the operation  $(r, p_2, x, 1)$  which is a read operation by  $p_2$ . Therefore, it violates neither process order nor writes-to order among the operations in its view. To be causally consistent, the view for each process must respect:

$$(\text{the transitive closure of } <_{PO} \cup \mapsto) | (*, i, *, *) \cup (w, *, *, *).$$

The definition for GPO already respects:

$$\text{the transitive closure of } (<_{PO} \cup \mapsto) | (*, i, *, *) \cup (w, *, *, *)$$

Note the different parentheses. The new property can be found in the set difference of these two relations. For an edge to be in the first relation and not the second, two operations in  $(*, i, *, *) \cup (w, *, *, *)$  must be transitively ordered by a chain of operation not in  $(*, i, *, *) \cup (w, *, *, *)$ . The only operations not in that set are reads by a process other than  $i$ . Reads cannot be ordered with each other by writes-to order, and if a chain of reads is ordered by process order they must all be by the same process, and the first and last reads in the chain will be ordered. So, any transitive chains of the sort we are interested in must have an operation,  $o_1$  ordered by process order or writes-to order before a read,  $r_1$ , possibly ordered by process order before another read,  $r_2$ , ordered by process order or writes-to order before an operation,  $o_2$ . All possibilities are summarized in Figure 10.

Cases 1, 2, 3, and 4 are impossible because a read cannot be on the left hand side of a writes-to relation. In cases 5 and 6, the two operations,  $o_1$  and  $o_2$ , are already ordered by process order. In case 7,  $r_1$  and  $o_2$  are ordered by process order so it reduces to case 8. Therefore, the only case that must be considered is case 8.

In case 8,  $o_1$  must be a write because it writes to  $r_1$ .  $o_2$  is in the set  $(*, i, *, *) \cup (w, *, *, *)$ .  $r_1$  is not in this set and so is not by process  $i$ .  $o_2$  is by the same process as  $r_1$  so it must be a write by another process. Therefore, only causal chains between two writes must be considered.

**Definition 4.8.** Two writes are ordered by *write-read-write order*,  $w_1 <_{WO} w_2$ , iff there exists a read,  $r$  such that  $w_1 \mapsto r <_{PO} w_2$ .

**Definition 4.9.** An execution is *Global Write-read-write Order (GWO)* iff

$$\forall i \in P \exists \text{ SerialView}(<_i \cup <_{WO} | (*, i, *, *) \cup (w, *, *, *))$$

**THEOREM 4.10.** *GPO + GWO is equivalent to causal consistency.*

**PROOF.** By the logic above, the transitive closure of  $<_{PO} \cup <_{WO} \cup \mapsto$   $| (*, i, *, *) \cup (w, *, *, *)$  is equivalent to  $<_{CR} | (*, i, *, *) \cup (w, *, *, *)$ . Any serial view respects  $\mapsto$ , and a view is a total order so if it respects a relation it respects the transitive closure of that relation. Also, any view that respects  $<_{PO}$  respects  $<_i$ .

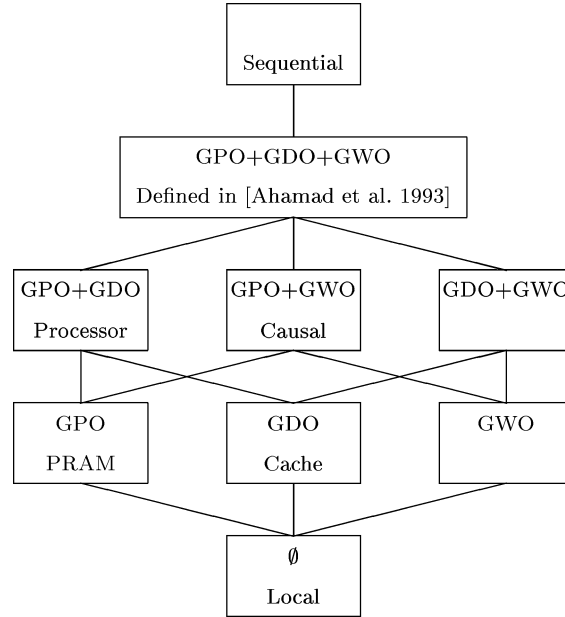


FIG. 11. A consistency model lattice including causal consistency.

$$\begin{aligned}
 &(w, p_1, x, 1) <_{PO} (r, p_1, y, \perp) <_{PO} (r, p_1, y, 2) \\
 &(w, p_2, y, 2) <_{PO} (r, p_2, x, \perp) <_{PO} (r, p_2, x, 1) \\
 &(w, \epsilon, y, \perp) \mapsto (r, p_1, y, \perp) \\
 &(w, p_2, y, 2) \mapsto (r, p_1, y, 2) \\
 &(w, \epsilon, x, \perp) \mapsto (r, p_2, x, \perp) \\
 &(w, p_1, x, 1) \mapsto (r, p_2, x, 1)
 \end{aligned}$$

FIG. 12. An execution that violates sequential consistency.

So a serial view respects  $<_i \cup <_{PO} \cup <_{WO} \mid (*, i, *, *) \cup (w, *, *, *)$  iff it respects  $<_{CR} \mid (*, i, *, *) \cup (w, *, *, *)$ . The first is the requirement for GPO + GWO. The second is the requirement for causal consistency.  $\square$

Adding GWO to the evolving lattice of consistency models results in Figure 11. The model GPO + GDO + GWO has been previously discovered. In Ahamad et al. [1993] the authors noticed that the definition of processor consistency allows executions that violate causality, and they developed an extension to processor consistency to prevent this. At this point, the lattice contains two new consistency models: GWO, and GDO + GWO.

**4.3. SEQUENTIAL CONSISTENCY AS A COMBINATION OF PROPERTIES.** GPO + GDO + GWO is weaker than sequential consistency. Consider the execution in Figure 12. The two writes are not by the same processor, nor to the same variable, and they are not causally related. These two writes could be seen to occur in either order, but to be sequentially consistent every process must see them in the same order. In this execution the following cycle exists:

$$\begin{aligned}
 &(w, p_1, x, 1) <_{PO} (r, p_1, y, \perp) <_{DO} (w, p_2, y, 2) <_{PO} (r, p_2, x, \perp) <_{DO} \\
 &(w, p_1, x, 1)
 \end{aligned}$$

But GPO + GDO + GWO requires separate views for processes  $p_1$  and  $p_2$ , and each view includes only its own read operations. So the following views are acceptable:

$$\begin{aligned} p_1 : (w, p_1, x, 1) <_{p_1} (r, p_1, y, \perp) <_{p_1} (w, p_2, y, 2) <_{p_1} (r, p_1, y, 2) \\ p_2 : (w, p_2, y, 2) <_{p_2} (r, p_2, x, \perp) <_{p_2} (w, p_1, x, 1) <_{p_2} (r, p_2, x, 1) \end{aligned}$$

For this execution to be prohibited, there must be another order that takes a cycle that includes read operations and creates a cycle among only write operations. In Figure 10, there were eight cases of a causal transitive chain. Four of them were deemed impossible because a read could not be on the left-hand side of a writes-to order. These cases are made possible by using data order as a generalization of writes-to order. A read may not be able to write to another operation, but it may be able to prove that it happened first. These four cases are the basis of a new consistency property called *anti-order*. The name anti order comes from parallel compiler optimization. When a program contains a read and later a write to the same variable, their orders cannot be reversed. This is called an anti-dependency, and is similar to this situation where a read can prove, through data order, that a write happened after it. It is at least similar enough to borrow the name.

The purpose of Global Anti-Order (GAO) is to complete the set of consistency properties so that, together, they simulate sequential consistency. To do this, anti order must take cycles involving read operations, and short circuit them to produce cycles involving only write operations. Therefore, anti order is limited to the case where  $o_1$  and  $o_2$  (in Figure 10) are writes. This weakens anti order, and our desire is to produce the weakest relation that supports the assertion that GPO + GDO + GWO + GAO is equivalent to sequential consistency. From Figure 10, case 1 seems necessary because the writes may only be ordered through the reads. Case 2 seems unnecessary because the writes are already ordered by data order, but it will be needed as explained later. Case 3 reduces to case 4, and case 4 solves the problem of Figure 12 since

$$\begin{aligned} (w, p_1, x, 1) <_{PO} (r, p_1, y, \perp) <_{DO} (w, p_2, y, 2), \text{ and} \\ (w, p_2, y, 2) <_{PO} (r, p_2, x, \perp) <_{DO} (w, p_1, x, 1), \text{ so} \\ (w, p_1, x, 1) <_{AO} (w, p_2, y, 2) <_{AO} (w, p_1, x, 1) \end{aligned}$$

So, an initial idea is to base anti order on only cases 1 and 4. However, this solution is not complete. The execution in Figure 12 can be modified by removing the final read of each process. This means that condition 3 of data order no longer applies and the writes are not data ordered after  $(r, p_1, y, \perp)$  and  $(r, p_2, x, \perp)$ . There is no anti-order cycle, and the execution is no longer rejected by anti order even though it still violates sequential consistency. The problem is with a limitation of data order. If a read does not read from a write to the same variable this is not enough to deduce that the write happened after the read. It could have happened very early and been overwritten. However, it could not have happened between the read and the write that wrote-to the read. This ordering restriction is not present in data order. Capturing this restriction requires a nondeterministic order called *serial order*. One can think of serial order as “pseudo-data order” that can replace writes-to order in the cases given in Figure 10. We now need to include case 2 because  $w_1 \mapsto r_1 <_{SO} w_2$  does not guarantee that the writes are data ordered.



*Definition 4.11.* A *Serial Order*,  $<_{SO}$ , is a minimal set of edges that enforces the following condition:

$\forall_{w,r \in O}$  such that  $w$  and  $r$  are to the same variable and do not have the same value either  $w <_{SO} w' \mapsto r$  or  $r <_{SO} w$

So the final definition of anti order is as follows.

*Definition 4.12.* *Anti-Order*,  $<_{AO(<_{SO})}$ ,

Given a serial order,  $<_{SO}$ ,

$\forall_{w_1, w_2 \in O} w_1 <_{AO} w_2$  iff

$\exists r_1, r_2$  such that

$w_1 \mapsto r_1 <_{PO} r_2 <_{DO} w_2$ , or

$w_1 \mapsto r_1 <_{PO} r_2 <_{SO} w_2$ , or

$w_1 \mapsto r_1 <_{SO} w_2$ , or

$w_1 <_{PO} r_1 <_{DO} w_2$ , or

$w_1 <_{PO} r_1 <_{SO} w_2$

To define global anti-order there must be serial views that respect anti order for some definition of serial order. However, this is still not enough. In the example of Figure 12 with the final reads removed serial order could be defined as:

$$(w, p_1, x, 1) <_{SO} (w, \epsilon, x, \perp)$$

$$(w, p_2, y, 2) <_{SO} (w, \epsilon, y, \perp)$$

There would be no anti order links. The views could then be written:

$$p_1 : (w, p_1, x, 1) <_{p_1} (r, p_1, y, \perp) <_{p_1} (w, p_2, y, 2)$$

$$p_2 : (w, p_2, y, 2) <_{p_2} (r, p_2, x, \perp) <_{p_2} (w, p_1, x, 1)$$

These views respect process order, data order, write-read-write order, and even anti order for some definition of serial order. They also respect some definition of serial order, but not the same definition that was used to construct anti-order. This is the crucial last piece of the puzzle. The views must respect the same definition of serial order that was used to construct anti-order.

*Definition 4.13.* An execution is *Global Anti-Order (GAO)* iff  $\exists <_{SO}$  such that

$$\forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{SO} \cup <_{AO(<_{SO})} | (*, i, *, *) \cup (w, *, *, *))$$

Serial order is a nondeterministic order in the sense that it may have many possible definitions, and if any one of the definitions accepts the execution then the execution is accepted. The number of possible serial orders for any execution is not infinite. In fact, for each pair of a read and a write with the same variable and a different value, there is exactly one edge in serial order, and this edge is chosen from two choices. Therefore, the number of serial orders for an execution is exactly  $2^x$  where  $x$  is the number of such read-write pairs. When accepting executions, an implementation of anti-order could be conservative, and only consider a subset of possible serial orders. It could even deterministically chose a single serial order on which to accept executions. This way, the implementation could be more efficient without accepting any unacceptable executions. However, it might reject some acceptable executions. From now on, for purposes of brevity, we will use serial order as if it were a single

order. Any definition using serial order can be read “There exists a serial order such that. . .”

It would be desirable if all four properties were orthogonal, but this is not the case. GAO is strictly stronger than GDO which is proved in Appendix C. One goal of this work was to develop GAO to be as weak as possible while still supporting the assertion that  $GPO + GDO + GWO + GAO$  is equivalent to sequential consistency. Every candidate definition of GAO that we tried that was not stronger than GDO did not support equivalence to sequential consistency. This may reveal some fundamental aspect of consistency models, or it may merely require further research to develop such a definition. As a result,  $GDO + GAO$  is equivalent to just GAO.

All that remains is to show that the four properties together make up sequential consistency. Since GAO is stronger than GDO, we will leave it out and prove that  $GPO + GWO + GAO$  is equivalent to sequential consistency.

LEMMA 4.14. *Every sequentially consistent execution is  $GPO + GWO + GAO$ .*

PROOF. A sequentially consistent execution has a single, serial view on all operations that respects  $<_{PO}$ . Call this view  $<_{seq}$ . By definition,  $<_{seq}$  respects  $<_{PO}$ . If  $w_1 <_{WO} w_2$  then  $\exists r$  such that  $w_1 \mapsto r <_{PO} w_2$ .  $<_{seq}$  respects  $<_{PO}$  and is serial so it respects  $\mapsto$  and therefore it respects  $<_{WO}$ .

Now we will show that a sequentially consistent execution respects  $<_{DO}$ . This is not strictly required by the theorem, but will make it easier to prove that the execution satisfies  $<_{AO}$ .  $<_{seq}$  respects  $<_{PO}$  and is serial, and so respects the  $<_{PO}$  and  $\mapsto$  conditions of  $<_{DO}$ . If  $o_1 <_{PO} r$ , and  $o_2 \mapsto r$ , and  $o_1$  has a different value than  $r$  then  $o_1$  must come before  $o_2$  in  $<_{seq}$ , or the view will not be serial. If this were not so, then  $o_1$  must come between  $o_2$  and  $r$  because  $o_1 <_{PO} r$  and  $<_{seq}$  respects  $<_{PO}$ . There are two cases,  $o_1$  is either a write or a read. If  $o_1$  is a write, then  $r$  does not read from the most recent write and  $<_{seq}$  is not serial. If  $o_1$  is a read, then either  $o_1$  does not read from the most recent write, or there is a write to the same variable with the same value as  $o_1$  between  $o_2$  and  $o_1$  in which case  $r$  does not read from the most recent write and  $<_{seq}$  is not serial. Therefore,  $<_{seq}$  respects Condition 3 of  $<_{DO}$ .  $<_{seq}$  is a total order. Since it respects the first three conditions of  $<_{DO}$ , it will respect the transitive closure condition.

To prove that a sequentially consistent execution is GAO, define a serial order,  $<_{SO}$ , with edges in the same order as  $<_{seq}$ . This is possible because if  $\exists w, w', r$  such that  $w' \mapsto r$  and  $w \neq w'$ , then it cannot be that  $w' <_{seq} w <_{seq} r$  because then  $<_{seq}$  would not be serial.  $w$  must be ordered either before  $w'$  or after  $r$ . If  $\exists w_1, w_2$  such that  $w_1 <_{AO} w_2$ , then  $\exists r_1, r_2$  such that  $w_1 \mapsto r_1 <_{PO} r_2 <_{DO} w_2$ , or  $w_1 \mapsto r_1 <_{PO} r_2 <_{SO} w_2$ , or  $w_1 \mapsto r_1 <_{SO} w_2$ , or  $w_1 <_{PO} r_1 <_{DO} w_2$ , or  $w_1 <_{PO} r_1 <_{SO} w_2$ .  $<_{seq}$  respects  $\mapsto$ ,  $<_{PO}$ ,  $<_{DO}$ , and  $<_{SO}$  so therefore respects  $<_{AO(<_{SO})}$ .

So  $<_{seq}$  respects  $<_{PO}$ ,  $<_{WO}$ ,  $<_{SO}$ ,  $<_{AO(<_{SO})}$ , is serial, and contains all operations so it can be used to construct the required per-process views for all processes:

$$\forall i \in P \exists \text{SerialView}(<_i \cup <_{PO} \cup <_{WO} \cup <_{SO} \cup <_{AO(<_{SO})})$$

$|(*, i, *, *) \cup (w, *, *, *)|$  so the execution is  $GPO + GWO + GAO$ .  $\square$

LEMMA 4.15. *For any  $GPO + GWO + GAO$  execution, the per-process views can be constructed where all write operations occur in the same order in all views.*

PROOF. Because  $<_i$  is a subset of  $<_{PO}$  we will ignore it and just show that the constructed views respect  $<_{PO}$ ,  $<_{WO}$ ,  $<_{SO}$ ,  $<_{AO(<_{SO})}$ , and are serial. There

must be an initial definition of serial order for which the execution satisfies  $GPO + GWO + GAO$ . This definition of serial order is not changed throughout this proof. That is, the final constructed views satisfy  $GPO + GWO + GAO$  for the same definition of serial order as the initial views. All initial writes must be ordered first in all views because all initial writes are ordered before any other operation by  $<_{PO}$ . These initial writes can come in any order because they are not ordered with respect to each other, and there are no reads between them, so place them in the same order in all views. For any two views  $<_i$  and  $<_j$ , the first write that is not an initial write in  $<_i$  can be placed first in  $<_j$ .

Then, the next write in  $<_i$  can be placed second in  $<_j$ , and so on. We will use an inductive proof to show that this reordering can be done and the resulting views will still respect  $<_{PO}$ ,  $<_{WO}$ ,  $<_{SO}$ ,  $<_{AO(<_{SO})}$ , and be serial. The inductive proof uses the following definitions and invariants:

- (1) The order  $<$  is defined as  $<_{PO} \cup <_{WO} \cup <_{SO} \cup <_{AO(<_{SO})}$ .
- (2) The views  $<_i$  and  $<_j$  respect  $<$  and are serial.
- (3) The write operation being moved is called  $w_1$ .
- (4) Point A is the place in  $<_j$  where  $w_1$  will be moved to.
- (5) Point B is the place in  $<_j$  where  $w_1$  is being moved from.
- (6) Point B is after point A in  $<_j$ .
- (7) All write operations ordered before  $w_1$  in  $<_i$  are before point A in  $<_j$ .
- (8) Corollary: All write operations ordered before  $w_1$  by  $<$  are before point A in  $<_j$  because  $<_i$  respects  $<$ .

The execution is  $GPO + GWO + GAO$  so there must exist initial views  $<_i$  and  $<_j$  that respect  $<$  and are serial. In the initial case, point A is just after the initial writes of  $<_j$ .  $w_1$  is the first noninitial write in  $<_i$  so only the initial writes are ordered before it in  $<_i$  and they are all before point A in  $<_j$ .  $w_1$  is after the initial writes in  $<_j$  so point B is after point A in  $<_j$ .

Consider all the operations between A and B. These must all be either read operations by process  $j$ , or write operations not ordered before  $w_1$  by  $<$ . Construct the set of prior reads as follows. The variable that  $w_1$  operates on will be referred to as  $x$ . Any read between A and B to variable  $x$  is a prior read. Also, any read between A and B ordered by process order before  $w_1$  or a prior read is a prior read. Then construct the set of remaining operations as all reads between A and B that are not prior reads plus all writes between A and B. Now, we will show that  $w_1$  or any prior read can not be ordered after any remaining operation.

*Case 1.  $w_1$  was submitted by process  $j$ .* Every read between A and B is a prior read. The remaining operations are all writes and cannot be ordered by  $<$  before  $w_1$  by the invariant. The remaining operations also cannot be ordered before any prior read by  $<$ . They cannot be ordered by  $<_{PO}$  because the write would be by process  $j$  and so would be ordered before  $w_1$  which is a contradiction. A read and a write cannot be ordered by  $<_{WO}$  or  $<_{AO(<_{SO})}$  because those two orders only occur between pairs of write operations. Also, a read cannot be ordered after a write by  $<_{SO}$ .

*Case 2.  $w_1$  was not submitted by process  $j$ .* If a remaining operation is a read, it is by process  $j$  so it cannot be ordered before  $w_1$  by  $<_{PO}$ . The remaining read also cannot be ordered before  $w_1$  by  $<_{WO}$  or  $<_{AO(<_{SO})}$  because those orders only occur

between pairs of write operations. The remaining read cannot be ordered before  $w_1$  by  $<_{SO}$  because the read would be to the same variable as  $w_1$ , and so would be a prior read. The remaining read cannot be ordered before any prior read because all reads are by process  $j$  so it would be ordered before a prior read by  $<_{PO}$  making it a prior read.

If the remaining operation is a write, it cannot be ordered by  $<$  before  $w_1$  by the invariant. It cannot be ordered before a prior read by  $<_{WO}$  or  $<_{AO(<_{SO})}$  because those only order pairs of writes. It cannot be ordered before a prior read by  $<_{SO}$  because a read cannot be ordered after a write by  $<_{SO}$ . All that remains is to show that a remaining operation which is a write cannot be ordered before a prior read by  $<_{PO}$ . Any prior read,  $r$ , comes before  $w_1$  in  $<_j$ . The write,  $w_2$ , which wrote to  $r$  must also come before  $w_1$  because  $<_j$  is serial. If  $r$  is to the same variable as  $w_1$ , then either,  $w_1 <_{SO} w_2$ , or  $r <_{SO} w_1$ . Since  $<_j$  respects  $<_{SO}$ , it must be the case that  $r <_{SO} w_1$ . If a remaining operation  $w_3$  is ordered before a prior read,  $r_1$ , by  $<_{PO}$ , then either  $r_1$  is to the same variable as  $w_1$  in which case  $r_1 <_{SO} w_1$ , or  $r_1$  is ordered by  $<_{PO}$  before  $r_2$  which is to the same variable as  $w_1$  in which case  $r_2 <_{SO} w_1$ . Therefore,  $w_3 <_{PO} (r_1 <_{PO})r_2 <_{SO} w_1$  so  $w_3 <_{AO} w_1$  which is a contradiction of the invariant.

In either case,  $w_1$  and all prior reads are not ordered after any remaining operations by  $<$ . Now  $<_j$  is changed as follows: All prior reads are placed immediately before point A preserving their order followed by  $w_1$ . All other operations preserve their order. For all pairs of operations that change their relative position, one must be  $w_1$  or a prior read. The other must be a remaining operation. These pairs cannot be ordered by  $<$  so the view still respects  $<$ .

Before the move,  $<_j$  was serial so each prior read must have read from the most recent write to that variable. That write must have been before point A because it is anti ordered before  $w_1$ . The write must still be the most recent write to the same variable because the moved read is after all writes before point A, and every write between the two was there before the move when  $<_j$  was serial. Remaining reads maintained their relative position with all writes except  $w_1$ . Remaining reads cannot be to variable  $x$ , and so they too must still read from the most recent write. No other pairs of reads and writes changed relative position so  $<_j$  must still be serial.

Now, move point A to immediately after  $w_1$ . The next write in  $<_i$  becomes the new  $w_1$ . This write has not been moved to before point A in  $<_j$  so point B is still after point A. The set of writes before  $w_1$  in  $<_i$  have all been moved to before point A in  $<_j$ , so the invariants are satisfied. Therefore, by induction, one can create views for all processes that respect  $<$  and have the write operations in the same order in all views.  $\square$

**LEMMA 4.16.** *For any  $GPO + GWO + GAO$  execution, it is possible to construct a single view containing all operations that respects process order and is serial.*

**PROOF.** From Lemma 4.15, create views which all have the write operations in the same order. These orders respect  $<_{PO}$  and are serial. Then take one of these views and add the read operations of all other processes in the same relative position to the writes as they occur in their own view. The read operations must all be ordered by  $<_{PO}$  correctly with respect to all writes because the writes occur in the same order in every view. Reads ordered with respect to each other by  $<_{PO}$  come from the

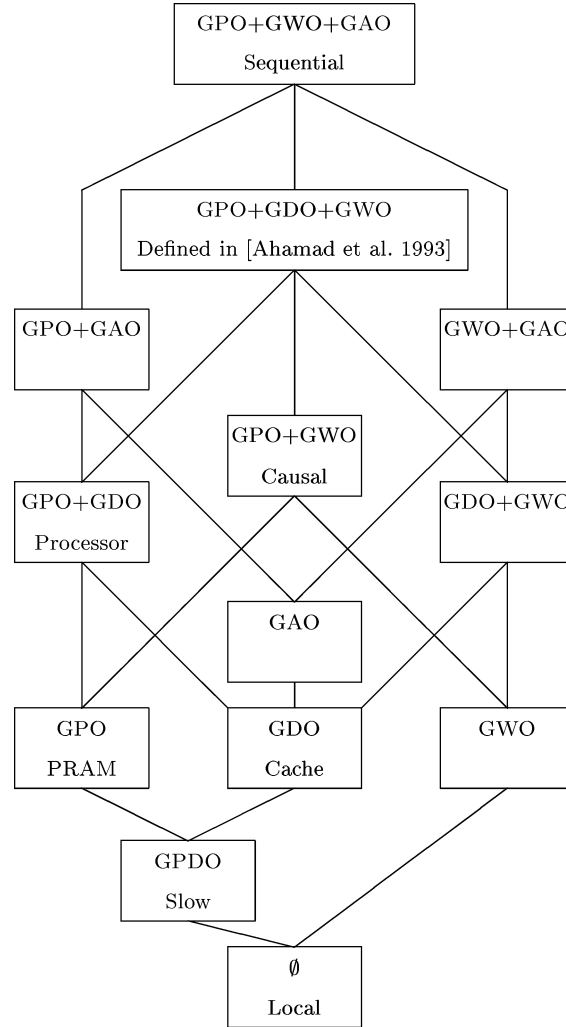


FIG. 13. The complete lattice of consistency models.

same view, and so they are placed in that order in the new view. The serial property is not affected by the relative position of pairs of reads, and every read operation is in the same position relative to all writes, so the view must be serial.  $\square$

**THEOREM 4.17.**  *$GPO + GWO + GAO$  is equivalent to sequential consistency.*

**PROOF.** Follows directly from Lemmas 4.14 and 4.16.  $\square$

Adding GAO almost completes the lattice as shown in Figure 13. Since GAO is stronger than GDO any box labeled with GAO will also enforce GDO, but that is not shown for brevity. The lattice now has three additional new consistency models: GAO,  $GPO + GAO$ , and  $GWO + GAO$ . The lattice is almost complete, but it does not yet contain slow consistency. Slow consistency would be located below both PRAM and cache, and above local.

4.4. SLOW CONSISTENCY AS A COMBINATION OF PROPERTIES. In slow consistency [Hutto and Ahamad 1990], two operations must maintain their order only if they are by the same process and to the same variable. This leads to the following definitions.

*Definition 4.18.* Two operations are ordered by *process-data order*,  $o_1 <_{PDO} o_2$ , iff  $o_1 <_{PO} o_2$ , and  $o_1 <_{DO} o_2$ .

*Definition 4.19.* An execution is *Global Process-Data Order (GPDO)* iff

$$\forall i \in P \exists \text{SerialView}(<_i \cup <_{PDO} | (*, i, *, *) \cup (w, *, *, *)).$$

**THEOREM 4.20.** *GPDO is equivalent to slow consistency.*

**PROOF.** For any GPDO execution, take the view for a single processor. Divide this view into separate views, one for each variable by restricting the set of operations to operations on a single variable, but maintaining their relative order. Process-data order contains all edges in process order between operations to the same variable. These views respect process-data order, and contain only operations to a single variable so they respect process order. These views are exactly what is required to satisfy slow consistency.

For any slow consistent execution, gather together the views over all variables for a particular processor. By similar logic to Lemma B.5, the union of these views and  $<_i$  must be acyclic. The union of the views must contain every edge in process-data order. Therefore, any topological sort of the union of the views and  $<_i$  must respect  $<_i \cup <_{PDO}$ . Also, each view is serial. In the topological sort, every pair of operations to the same variable must preserve their relative position so the topological sort must be serial. The topological sort is exactly what is required to satisfy GPDO.  $\square$

GPDO is more than just a new statement of slow consistency. It represents a new way of combining consistency properties. We have already seen  $GPO + GDO$  as a way to combine two models to produce a stronger model. Now, GPDO combines two models to produce a weaker model. This could be done for any pair of properties. For example, process anti-order orders only operations that are ordered by both process order and anti order. GPAO would be weaker than both GPO and GAO. However, it is questionable how useful models this weak would be. Slow consistency is essentially only valuable in defining synchronized models. Perhaps these models would be usable with a transition theory, and higher consistency operations between them for synchronization.

4.5. A LATTICE OF CONSISTENCY MODELS. The result of these composable consistency properties is the lattice of consistency models shown in Figure 13. Every possible combination of properties produces a model represented by a box in the lattice. The top of the lattice is sequential consistency, and the bottom is local consistency. Every pair of models has a unique least upper bound and greatest lower bound. There are other combinations of properties demonstrated in this work such as  $GPO \cap GDO$ , and GPAO. These are not shown in the lattice for brevity, and because their utility is unknown. GPDO is shown in the lattice because slow consistency is a well known and widely used model.

One can think of every box in the lattice as representing a set of executions that satisfies that model, and no stronger model in the lattice. To show that every box of the lattice is nonempty we provide example executions that violate each

$p_5$	$(w, p_5, a, 1)$	$(r, p_5, a, 2)$
$p_6$	$(w, p_6, a, 2)$	$(r, p_6, a, 1)$

FIG. 14. An execution that violates GDO.

$p_7$	$(r, p_7, b, 2)$	$(w, p_7, c, 1)$
$p_8$	$(r, p_8, c, 1)$	$(w, p_8, b, 2)$

FIG. 15. An execution that violates GWO.

of the four consistency properties. To derive an example execution for a particular box, combine the executions violating all the properties not contained in that box. Figure 12 given when defining anti-order in Section 4.3 provides an execution that violates GAO without violating any of the other three properties.

Figure 14 provides an execution that violates GDO (and thus GAO), but does not violate GPO or GWO. From condition 3 of data order:

$$(w, p_5, a, 1) <_{DO} (w, p_6, a, 2) <_{DO} (w, p_5, a, 1).$$

Therefore, there is a cycle in  $<_{DO}$  so the execution is not GDO. However, write-read-write order is empty. The following views satisfy  $<_{PO}$  and  $<_{WO}$ , and are serial:

$$\begin{aligned} p_5 &: (w, p_5, a, 1) <_5 (w, p_6, a, 2) <_5 (r, p_5, a, 2) \\ p_6 &: (w, p_6, a, 2) <_6 (w, p_5, a, 1) <_6 (r, p_6, a, 1). \end{aligned}$$

Figure 15 provides an execution that violates GWO, but does not violate GPO, GDO, or GAO. The following cycle exists.

$$(w, p_7, c, 1) <_{WO} (w, p_8, b, 2) <_{WO} (w, p_7, c, 1).$$

These two writes must be present in all views, so no view can respect  $<_{WO}$ . Each write is data ordered before the read it writes-to. Serial order and anti-order are empty. The following views satisfy  $<_{PO}$ ,  $<_{DO}$ ,  $<_{AO(<_{SO})}$ ,  $<_{SO}$ , and are serial:

$$\begin{aligned} p_7 &: (w, p_8, b, 2) <_7 (r, p_7, b, 2) <_7 (w, p_7, c, 1) \\ p_8 &: (w, p_7, c, 1) <_8 (r, p_8, c, 1) <_8 (w, p_8, b, 2). \end{aligned}$$

To produce an execution that satisfies only GPO and no stronger model in the lattice, define an execution containing  $p_5$  and  $p_6$  from Figure 14 and  $p_7$  and  $p_8$  from Figure 15. Likewise, to create an execution satisfying only GPO + GDO combine Figure 12 with Figure 15, and so forth.

Figure 16 provides an execution that violates GPO, but does not violate GDO, GWO, or GAO. In order for the view for  $p_{10}$  to be serial,  $(w, p_9, e, 2)$  must come before  $(r, p_{10}, e, 2)$ , and  $(w, p_9, d, 1)$  must come after  $(w, p_{10}, d, 3)$ . In order to respect local order,  $(r, p_{10}, e, 2)$  must come before  $(w, p_{10}, d, 3)$ . Therefore,  $(w, p_9, e, 2)$  must come before  $(w, p_9, d, 1)$  which does not respect  $<_{PO}$ .

The following are the definitions of  $<_{DO}$  and  $<_{WO}$  for this execution:

$$\begin{aligned} (w, p_{10}, d, 3) &<_{DO} (w, p_9, d, 1) <_{DO} (r, p_{10}, d, 1) \\ (w, p_9, e, 2) &<_{DO} (r, p_{10}, e, 2) \\ (w, p_9, e, 2) &<_{WO} (w, p_{10}, d, 3). \end{aligned}$$

$p_9$	$(w, p_9, d, 1)$	$(w, p_9, e, 2)$	
$p_{10}$	$(r, p_{10}, e, 2)$	$(w, p_{10}, d, 3)$	$(r, p_{10}, d, 1)$

FIG. 16. An execution that violates GPO.

$p_{11}$	$(w, p_{11}, f, 1)$	$(w, p_{11}, g, 4)$	$(r, p_{11}, g, 3)$
$p_{12}$	$(w, p_{12}, g, 3)$	$(w, p_{12}, f, 2)$	$(r, p_{12}, f, 1)$

FIG. 17. An execution that satisfies GWO + GAO.

With the following definition of serial order, anti order is empty.

$$(w, p_{10}, d, 3) <_{SO} (w, p_9, d, 1)$$

The following view for  $p_{10}$  satisfies  $<_{DO}$ ,  $<_{WO}$ ,  $<_{AO(<_{SO})}$ ,  $<_{SO}$ ,  $<_{p_{10}Local}$ , and is serial:

$$p_{10} : (w, p_9, e, 2) <_{10} (r, p_{10}, e, 2) <_{10} (w, p_{10}, d, 3) <_{10} (w, p_9, d, 1) <_{10} (r, p_{10}, d, 1).$$

However, the view for  $p_9$  is not as simple. The following cycle exists:

$$(w, p_{10}, d, 3) <_{SO} (w, p_9, d, 1) <_{p_9Local} (w, p_9, e, 2) <_{WO} (w, p_{10}, d, 3).$$

No view can be written for  $p_9$  that satisfies GWO + GAO. However, separate views can be written, one that satisfies GWO, and one that satisfies GAO.

$$\begin{aligned} p_9(GWO) : (w, p_9, d, 1) <_9 (w, p_9, e, 2) <_9 (w, p_{10}, d, 3) \\ p_9(GAO) : (w, p_{10}, d, 3) <_9 (w, p_9, d, 1) <_9 (w, p_9, e, 2). \end{aligned}$$

Therefore, this execution satisfies GAO, and no stronger model in the lattice. It also satisfies GWO, and no stronger model in the lattice. By combining this execution with Figure 12, we achieve an execution that satisfies only GDO. All that remains is to find executions that satisfy GWO + GAO and GDO + GWO.

Figure 17 satisfies GWO + GAO, but not GPO + GWO + GAO. Below is the definition of  $<_{DO}$ :

$$\begin{aligned} (w, p_{12}, f, 2) <_{DO} (w, p_{11}, f, 1) <_{DO} (r, p_{12}, f, 1) \\ (w, p_{11}, g, 4) <_{DO} (w, p_{12}, g, 3) <_{DO} (r, p_{11}, g, 3). \end{aligned}$$

The following definition of serial order must be chosen.

$$(w, p_{11}, g, 4) <_{SO} (w, p_{12}, g, 3)$$

If not; then  $(w, p_{11}, g, 4)$  must be ordered after  $(r, p_{11}, g, 3)$ , which violates the order  $<_{p_{11}Local}$ . Likewise, for  $(w, p_{12}, f, 2)$  and  $(r, p_{12}, f, 1)$ .  $<_{WO}$  and  $<_{AO(<_{SO})}$  are empty. The following cycle exists.

$$(w, p_{11}, g, 4) <_{SO} (w, p_{12}, g, 3) <_{PO} (w, p_{12}, f, 2) <_{SO} (w, p_{11}, f, 1) <_{PO} (w, p_{11}, g, 4).$$

Therefore, it is impossible for any view to respect both  $<_{PO}$  and  $<_{SO}$ . So the execution is not GPO + GAO, and hence it is not GPO + GWO + GAO. However,



this execution is GWO + GAO as the following views demonstrate.

$$\begin{aligned}
 p_{11} : (w, p_{12}, f, 2) &<_{11} (w, p_{11}, f, 1) <_{11} (w, p_{11}, g, 4) <_{11} \\
 (w, p_{12}, g, 3) &<_{11} (r, p_{11}, g, 3) \\
 p_{12} : (w, p_{11}, g, 4) &<_{12} (w, p_{12}, g, 3) <_{12} (w, p_{12}, f, 2) <_{12} \\
 (w, p_{11}, f, 1) &<_{12} (r, p_{12}, f, 1)
 \end{aligned}$$

To create an execution that satisfies GDO + GWO and no stronger model combine this execution with Figure 12. The complete lattice as shown in Figure 13 is a powerful new way to describe and organize consistency models. Every non-synchronized model described in Section 2 is encompassed by the lattice model. In addition, five new consistency models are uncovered by the symmetry of the lattice. Every model in the lattice has a nonempty set of executions which satisfy that model and no stronger model in the lattice. Finally, new consistency properties would be easy to integrate into the lattice if they are discovered. Synchronized models are not covered directly by the lattice. Instead, synchronized models can be viewed as processes submitting some operations under one consistency model, and some operations under another consistency model, that is, a consistency transition. Synchronized models will be covered in Section 5 on consistency transitions. The lattice model facilitates the definition of consistency transitions because any two models are easily compared by the properties they enforce.

### 5. Consistency Transitions

Our final generalization of consistency models is the idea of consistency transitions. In synchronized consistency models, a program executes ordinary operations with a relaxed consistency model, usually slow consistency. Occasionally, the program executes synchronization operations with a stronger consistency model, usually sequential consistency. These synchronization operations enforce additional ordering restrictions between ordinary operations. This can be viewed as a consistency transition where the process executing a synchronization operation temporarily requests a stronger level of consistency. Our goal is to develop a general theory of consistency transitions between any two consistency models, not just slow and sequential. Synchronized models require the following:

- (1) All synchronization operations must be sequentially consistent.
- (2) All ordinary operations must be slow consistent.
- (3) The order  $<_D$  must be respected between synchronization and ordinary operations

Sequential consistency is equivalent to GPO + GWO + GAO. So the first condition can be satisfied with serial views on synchronization operations.

$$\begin{aligned}
 \forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{PO} \cup <_{WO} \cup <_{SO} \cup <_{AO(<_{SO})} \mid \\
 (sr, i, *, *) \cup (sw, *, *, *)).
 \end{aligned}$$

Weak consistency does not include acquire and release operations. Instead, synchronization operations are special read and write operations. To distinguish them, we use the operation types *sr* for synchronized read and *sw* for synchronized write. Remember that, for other synchronized models, the writes-to relation is defined

$p_1$	$(r, p_1, y, 2)$	$(sw, p_1, z, 3)$	$(w, p_1, x, 1)$
$p_2$	$(r, p_2, x, 1)$	$(sw, p_2, z, 4)$	$(w, p_2, y, 2)$

FIG. 18. An execution that violates weak consistency.

with acquire operations treated as reads, and release operations treated as writes. If an acquire is defined as an  $sr$  and a release as an  $sw$ , this definition is equally valid for every synchronized model.

Slow consistency is equivalent to GPDO. So the second condition can be satisfied by serial views on ordinary operations.

$$\forall i \in P \exists \text{ SerialView}(<_i \cup <_{PDO} \mid (or, i, *, *) \cup (ow, *, *, *)).$$

The operation type  $or$  is used for ordinary read, and  $ow$  for ordinary write. The views for synchronization and ordinary operations are very similar. They each have one view per processor, and each view contains the reads of that processor plus all writes. It would be nice to combine these views into a single view for each processor containing both synchronization and ordinary operations. The view would have to respect the ordering among synchronization operations,  $<_{synch}$ ,

$$<_{synch} \equiv <_{PO} \cup <_{WO} \cup <_{SO} \cup <_{AO(<_{SO})} \mid (sr, i, *, *) \cup (sw, *, *, *)$$

and the ordering among ordinary operations,  $<_{ord}$ ,

$$<_{ord} \equiv <_{PDO} \mid (r, i, *, *) \cup (w, *, *, *)$$

and it would have to respect  $<_i$  and  $<_D$ . However, this straightforward approach has some problems.

Figure 18 satisfies all of these properties and still does not satisfy weak consistency as the following views demonstrate.

$$\begin{aligned} p_1 : (sw, p_2, z, 4) <_1 (w, p_2, y, 2) <_1 (r, p_1, y, 2) <_1 (sw, p_1, z, 3) <_1 \\ & (w, p_1, x, 1) \\ p_2 : (sw, p_1, z, 3) <_2 (w, p_1, x, 1) <_2 (r, p_2, x, 1) <_2 (sw, p_2, z, 4) <_2 \\ & (w, p_2, y, 2) \end{aligned}$$

Notice that the synchronized writes are unordered by  $<_{synch}$ . They may occur in either order, but in this execution they are seen to occur in different orders by different processes. Does this violate the assertion that synchronization operations must be sequentially consistent? After all, the synchronization operations by themselves, ignoring ordinary operations, are sequentially consistent. The reason for this conundrum comes from a slight discrepancy between the intuitive definition and the formal definition of sequential consistency. The intuitive definition can be stated like this.

There is a single total order of events, and all processes agree that the events happened in that order.

However, the formal definition requires that there exist at least one order of events that every process can agree on. There may be more than one order of events that would satisfy every process, and there is no way to distinguish a single correct order from the sequentially consistent operations alone. This problem is not an artifact of our definition of  $GPO + GWO + GAO$ . It can still occur with the original definition

of sequential consistency. Below is a restatement of the definition given previously for synchronized consistency models except that the positions of  $\forall_{i \in P, x \in V}$  and  $\exists <_{seq} = \dots$  have been reversed.

$$\begin{aligned} &\forall_{i \in P, x \in V} \exists <_{seq} = \text{SerialView}(<_{PO} \mid (s, *, *, *)), \text{ and} \\ &<_S = \text{the transitive closure of } <_D \cup <_{seq}, \text{ and} \\ &\exists \text{ SerialView}(<_S \cup <_{PO} \mid (*, i, x, *) \cup (w, *, x, *)). \end{aligned}$$

The synchronized operations are sequentially consistent, but each process gets to choose its own definition of  $<_{seq}$ . This causes the same problem. The original definition resolved this problem by requiring that the definition of  $<_S$  for every process be based on a single definition of  $<_{seq}$ . This same strategy can be used with GPO + GWO + GAO to generate the definition given below. Note that all synchronized reads must be included in every view. This will be addressed later.

**THEOREM 5.1.** *The following definition is equivalent to synchronized model consistency*

$$\begin{aligned} &\forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{synch} \cup <_{ord} \cup <_D \mid (or, i, *, *) \cup \\ &(ow, *, *, *) \cup (sr, *, *, *) \cup (sw, *, *, *)) \text{ and all synchronization} \\ &\text{operations appear in the same order in every view.} \end{aligned}$$

**PROOF.** For an execution that satisfies the above views, construct the original definition synchronized consistent views as follows. The view  $<_{seq}$  is the total order on synchronization operations that occurs in every view. Any two ordinary operations ordered by  $<_S$  must be ordered by a transitive chain containing only synchronization operations. The per-process views contain all synchronization operations and respect  $<_D$  and  $<_{seq}$  so they must also respect  $<_S$ . Construct the per-process per-variable slow consistent views required by weak consistency from the per-process GPDO consistent views as shown in Theorem 4.20. The new views will respect the old views so they will respect  $<_S$ .

For an execution that satisfies the original definition of synchronized consistency, construct the above views as follows. Begin with all the synchronization operations in the order specified by  $<_{seq}$ . This is the order in which they will appear in every per-process view. The synchronization operations must respect  $<_{synch}$  because by Lemma 4.14 every sequentially consistent view respects process order, write-read-write order, serial order, and anti order. Any single per-process, per-variable slow consistent view can always be combined with the synchronization operations in a way that respects  $<_D$  because the view respects  $<_S$  which is the transitive closure of  $<_{seq}$  and  $<_D$ . Combine all slow consistent views with the synchronization operations in this way ignoring, for now, the order between operations from different slow consistent views. The resulting view will respect  $<_{synch}$ ,  $<_{ord}$ , and  $<_D$ . All that remains is to show that it respects  $<_i$ .

Between two synchronization operations, the ordinary operations can always be rearranged as a topological sort of  $<_{ord} \cup <_i$  which is acyclic by Lemma B.5. Two ordinary operations separated by synchronization operations cannot be out of order with respect to  $<_i$  because

$$o_1 <_D s_1 <_{seq} s_2 <_D o_2 <_i o_1.$$

This implies that  $s_2$  is process ordered before  $s_1$ , but appears after it in  $<_{seq}$ , which is a contradiction.  $\square$




FIG. 19. Linearizability for synchronization operations.

Should all processes be required to see the same total order of synchronization operations, or is it sufficient that the synchronization operations are sequentially consistent? We argue that sequential consistency of synchronization operations should be sufficient even if this allows different processes see different total orders. First, we feel that the intuitive definition is in fact enforcing a consistency model stronger than sequential. For example, linearizability [Herlihy and Wing 1990] assumes the existence of a global Newtonian clock. The processes may not have access to this clock, but it does exist. Each operation spans a certain period of time. A linearizable execution must be sequential, and in addition if two operations have non-overlapping time spans they must appear in the sequential view in that order. Perhaps this problem would be solved if synchronization operations were linearizable, and ordinary operations had defined time spans and were forced to respect certain linearizable restrictions with synchronization operations.

For example, Figure 19 shows how linearizability could solve this problem for Figure 18. Even if the time spans for  $(sw, p_1, z, 3)$  and  $(sw, p_2, z, 4)$  overlap, that is, they can be seen in either order, there is no way that  $(r, p_1, y, 2)$  and  $(w, p_2, y, 2)$  can overlap while  $(r, p_2, x, 1)$  and  $(w, p_1, x, 1)$  also overlap. The definitions given for synchronized consistency models explicitly state that synchronization operations must be sequentially consistent. However, the implementations given with those definitions implicitly enforce linearizability over synchronization operations. The authors of the various models did not appreciate the effect of this slight distinction.

Another reason not to require every process to see the same total order is once again the argument over the distinction between memory model and programming model. The reader may have noticed that Figure 18 does not implement any kind of mutual exclusion or barrier behavior. The program does not know in which order the synchronized writes occurred, but is relying on the fact that they occurred in the same order at all processes. If the program knows that two synchronization operations occurred in a particular order the problem disappears. If the operations are ordered by  $<_{synch}$  then they must appear in that order in all views. In our opinion, if the programmer needs two operations to occur in the same order in all views then the control and data flow of the program must be able to detect in what order they occurred. This is part of the programming model, not the consistency model. In particular, this problem does not occur for data-race-free programs because every pair of conflicting ordinary operations is separated by synchronization operations with control or data dependencies. That is, the synchronization operations must be ordered by  $<_{synch}$ . We propose to redefine  $<_S$  for synchronized consistency models. Rather than being the transitive closure of  $<_D \cup <_{seq}$  it should be the transitive closure of  $<_D \cup <_{synch}$ . Essentially, the synchronization operations must be sequentially consistent, and if the program can tell that two synchronization operations happened in a particular order then they must be placed in that order in all process' views. This leads to a revised definition of synchronized model consistency.

**Definition 5.2.** For a given definition of  $<_D$ , an execution is *synchronized model consistent* with the new definition  $<_S$  iff

$$\begin{aligned} \exists <_{seq} = \text{SerialView}(<_{PO} \mid (s, *, *, *)), \text{ and} \\ <_S = \text{the transitive closure of } <_D \cup <_{synch}, \text{ and} \\ \forall_{i \in P, x \in V} \exists \text{SerialView}(<_S \cup <_{PO} \mid (*, i, x, *) \cup (w, *, x, *)). \end{aligned}$$

**THEOREM 5.3.** *The following definition is equivalent to synchronized model consistency with the new definition  $<_S$*

$$\begin{aligned} \forall_{i \in P} \exists (<_i \cup <_{synch} \cup <_{ord} \cup <_D \mid (or, i, *, *) \cup \\ (ow, *, *, *) \cup (sr, *, *, *) \cup (sw, *, *, *)). \end{aligned}$$

**PROOF.** For an execution that satisfies the above views, construct the synchronized consistent views as follows. The order  $<_{seq}$  is taken from any of the views as they all contain all synchronization operations. Any two ordinary operations ordered by  $<_S$  must be ordered by a transitive chain containing only synchronization operations. The per-process views contain all synchronization operations and respect  $<_D$  and  $<_{synch}$  so they must also respect  $<_S$ . Construct the per-process per-variable slow consistent views required by weak consistency from the per-process GPDO consistent views as shown in Theorem 4.20. The new views will respect the old views so they will respect  $<_S$ .

For an execution that satisfies the new definition of synchronized consistency, construct the above views as follows. There must be at least one order of synchronization operations that respects  $<_{synch}$  because  $<_{seq}$  exists. Furthermore, each per-process, per-variable slow consistent view respects  $<_S$  so it can always be combined with an ordering of synchronization operations that respects  $<_{synch}$  and  $<_D$ . By similar logic as above, combine all operations for a single process into a single view, and the view will respect  $<_i$ .  $\square$

Now we will deal with the fact that every synchronized read must be placed in every view. The proof above relies on the fact that if  $o_1 <_S o_2$  then  $o_1$  and  $o_2$  must be placed in that order in every view in which they both occur. This is enforced by the fact that every view contains all synchronization operations and respects  $<_D$  and  $<_{synch}$ . If some views were not to contain some synchronized reads this might not hold. There are two cases in which ordinary operations can be ordered by  $<_S$ . Case 1,  $o_1$  and  $o_2$  are linked by a transitive chain containing at least one  $sw$ . In this case, the  $sw$  will be in every view so we can just link the ordinary operations to the synchronized write instead of any possible synchronized reads in the chain. Case 2,  $o_1$  and  $o_2$  are linked by a transitive chain containing only synchronized reads. In this case, we can link the ordinary operations to each other. This will be called transitive order,  $<_T$ . In Definition 5.4,  $<_{synch}^+$  refers to traversing one or more edges of  $<_{synch}$ .

**Definition 5.4.** Transitive order,  $<_T$ , is defined as

$$\begin{aligned} \text{if } o <_D sr <_{synch}^+ sw, \text{ then } o <_T sw \\ \text{if } sw <_{synch}^+ sr <_D o, \text{ then } sw <_T o \\ \text{if } o_1 <_D sr <_D o_2, \text{ then } o_1 <_T o_2 \\ \text{if } o_1 <_D sr_1 <_{synch}^+ sr_2 <_D o_2, \text{ then } o_1 <_T o_2. \end{aligned}$$

Now we have another equivalent definition of synchronized model consistency where each per-process view contains only its own reads whether ordinary or synchronized.

**THEOREM 5.5.** *The following definition is equivalent to synchronized model consistency with the new definition  $<_S$*

$$\forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{\text{synch}} \cup <_{\text{ord}} \cup <_D \cup <_T \mid (or, i, *, *) \cup (ow, *, *, *) \cup (sr, i, *, *) \cup (sw, *, *, *)).$$

**PROOF.** By Lemma 4.16, it must still be possible to construct the view  $<_{\text{seq}}$ . Also, the views must still respect  $<_S$  because any transitive chain in  $<_D$  and  $<_{\text{synch}}$  must be reflected in the operations present in each view through  $<_D$ ,  $<_{\text{synch}}$ , and  $<_T$ .  $\square$

Now this definition can be generalized. The definition says that sequential consistency operations must be sequentially consistent with each other, slow consistent operations must be slow consistent with each other, and operations of different consistency levels must respect  $<_D$  and  $<_T$  between them. There is no reason this definition has to be limited to sequential and slow consistency, or limited to just two consistency levels. Each operation can be submitted under a different consistency model; any model within the lattice. This leads to a generalized definition of memory consistency. Each operation is considered to be labeled with a subset of the consistency properties, and two operations must respect an order such as process order if they are both labeled with the global process order property.

**Definition 5.6.** Two operations are ordered by *synchronization order*  $o_1 <_{\text{synch}} o_2$  iff

both are labeled GPO and  $o_1 <_{PO} o_2$ , or  
 both are labeled GDO and  $o_1 <_{DO} o_2$ , or  
 both are labeled GWO and  $o_1 <_{WO} o_2$ , or  
 both are labeled GAO and  $o_1 <_{SO} o_2$  or  $o_1 <_{AO(<_{SO})} o_2$ , or  
 both are labeled GPDO and  $o_1 <_{PDO} o_2 \dots$

**Definition 5.7.** For a given definition of  $<_D$ , an execution satisfies *generalized memory consistency* iff

$$\forall_{i \in P} \exists \text{ SerialView}(<_i \cup <_{\text{synch}} \cup <_D \cup <_T \mid (r, i, *, *) \cup (w, *, *, *)).$$

So a consistency model is defined by specifying  $<_D$  and labeling operations with consistency properties. To simulate the nonsynchronized models,  $<_D$  is empty and all operations are labeled with the consistency properties of that model.  $<_{\text{synch}}$  reduces to the union of the orders representing the labeled properties. For example, if all operations are labeled GPO + GWO, this definition reduces to the original definition of causal consistency. To simulate the synchronized consistency models, use  $<_D$  given for that model. Synchronization operations are labeled with GPO + GWO + GAO, and ordinary operations are labeled with GPDO. This definition can also accommodate the variant of release consistency where synchronization operations respect processor consistency. To simulate location consistency all that is needed is to replace SerialView with SerialPartialView as given in Definition 3.11.

These new definitions also allow new ideas about what it means for a program to be data-race-free [Adve and Hill 1993]. A data-race-free program is one that

Process $p_1$	Process $p_2$
The initial value of $x$ is 0;	
$y = f(\text{input});$	
$x = 1; \text{synch}$	$\text{while}(x==0) \text{ wait}; \text{synch}$
	$\text{read}(y);$

FIG. 20. A data-race-free program.

will only produce sequential executions even when the memory system supports a particular consistency model weaker than sequential consistency. For example, Figure 20 contains a program that will only produce sequential executions when it is run under weak consistency. This program is said to be weak-sequential data-race-free. A program may be data-race-free for some nonsequential consistency models, and not data-race-free for others [Gharachorloo et al. 1990]. The operations on  $x$  are synchronization operations. In order to exit the loop,  $p_2$  must read 1 from  $x$ . Therefore, the following ordering restrictions exist.

$$(w, p_1, y, f(\text{input})) <_D (w, p_1, x, 1) \mapsto (r, p_2, x, 1) <_D (r, p_2, y, ?)$$

The view for  $p_2$  must contain all of these operations. If weak consistency is enforced, then  $<_D$  must be respected, and  $\mapsto$  must be respected because the view is serial. There are no other writes to  $y$ , so  $(r, p_2, y, ?)$  must return the value written by  $(w, p_1, y, f(\text{input}))$ . If this value is returned then the execution is also sequentially consistent. One goal of synchronized consistency models is to simulate sequential consistency in this manner. This work provides a new, formal definition of what it means to be a data-race-free program. A program is *data-race-free* if and only if, for any execution produced by the program,

Given the definition of  $<_D$  and labeling of operations required for weak consistency:

$$\exists <_{SO} \forall i \exists \text{ SerialView}(<_i \cup <_{\text{synch}} \cup <_D \cup <_T \mid (*, i, *, *) \cup (w, *, *, *))$$

implies

$$\exists <_{SO} \forall i \exists \text{ SerialView}(<_{PO} \cup <_{WO} \cup <_{AO(<_{SO})} \cup <_{SO} \mid (*, i, *, *) \cup (w, *, *, *))$$

This literally says that if the program produces a weak consistent execution, then that same execution is also sequentially consistent. If the program is run in an environment that only produces weak consistent executions, then the program will only produce sequentially consistent executions. This definition of data-race-free is very general, but may not be too helpful to programmers. It does not give insight on how to write a program that satisfies the condition, and it may be hard to prove that a particular program satisfies the condition. For example, it does not even require that the same definition of serial order be used to produce the weak consistent views as the sequentially consistent views. One could provide simpler, conservative definitions that are easier to implement and prove, but still enforce the above condition. For example, if every pair of operations ordered by  $<_{PO} \cup <_{WO} \cup <_{AO(<_{SO})} \cup <_{SO}$  were also ordered by  $<_i \cup <_{\text{synch}} \cup <_D \cup <_T$  then the condition would hold. A further restriction along these lines is to say that every pair of ordinary operations to the same variable must be separated by control

and data dependencies among synchronization operations which is the traditional definition of data-race-free. This new uniform notation may allow more precise, less conservative formulations of the class of data-race-free programs.

## 6. Conclusions and Future Work

The thesis of this work is that every useful shared memory consistency model (well known and often used models in the literature) can be described by a single unifying framework. This work presents such a framework in the form of a lattice of primitive consistency properties, and a theory of transitions within the lattice. Shared memory can be viewed as an abstract API of interprocess communication parameterized by its consistency model. This API can be implemented in environments with physically shared memory banks in hardware. Or in environments with no physically shared memory, as in distributed shared memory systems. This style of interprocess communication is appropriate for many types of applications which can leverage research done on memory implementations and memory consistency models.

The first contribution of this work is the discovery of four fundamental consistency properties. Global Process Order enforces the condition that all operations by a single process are seen everywhere in the system to occur in the order in which they were submitted. Global Data Order enforces the condition that for each variable, there exists at least one total order of operations which every process can agree could have been the actual order of those operations. Combining these two orders produces another consistency model,  $GPO + GDO$ , very similar to processor consistency. The difference arises in the fact that there may be more than one possible total order on each variable which satisfies data order. However, data order can be augmented to be a total order on operations to each variable. Processor consistency is equivalent to process order plus this augmented data order. This method of combining consistency properties is a general method which can be used to create a lattice of consistency models. Any two properties can be combined in this way to produce a consistency model stronger than either property alone. This work has also identified another combination operator which produces a new model weaker than either property alone. In this case,  $GPDO$  produces slow consistency. Thus, all possible combinations of consistency properties produce a lattice of models.

The third property, Global Write-read-write Order enforces aspects of causality. It is defined such that  $GPO + GWO$  is equivalent to causal consistency. It is the weakest property (the smallest set of edges) for which this is true. The fourth property is Global Anti Order. Anti order is defined such that all four properties combined produce a model equivalent to sequential consistency. To accomplish this, Global Anti-Order requires two ordering relations among operations, anti order and serial order. Serial order captures the restriction that every read must read from the most recent write. Anti order is based on both serial order and data order. This complexity is required as any weaker definition of anti order was not sufficient to enforce equivalence to sequential consistency. Another side effect of this complexity is that Global Anti-Order is not orthogonal to all three other properties. It is strictly stronger than data order.

The second contribution of this work is the concept of a consistency lattice. As stated before, enumerating every combination of the four consistency properties with both combination operators produces a lattice of consistency models.



The strongest model in the lattice is sequential consistency, and the weakest is local consistency. Every nonsynchronized consistency model described in Section 2.2 is equivalent to a node in this lattice. The lattice model validates the derived consistency properties as necessary and sufficient to describe all such models. Furthermore, for every consistency model in the lattice there exists a non-empty set of executions accepted by that model and no stronger model in the lattice.

The third contribution of this work is that the lattice includes five previously unnamed, nonempty consistency models: GWO, GAO, GDO + GWO, GPO + GAO, GWO + GAO. We believe the most promising of these is GDO + GWO. It is a data-centric version of causality where operations are placed in causal order when they are applied to their variable, not when they are issued by their process.

The fourth contribution of this work is a transition theory over the consistency lattice. The uniform lattice framework assists in the development of the transition theory because any two models can be compared by their properties, and transitions can be viewed as adding or removing properties. The transition theory was evaluated against synchronized consistency models, and every synchronized model described in Section 2.3 can be modeled by this transition theory. This led to the development of a single statement of consistency called generalized consistency. Under generalized consistency, every operation is labeled with a set of consistency properties. Consistency requirements among operations depend on their labelings. If every operation is labeled with the same set of properties, generalized consistency simulates the nonsynchronized consistency model represented by the combination of those properties. Various other labelings simulate the transitions equivalent to the synchronized models.

In the future, this work can be extended in several directions. In the lattice, the five new consistency models need to be examined to determine intuitive definitions of the effects enforced by those models, and whether existing applications may be able to take better advantage of the new models. The space of consistency models around processor consistency needs to be explored in more detail as well as other methods of combining properties such as  $GPO \cap GDO$  and  $GPAO$ . Finally, efficient implementations could be examined with regards to what consistency properties they enforce. A lattice of implementations related to the lattice of consistency models would be helpful in automating selection of memory implementations.

## Appendix

### A. Shared Memory Formalism

*Definition A.1.* An *execution* is a set of *processes*,  $P$ , a set of *shared variables*,  $V$ , a set of *operations*,  $O$ , and two partial orders on  $O$ , *process order*,  $<_{PO}$ , and *writes-to order*,  $\mapsto$ .

*Definition A.2.* An *operation* is a tuple  $(op, i, x, v)$  where  $op$  is  $r$  for a read or  $w$  for a write.  $i \in P$  is the process submitting the operation.  $x \in V$  is the variable to which the operation is applied, and  $v$  is a valid value for the variable  $x$ . We also define the following convenience functions:

$$\begin{aligned} operation((op, i, x, v)) &\equiv op, \\ process((op, i, x, v)) &\equiv i, \end{aligned}$$

$$\begin{aligned} \text{variable}((op, i, x, v)) &\equiv x, \text{ and} \\ \text{value}((op, i, x, v)) &\equiv v. \end{aligned}$$

**Definition A.3.** An *operation pattern* is a way of representing a subset of  $O$ . It is written as a tuple  $(op, i, x, v)$  like an operation, and the same convenience functions are defined as for operations, but one or more of its elements has the value  $*$ . An operation pattern represents the set of all operations in  $O$  that match the pattern in all attributes that are not  $*$ . Formally, if  $O'$  is an operation pattern and  $o$  is an operation, then

$$\begin{aligned} o \in O' \text{ iff } o \in O \wedge & (\text{operation}(O') = \text{operation}(o) \vee \text{operation}(O') = *) \\ & \wedge (\text{process}(O') = \text{process}(o) \vee \text{process}(O') = *) \\ & \wedge (\text{variable}(O') = \text{variable}(o) \vee \text{variable}(O') = *) \\ & \wedge (\text{value}(O') = \text{value}(o) \vee \text{value}(O') = *). \end{aligned}$$

For example,  $(r, p_1, x, 5)$  denotes that process  $p_1$  read the variable  $x$ , and received the value 5.  $(w, *, *, *)$  denotes the set of all write operations.

**Definition A.4.**  $O$  is a set of operations such that

$$(\forall_{r \in (r, *, *, *)} | \{ (w, *, \text{variable}(r), \text{value}(r)) \} | = 1) \bigwedge (\forall_{x \in V} (w, \epsilon, x, \perp) \in O),$$

where  $\epsilon$  is a special symbol not used to denote any process, and  $\perp$  is a special value that cannot be written by any process. The operation  $(w, \epsilon, x, \perp)$  is called the *initial write* of  $x$ .

**Definition A.5.** *Local order* for process  $i$ ,  $<_i$ , is a relation such that

$$(\forall_{o_i, o_j \in (*, i, *, *)} o_i <_i o_j \oplus o_j <_i o_i) \wedge (\forall_{x \in V, o \in (*, i, *, *)} (w, \epsilon, x, \perp) <_i o).$$

**Definition A.6.** *Process order*,  $<_{PO}$ , is a relation such that

$$<_{PO} \equiv \bigcup_{i \in P} <_i.$$

**Definition A.7.** *Writes-to order*,  $\mapsto$ , is a relation such that

$$\forall_{r \in (r, *, *, *)} \exists_{i \in P} (w, i, \text{variable}(r), \text{value}(r)) \mapsto r.$$

These definitions say that the set  $O$  includes a unique write for every read and an initial write for every variable. Operations by a single process are totally ordered and are ordered after all initial writes by local order. Process order is the union of all local orders. Without loss of generality, we assume that every variable has an initial write, and writes are uniquely valued. As a consequence of this, for every read there exists exactly one write that writes-to that read. Writes-to order is redundant with the values returned by read operations. Knowing either one determines the other, but both are defined for convenience.

An execution defines the operations that were submitted to a memory system and the externally visible behavior of the memory system which consists of the values returned to all reads. Now we need to relate the behavior of the memory system to correctness with respect to a consistency model. Consider Figure 21. Execution (a) corresponds to a sequentially consistent execution. From the set of operations,  $O$ , and the process order we see that  $p_1$  wrote  $x$  and then read  $y$ , and  $p_2$  read  $x$  and then wrote  $y$ . From the writes-to order we see that  $p_2$  read  $p_1$ 's write, and  $p_1$  read

$ \begin{aligned} P &= \{p_1, p_2\} \\ V &= \{x, y\} \\ O &= \{(w, \epsilon, x, \perp), (w, \epsilon, y, \perp), \\ &\quad (w, p_1, x, 1), (r, p_1, y, 2), \\ &\quad (r, p_2, x, 1), (w, p_2, y, 2)\} \\ &(w, \epsilon, x, \perp) <_{PO} (w, p_1, x, 1) <_{PO} (r, p_1, y, 2) \\ &(w, \epsilon, y, \perp) <_{PO} (w, p_1, x, 1) <_{PO} (r, p_1, y, 2) \\ &(w, \epsilon, x, \perp) <_{PO} (r, p_2, x, 1) <_{PO} (w, p_2, y, 2) \\ &(w, \epsilon, y, \perp) <_{PO} (r, p_2, x, 1) <_{PO} (w, p_2, y, 2) \\ &(w, p_1, x, 1) \mapsto (r, p_2, x, 1) \\ &(w, p_2, y, 2) \mapsto (r, p_1, y, 2) \end{aligned} $	$ \begin{aligned} P &= \{p_1\} \\ V &= \{x\} \\ O &= \{(w, \epsilon, x, \perp), (w, p_1, x, 1), \\ &\quad (w, p_1, x, 2), (r, p_1, x, 1)\} \\ &(w, \epsilon, x, \perp) <_{PO} (w, p_1, x, 1) <_{PO} \\ &\quad (w, p_1, x, 2) <_{PO} (r, p_1, x, 1) \\ &(w, p_1, x, 1) \mapsto (r, p_1, x, 1) \end{aligned} $
(a)	(b)

FIG. 21. Two executions.

$p_2$ 's write. This corresponds to a sequential order of:

$$\begin{aligned}
(w, \epsilon, x, \perp) &< (w, \epsilon, y, \perp) < (w, p_1, x, 1) < (r, p_2, x, 1) \\
&< (w, p_2, y, 2) < (r, p_1, y, 2).
\end{aligned}$$

where  $<$  denotes an unnamed total order. Execution (b), however, is a little disconcerting. There is one process.  $p_1$  wrote 1 to  $x$ , then wrote 2 to  $x$ , and then read  $x$ . Unfortunately, the read returned the value 1 from the first write, and not 2 from the second. When we try to create a total order we run into a contradiction. If the order is:

$$(w, \epsilon, x, \perp) < (w, p_1, x, 1) < (w, p_1, x, 2) < (r, p_1, x, 1),$$

then the read does not read from the most recent write, but if the order is:

$$(w, \epsilon, x, \perp) < (w, p_1, x, 1) < (r, p_1, x, 1) < (w, p_1, x, 2),$$

then this violates process order. The important thing to note is that this does qualify as an execution. Imagine a computer with out of order instruction dispatching. If this dispatching mechanism were buggy it might accidentally switch the order of a read and write to the same variable. Execution (b) exactly models this sort of phenomenon. However, it is not likely that this execution will be deemed correct by any consistency model. The problems we just saw with creating a total order also give us a hint about how to define a consistency model in terms of allowable executions.

**Definition A.8 (The view property).** If  $V$  is a relation and  $O'$  is a set of operations, then

$$\begin{aligned}
\text{view}(V, O') \text{ iff } O' \subseteq O \wedge V(o_i, o_j) \rightarrow o_i, o_j \in O' \wedge \forall_{o_i, o_j \in O'} V(o_i, o_j) \oplus \\
V(o_j, o_i).
\end{aligned}$$

**Definition A.9 (The serial Property).**

$$\begin{aligned}
\text{serial}(V, O') \text{ iff } \text{view}(V, O') \wedge \forall_{r \in (r, *, *, *) \cap O'} (\exists_{w_i \mapsto r} V(w_i, r) \wedge \\
\neg \exists_{w_j \in (w, *, \text{variable}(r), *) \cap O'} (V(w_i, w_j) \wedge V(w_j, r)))
\end{aligned}$$

A view is merely a total order on a subset of  $O$  intended to represent one process' view of the order of read and write operations. A view may only contain a subset of  $O$  because a process might not see every operation. In particular, processes usually do not see each other's reads. We explicitly define it because the correctness of an

execution depends on whether it is possible to construct a view that might have happened which is compatible with the rules of the consistency model. A view is serial if every read reads from the most recent write in that view. This is considered a fundamental correctness requirement for all consistency models.

*Definition A.10 (The respects Property).* If  $V$  and  $R$  are relations and  $O'$  is a set of operations, then

$$\text{respects}(V, R, O') \text{ iff } \text{view}(V, O') \wedge \forall_{o_i, o_j \in O'} R(o_i, o_j) \rightarrow V(o_i, o_j).$$

*Definition A.11.* A relation,  $R$ , can be *restricted* to a subset of operations,  $O'$  by the *restrict* function. If  $R' = \text{restrict}(R, O')$ , then  $R'$  is a relation such that

$$R'(o_i, o_j) \text{ iff } R(o_i, o_j) \wedge o_i, o_j \in O'.$$

The notation,  $\text{SerialView}(R|O')$ , is used as a shorthand to denote a view  $V$  such that  $\text{serial}(V, O') \wedge \text{respects}(V, \text{restrict}(R, O'), O')$ . Usually,  $O'$  will be defined in terms of operation patterns, or if  $O' = O$  the shorthand  $\text{SerialView}(R)$  will be used.

The differences between this formalism and that used in Collier [1992] are:

- (1) Our model specifies the initial values of variables as explicit initial writes.
- (2) The output of the memory system in our model is the values read by every read operation instead of just the terminal value of each variable.
- (3) Our operation tuples omit the elements of line number and store number.
- (4) Our model uses views while Collier's uses graph sets.
- (5) Collier assumes a program is specified as a list of assignment statements implemented by reads and writes while in our model we assume a program is specified as reads and writes.

The differences are cosmetic. Initial values for variables and guaranteed initial writes are interchangeable. For the memory system output we can simulate Collier's model by reading all variables as the last act of any program, and Collier can simulate our model by saving the results of each read in a unique variable. The examples in Collier [1992] have a single terminal value for each variable which seems to imply that the values of all stores must agree at the end of the program. This would prevent modeling of some non cache consistent models, but the problem can easily be solved by assuming a separate terminal value for each variable in each store.

For the third difference, we assume operations are identifiable without line number information. In particular, we assume uniquely valued writes. Collier models nonatomic writes as multiple writes, one to the store of each processor. We model nonatomic writes by allowing each processor to create a separate view of the same write operations in a different order. As a result we have no need to specify stores. A graph set is a set of all possible orders of operations with dependency edges. At least one graph in the set must be acyclic for the execution to satisfy a consistency model. Instead of enumerating every possible order of events we use the existential quantifier. Our views are equivalent to Collier's graph sets.

Assuming that reads and writes are joined in assignment statements creates Collier's SRW rule of computation. In our model, if a write depends on a read for its value, it is modeled as a local order dependency. The combined effect of the other (CWR, CRW, CWW) rules of computation and the value function is to eliminate orderings where the terminal values of the variables are incompatible

with reads reading from the most recent write. In our model this is accomplished by requiring views to be serial.

### B. *Proof that GDO is Equivalent to Cache Consistency*

LEMMA B.1. *If an execution is Cache Consistent, then Data Order is acyclic.*

PROOF. Data order only contains edges between pairs of operations to the same variable. Therefore, if data order were cyclic, the cycle would have to involve only operations to a single variable. Cache consistency requires for every variable a serial view respecting process order on all the operations to that variable. We will show that these views must also respect data order, and so data order is acyclic.

The cache consistent view for a variable respects process order by definition and writes-to order because it is serial so it respects the first two conditions of data order. The third condition of data order must also be respected. If  $o_1$  is process ordered before  $r$ , it must come before  $r$ , in the view. If, in addition,  $o_1$  has a different value than  $r$ , and  $o_2$  writes to  $r$ , then  $o_1$  must come before  $o_2$  in the view. If not and  $o_1$  is a write, then  $r$  does not read from the most recent write so the view is not serial. If not and  $o_1$  is a read then either  $o_2$  is the most recent write before  $o_1$ , in which case  $o_1$  does not read from the most recent write, or there is another write between  $o_2$  and  $o_1$ . This write is also between  $o_2$  and  $r$ , and it has the same value as  $o_1$  which is different than  $r$ , so  $r$  does not read from the most recent write and the view is not serial. Since the view is a total order and it respects the first three conditions of data order, it must respect their transitive closure which is the fourth condition of data order.

The views required for cache consistency must respect data order. If data order contained a cycle, then the view for some variable could not be constructed and the execution would not be cache consistent. Therefore, if the execution is cache consistent, data order is acyclic.  $\square$

LEMMA B.2. *If two reads are ordered by data order then either they are by the same process, or they are ordered by a transitive chain containing a write.*

PROOF. Two reads cannot be ordered by writes-to, or by having one write to a read that the other is process ordered before. So the only way two reads can be data ordered is by process order, or a transitive chain. If two reads are not by the same process, and are data ordered, take the last operation in the transitive chain. If this operation is a write, it satisfies the lemma. Otherwise, it must be a read by the same process as the final read. By the same logic the next to last operation in the chain must also be a write, or a read by the same process as the final read. By induction, if there is no write in the chain, then the first operation in the chain must be a read by the same process as the final read. Therefore, if the two reads are not by the same process, there must be a write in the transitive chain.  $\square$

LEMMA B.3. *If an execution is GDO, then data order is acyclic.*

PROOF. GDO requires a view for every process that is serial and respects data order over the subset of all operations by that process plus all writes. If these views are constructible, then data order must be acyclic at least on the subsets of operations in each view. Therefore, if data order is cyclic, then the cycle must contain at least two read operations by different processes,  $r_1$  and  $r_2$ , such that

$r_1 <_{DO} r_2$  and  $r_2 <_{DO} r_1$ . By Lemma B.2, these two reads must be ordered by two transitive chains, and each chain must contain a write. Because data order is a transitive closure, there must be a cycle between the writes in the two chains. This makes it impossible to construct the views required for GDO because every view must include all writes. If data order is cyclic, then the views required for GDO cannot be constructed. Therefore, if an execution is GDO, then data order is acyclic.  $\square$

LEMMA B.4. *If data order is acyclic, then*

$$\forall_{x \in V} \exists <_x = \text{SerialView}(<_{DO} | (*, *, x, *)).$$

PROOF. First, collect all the operations on a single variable and place them into groups where each group contains a write and all reads that the write writes-to. Order the operations in each group in an order that respects data order. This is possible because data order is acyclic. The reads in a group all read from the write in that group so the write will be ordered first in each group. The serial view for that variable is constructed by ordering the groups with no interleaving of operations between different groups. For every read, the most recent write to the same variable must be the write from its group which is the write which wrote-to it so the view must be serial. Any order of the groups with no interleaving will produce a serial view.

If  $G_1$  and  $G_2$  are two groups, then define group order,  $<_{GO}$  as:  $G_1 <_{GO} G_2$  iff  $\exists o_1 \in G_1, o_2 \in G_2$  such that  $o_1 <_{DO} o_2$ . If group order is acyclic, then any topological sort on group order will produce a view that respects data order and is serial. Assume there is a cycle in group order, but not in data order. Take any two ordered groups from the cycle,  $G_1 <_{GO} G_2$ . We will show that the writes from the groups,  $w_1$  and  $w_2$ , must be ordered by data order. Therefore, any cycle in group order must be accompanied by a cycle in data order. So if data order is acyclic, then group order must be acyclic and the views can be constructed.

There must be operations from the two groups such that  $o_1 <_{DO} o_2$ . Either  $o_1$  is  $w_1$ , or  $o_1$  is a read that  $w_1$  writes-to. So  $w_1 <_{DO} o_2$ . Also, either  $o_2$  is  $w_2$  in which case  $w_1 <_{DO} w_2$ , or  $o_2$  is a read that  $w_2$  writes-to. If  $o_2$  is a read, consider how it came about that  $w_1$  is data ordered before  $o_2$ .  $w_1$  did not write to  $o_2$  so either  $w_1 <_{PO} o_2$ , or they are ordered by a transitive chain. If  $w_1 <_{PO} o_2$ , then  $w_1 <_{DO} w_2$  because  $w_2 \mapsto o_2$ . If not, let  $o$  be the last operation in the transitive chain so  $w_1 <_{DO} o <_{DO} o_2$ . Either  $o \mapsto o_2$  in which case  $o$  is  $w_2$  and  $w_1 <_{DO} w_2$  or  $o <_{PO} o_2$  in which case  $o <_{DO} w_2$  because  $w_2 \mapsto o_2$  so  $w_1 <_{DO} w_2$ .

Therefore, if  $G_1 <_{GO} G_2$ , then  $w_1 <_{DO} w_2$ . Any cycle in group order will be reflected in data order by the writes. If data order is acyclic, then there can be no cycle in group order, and a topological sort of the groups respecting group order will produce the required serial view respecting data order for that variable.  $\square$

LEMMA B.5. *If the serial views,  $<_x$ , defined in Lemma B.4 exist, then*

$$\forall_{i \in P} \cup_{x \in V} <_x \cup <_i \text{ is acyclic.}$$

PROOF. Assume that a cycle exists for some process,  $i$ . The views,  $<_x$ , are acyclic, and their union cannot contain a cycle because no operation is in more than one view. Therefore, the cycle must have an edge in the local order, and thus include operations by process  $i$ . Pick any operation by process  $i$  in the cycle. Call it  $o$ . Follow the edges that make up the cycle. If you follow an edge in local order,

then you must reach an operation by process  $i$  that occurs after  $o$  in local order. If you follow an edge not in local order, then you must reach an operation to the same variable as  $o$  by a process other than  $i$ . Operations by other processes are not ordered by local order, and thus the cycle must proceed through operations to the same variable following edges of  $<_x$  for that variable until reaching an operation by process  $i$ . This operation must be to the same variable as  $o$ , and must be ordered after  $o$  by local order. Otherwise, the view for that variable would not respect data order.

In any case, the first operation by process  $i$  encountered after  $o$  in the cycle must be after  $o$  in local order. Call this operation  $o'$ . By the same logic, the next operation by process  $i$  after  $o'$  in the cycle must be ordered after  $o'$  and  $o$  by local order. By induction, every operation by process  $i$  in the cycle must be after  $o$  in local order. Eventually, the cycle will reach  $o$  itself showing that there is a cycle in local order which is a contradiction.  $\square$

LEMMA B.6. *If the views,  $<_x$ , defined in Lemma B.4 exist, then the execution is GDO.*

PROOF. Construct the view for process  $i$  required for GDO as any topological sort of  $(*, i, *, *) \cup (w, *, *, *)$  respecting  $\cup_{x \in V} <_x \cup <_i$ . This is possible because by Lemma B.5 the relation is acyclic. The views will be serial because the views,  $<_x$ , are serial, and the relative position of all pairs of operations to the same variable is preserved. Data order only contains edges between two operations to the same variable and so is a subset of  $\cup_{x \in V} <_x$ . Therefore, the constructed views respect local order and data order, and they are serial so the execution is GDO.  $\square$

LEMMA B.7. *If the views,  $<_x$ , defined in Lemma B.4 exist, then the execution is cache consistent.*

PROOF. Process order restricted to the set of operations on a single variable is a subset of data order. Therefore, any view on  $(*, *, x, *)$  that respects data order will also respect process order. Therefore, the views defined in Lemma B.4 respect process order, and so prove that the execution is cache consistent.  $\square$

THEOREM B.8. *An execution is Cache Consistent iff it is GDO iff data order is acyclic.*

PROOF. Follows directly from Lemmas B.1, B.3, B.4, B.6, and B.7.  $\square$

### C. Proof that GAO Is Strictly Stronger than GDO

LEMMA C.1. *If data order has a cycle, then the execution is not GAO.*

PROOF

*Case 1. The Cycle Has a Read.* Take the operation immediately before the read in the cycle. If it is linked by a transitive chain, add that transitive chain to the cycle. Repeat until the operation immediately before the read is linked directly without a transitive chain. This is either a write, or by Lemma B.2, it is a read ordered by process order. If it is a read, repeat until a write is reached. A write must be reached because otherwise the cycle will return to the original read which must be ordered before itself by process order which is a contradiction. The write that is reached is directly ordered by data order before the next operation in the cycle which is a read. They cannot be ordered by condition 3 of data order because this would imply

that there exists a third operation such that the write is process ordered before that operation, and the read writes to that operation. This is impossible since a read cannot write to another operation. So the write must be ordered before the read by process order or writes-to order. Also, the operations are in a cycle in data order so the read is data ordered before the write. In either case, the write is anti ordered before itself. The serial views for GAO must all contain this write, so they cannot respect this cycle in anti order. Therefore, the execution is not GAO.

*Case 2. The Cycle Has no Reads.* Once again, expand the cycle so that no link is a transitive chain. If the transitive chain includes a read refer to Case 1. None of the links can result from writes-to order because a write cannot write to another write. The cycle must contain writes from at least two processes. If not, a write must be ordered by data order before another write earlier in process order. This must have come about by condition 3 of data order. Therefore, the following condition exists,  $w_1 <_{PO} w_2 <_{PO} r$ , and  $w_1 \mapsto r$ . all of these operations are to the same variable. It is impossible for this process' view to be serial and respect local order, so the execution is not GAO. So there must be some links that result from condition 3 of data order between writes by different processes. Pick one write,  $w_1$  and follow the cycle along process order links until a link resulting from condition 3 is reached. In this case, a write,  $w_2$  is process ordered before a read,  $r$ , which is written-to by another write,  $w_3$ , creating the link  $w_2 <_{DO} w_3$ .  $w_1$  must also be process ordered before  $r$  because either it is process ordered before  $w_2$ , or it is  $w_2$ . So  $w_1 <_{DO} w_3$ . Now,  $w_1$  does not write to  $r$ , so it must be ordered by serial order either  $w_1 <_{SO} w_3 \mapsto r$ , or  $r <_{SO} w_1$ . The second case is impossible. The view for the process that submitted  $w_1$  and  $r$  must contain both operations and respect local order. The assignment of  $r <_{SO} w_1$  would prevent this, and so this assignment could never be used to show that the execution is GAO. Therefore, the assignment must be  $w_1 <_{SO} w_3$ . By the same logic, follow the chain from  $w_3$  to the next link that results from Condition 3. There must be another write serial ordered after  $w_3$ . Every time the cycle switches to an operation by a different process, the first operation by the new process must be serial ordered after  $w_3$ . Continue around the cycle. At some point the cycle will change processes for the last time before reaching  $w_1$ . The first operation by this new process is either  $w_1$ , or a write process ordered before  $w_1$ . This write must also be serial ordered before  $w_3$ . Either it is  $w_1$ , or it is process ordered before  $r$ , and the same reasoning applies. This assignment of serial order has a cycle involving only writes, and so no process' view could respect it. We have previously shown that any alternate assignment would also prevent the execution from satisfying GAO. Therefore, the execution is not GAO.  $\square$

**THEOREM C.2.** *GAO is strictly stronger than GDO.*

**PROOF.** GAO is shown to be stronger by Theorem B.8 and Lemma C.1. GAO is shown to be strictly stronger by the fact that the execution in Figure 12 satisfies GDO and not GAO.  $\square$

## REFERENCES

- ADVE, S. V., AND GHARACHORLOO, K. 1996. Shared memory consistency models: A tutorial. *IEEE Comput. Mag.* 29, 12 (Dec.), 66–76.
- ADVE, S. V., AND HILL, M. D. 1993. A unified formalization of four shared-memory models. *IEEE Trans. Paral. Distrib. Syst.* 4, 6 (June), 613–624.



- AHAMAD, M., BAZZI, R., JOHN, R., KOHLI, P., AND NEIGER, G. 1993. The power of processor consistency. In *Proceedings of the 5th ACM Symposium on Parallel Algorithms and Architectures*. ACM, New York, 251–260.
- AHAMAD, M., BURNS, J. E., HUTTO, P. W., AND NEIGHER, G. 1991. Causal memory. In *Proceedings of the 5th International Workshop on Distributed Algorithms*, S. Toueg, G. Spirakis, P., and L. Kirousis, Eds. Lecture Notes in Computer Science, vol. 579. Springer-Verlag, New York, 9–30.
- AMZA, C., COX, A. L., DWARKADAS, S., KELEHER, P., LU, H., RAJAMONY, R., YU, W., AND ZWANEOEL, W. 1996. Treadmarks: Shared memory computing on networks of workstations. *IEEE Comput. Mag.* 29, 2 (Feb.), 18–28.
- BATALLER, J., AND BERNABEU, J. 1997. Synchronized DSM models. In *Proceedings of the 3rd International Euro-Par Conference*, C. Lengauer, M. Griebel, and S. Gorlatch, Eds. Springer, Berlin, Germany, 468–475.
- BATALLER, J., AND BERNABEU-AUBAN, J. M. 1998. Adaptable distributed shared memory: A formal definition. In *Proceedings of the 4th International Euro-Par Conference*, D. Pritchard and J. Reeve, Eds. Springer, Berlin, Germany, 887–891.
- BENNETT, J. K., CARTER, J. K., AND ZWAENEOEL, W. 1990. Munin: Distributed shared memory based on type-specific memory coherence. In *Proceedings of the 2nd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*. ACM, New York, 168–176.
- BENNETT, J. K., CARTER, J. K., AND ZWAENEOEL, W. 1995. Techniques for reducing consistency-related communication in distributed shared memory systems. *ACM Trans. Comput. Syst.* 13, 3 (Aug.), 205–243.
- BERSHAD, B. N., AND ZEKAUSKAS, M. 1991. Midway: Shared memory parallel programming with entry consistency for distributed memory multiprocessors. Tech. Rep. CMU-CS-91-170, Carnegie-Mellon University, Pittsburgh, Penn.
- BERSHAD, B. N., ZEKAUSKAS, M. J., AND SAWDON, W. A. 1993. The midway distributed shared memory system. In *Proceedings of the IEEE COMPCON Conference*. IEEE Computer Society Press, Los Alamitos, Calif., 528–537.
- BLUMOFF, R. D., FRIGO, M., JOERG, C. F., LEISEN, C. E., AND RANDALL, K. H. 1996. Dag-consistent distributed shared memory. In *Proceedings of the 10th International Parallel Processing Symposium*.
- COLLIER, W. W. 1992. *Reasoning About Parallel Architectures*. Prentice-Hall, Englewood Cliffs, N.J.
- DUBOIS, M., AND SCHEURICH, C. 1990. Memory access dependencies in shared memory multiprocessors. *IEEE Trans. Softw. Eng.* 16, 6 (June), 660–673.
- DUBOIS, M., SCHEURICH, C., AND BRIGGS, F. 1986. Memory buffering in multiprocessors. In *Proceedings of the 13th Annual Symposium on Computer Architecture*. IEEE Computer Society Press, Los Alamitos, Calif., 434–442.
- FRIGO, M. 1997. The weakest reasonable memory model. Ph.D. dissertation, Massachusetts Institute of Technology, Cambridge, Mass.
- GAO, G. R., AND SARKAR, V. 2000. Location consistency—A new memory model and cache consistency protocol. *IEEE Trans. Comput.* 49, 8 (Aug.), 798–813.
- GHARACHORLOO, K., LENOSKI, D., LAUDON, J., GIBBONS, P., GUPTA, A., AND HENNESSY, J. 1990. Memory consistency and event ordering in scalable shared-memory multiprocessors. In *Proceedings of the 17th Annual International Symposium on Computer Architecture*. IEEE Computer Society Press, Los Alamitos, Calif., 15–26.
- GNIADY, C., FALSAFI, B., AND VIJAYKUMAR, T. N. 1999. Is sc + ilp = rc? In *Proceedings of the 26th Annual International Symposium on Computer Architecture*. IEEE Computer Society Press, Los Alamitos, Calif., 162–171.
- GOODMAN, J. R. 1989. Cache consistency and sequential consistency. Tech. Rep. 61. IEEE Scalable Coherent Interface Working Group. Mar.
- GOSLING, J., JOY, B., AND STEELE, G. 1996. *The Java Language Specification*. Addison-Wesley, Reading, Mass.
- HERLIHY, M. P., AND WING, J. M. 1990. Linearizability: A correctness condition for concurrent objects. *ACM Trans. Program. Lang. Syst.* 12, 3 (July), 463–492.
- HUTTO, P. W., AND AHAMAD, M. 1990. Slow memory: Weakening consistency to enhance concurrency in distributed shared memories. In *Proceedings of the 10th International Conference on Distributed Computing Systems*. 302–309.
- IFTODE, L., SINGH, J. P., AND LI, K. 1996. Scope consistency: A bridge between release consistency and entry consistency. Tech. Rep., Princeton University, Princeton, N.J.

- KELEHER, P., COX, A. L., AND ZWAENPOEL, W. 1992. Lazy release consistency for software distributed shared memory. In *Proceedings of the 19th Annual International Symposium on Computer Architecture*. IEEE Computer Society Press, Los Alamitos, Calif., 13–21.
- LAMPORT, L. 1978. Time, clocks, and the ordering of events in a distributed system. *Commun. ACM* 21, 7 (July), 558–565.
- LAMPORT, L. 1979. How to make a multiprocessor computer that correctly executes multiprocess programs. *IEEE Trans. Comput.* C-28, 9 (Sept.), 690–691.
- LAMPORT, L. 1986. On interprocess communication; Part II: Algorithms. *Distrib. Comput.* 1, 2 (Apr.), 86–101.
- LENOSKI, D., LAUDON, J., GHARACHORLOO, K., GUPTA, A., AND HENNESSY, J. 1990. The directory-based cache coherence protocol for the dash multiprocessor. In *Proceedings of the 17th Annual International Symposium on Computer Architecture*. IEEE Computer Society Press, Los Alamitos, Calif., 148–159.
- LI, K. 1986. Shared virtual memory on loosely coupled multiprocessors. Ph.D. dissertation, Yale University.
- LI, K., AND HUDAK, P. 1989. Memory coherence in shared virtual memory systems. *ACM Trans. Comput. Syst.* 7, 4 (Nov.), 321–359.
- LIPTON, R. J., AND SANDBERG, J. S. 1988. Pram: A scalable shared memory. Tech. Rep. CS-TR-180-88, Princeton University. Sept.
- MOSBERGER, D. 1993. Memory consistency models. *ACM SIGOPS Operating Systems Review* 27, 1 (Jan.), 18–27.
- RANGANATHAN, P., PAI, V. S., ABDEL-SHAFI, H., AND ADVE, S. V. 1997a. The interaction of software prefetching with ilp processors in shared-memory systems. In *Proceedings of the 24th Annual International Symposium on Computer Architecture*. IEEE Computer Society Press, Los Alamitos, Calif., 144–156.
- RANGANATHAN, P., PAI, V. S., AND ADVE, S. V. 1997b. Using speculative retirement and larger instruction windows to narrow the performance gap between memory consistency models. In *Proceedings of the 9th ACM Symposium on Parallel Algorithms and Architectures*. ACM, New York, 199–210.
- ROYCHOUDBURY, A., AND MITRA, T. 2002. Specifying multithreaded java semantics for program verification. In *Proceedings of the 24th International Conference on Software Engineering*. 489–499.
- TANENBAUM, A. S. 1995. *Distributed Operating Systems*. Prentice-Hall, Englewood Cliffs, N.J.

RECEIVED AUGUST 2002; REVISED MARCH 2004; ACCEPTED MARCH 2004