

Analyzing Massive Data Sets

Summer Semester 2019

Prof. Dr. Peter Fischer

Institut für Informatik

Lehrstuhl für Datenbanken und Informationssysteme

Chapter 4: Finding Similar Items

High-Dimensional Data and Similarity

- First *conceptual* and *algorithmic* part of the lecture
- Two core concepts:
 - **High-Dimensional Data**: Data items represented by many data points (hundreds, thousands, ... possibly out of a much large space)
 - Analyzing a *single* or *few* dimensions insufficient to understand items
 - **Similarity/Distance**: Expressing pair-wise similarity over all features
- Applications:
 - **Finding Similar Items**: pairwise (this chapter)
 - **Clustering**: Identify structure / groups using similarity
 - **Retrieval**: Similarity between search expression and data set
- Strategies for massive volumes:
 - **Exact solutions** are **costly** – but there are several strategies to help
 - **Approximate solutions** more **feasible** – e.g., multiple hashes

Similar Items - A Common Metaphor

- **Many problems can be expressed as finding “similar” sets:**
 - Find *near-neighbors* in *high-dimensional* space
- **Examples:**
 - **Pages with similar words (around 1.7% - 7% pages on the web)**
 - Mirror pages
 - Common source pages
 - Plagiarism
 - Classification by topic
 - **Customers who purchased similar products**
 - Foundation of recommendation systems (think Amazon)
 - Products with similar customer sets
 - **Media with similar content (images, music, videos)**
 - Duplicate removal
 - Recommendations

Defining Similarity

- Representing the data items:

*What is **specific** about items in a collections?*

- Text: Character Frequency, Lexical, Structure (sentence, chapters), Semantics/Meaning
- Images: Colors, Structure, Objects, ...
- Music: Pitch, Melody, Metric/Rhythm, Modulation, ...

- Expressing Similarity:

*Given the features, what makes items **close**?*

- Shared features
- Numerically similar features values
- Relevance of certain features
- Combination of features

Computing Similarity


- **Given: High dimensional data points x_1, x_2, \dots**
 - **For example:** Image is a long vector of pixel colors
$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 2 & 1 \\ 0 & 1 & 0 \end{bmatrix} \rightarrow [1 \ 2 \ 1 \ 0 \ 2 \ 1 \ 0 \ 1 \ 0]$$
- **And some distance function $d(x_1, x_2)$**
 - Which quantifies the “distance” between x_1 and x_2
- **Goal:** Find **all pairs of data points (x_i, x_j)** that are within some distance threshold $d(x_i, x_j) \leq s$
- **Note:** Naïve solution would take $O(N^2)$ ☹
where N is the number of data points
- Documents are so large or so many that they cannot fit in main memory
- **An exact solution is possible in $O(N \log N)$**
- **An approximate version can be done in $O(N)!!$**
- **How?**

Easier problem: Identical copies

- Relatively straightforward task
- Naive strategy:
 - Enumerate all pairs $\frac{n^2}{2} \rightarrow O(n^2)$
 - Do a bitwise comparison, e.g. `cmp`: True or False
- Can be solved with cost $O(n)$
- How?
- Apply a hash function on every element
- Elements with same content will be in the same bucket
- Caveat: Other direction may not hold:
same hash value does not imply same content
- -> Need to check for collisions
- Scales well: hashing can be done fully parallel
- Hash tables can be distributed

Similar problem: Join processing

- $A \times B + \sigma_{A.a \text{ comp } B.b}$
- For all pairs, apply a comparison \Rightarrow again $O(n^2)$
- Joins can reduce the search space for certain predicates:



Strategies to reduce search space

- Join search space reduction
 - Indexing: Allow faster access to matching "inner" elements
 - Sorting: order one or both sides
 - Hashing: reduce space to items with same hash value
- Will not get $O(n)$ complexity in the general case
 - Common Indexes (B+-Trees) have $O(n \log n)$ creation and access cost
 - Sorting costs $O(n \log n)$
 - Indexes and sorting don't work well with many dimensions
 - Hashing only works for equality
- Is this a dead-end?
- No, we but can play with these concepts!

Expressing Similarity

Distance Measures (first take)


■ Goal: Find near-neighbors in high-dim. space

- We formally define “near neighbors” as points that are a “small distance” apart
- For each application, we first need to define what “**distance**” means
- Common for texts: Jaccard distance/similarity

- The **Jaccard similarity** of two **sets** is the size of their intersection divided by the size of their union:

$$\text{sim}(C_1, C_2) = \frac{|C_1 \cap C_2|}{|C_1 \cup C_2|}$$

- **Jaccard distance:** $d(C_1, C_2) = 1 - \frac{|C_1 \cap C_2|}{|C_1 \cup C_2|}$



3 in intersection

8 in union

Jaccard similarity = 3/8

Jaccard distance = 5/8

Requirements for distance functions

Two data instances/points:

- $d(x, y) \geq 0$ (no negative distances)
- $d(x, y) = 0$ if and only if $x = y$ (distances are positive, except for the distance from a point to itself)
- $d(x, y) = d(y, x)$ (distance is symmetric)
- $d(x, y) \leq d(x, z) + d(z, y)$ (the triangle inequality)
- Triangle Inequality:
 - no gain from a "detour", distance describes the shortest path
 - Hardest to prove, most often violated by candidates

"Base type" for measurements

- Set Membership / Binary Variables
(does the other data items contain the same features)
- Vector (Ordered Set)
- Spatial/Numeric values in Vector space
- (String) Editing: number of operations to transform data item into the other
- Graphs: common nodes and edges? Common label names?
- Time Series: sequences with timestamps, align shapes, time shifts, value shifts, ...
- ...

Set Membership

- Basic Idea:
How many members of one set are also member of the other set?
- Also used for binary variables
(0/1, True/False, Present/Absent)
- Most well-known measure - Jaccard

$$J(C_1, C_2) = 1 - \frac{|C_1 \cap C_2|}{|C_1 \cup C_2|}$$


- Many variants, not all true distance functions
 - **Sørensen–Dice**: denominator sum of sizes (~**F1 score**)
 - **Simple Matching coefficient, Rand**: also mutual absence
- Alternative: Hamming - symmetric difference

$$H(x, y) = |(x - y) \cup (y - x)|$$

Numeric / "Spatial" Distances (1)

- Idea:
 - Observe component-wise difference
 - Scale/normalize when needed
- Euclidean aka "Straight-Line":

$$E(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$



- Based on Pythagorean formula for triangles: $a^2 + b^2 = c^2$

Numeric / "Spatial" Distances (2)

- Manhattan (aka city block, taxicab)

$$M(x, y) = \sum_{i=1}^n |x_i - y_i|$$


- Weighted version: Canberra


$$CB(x, y) = \sum_{i=1}^n \frac{|x_i - y_i|}{|x_i| + |y_i|}$$

- Chebyshev/Chessboard (maximum component):

$$C(x, y) = \max_i (|x_i - y_i|)$$

- Observe a pattern?



	a	b	c	d	e	f	g	h	
8	5	4	3	2	2	2	2	2	8
7	5	4	3	2	1	1	1	2	7
6	5	4	3	2	1		1	2	6
5	5	4	3	2	1	1	1	2	5
4	5	4	3	2	2	2	2	2	4
3	5	4	3	3	3	3	3	3	3
2	5	4	4	4	4	4	4	4	2
1	5	5	5	5	5	5	5	5	1
	a	b	c	d	e	f	g	h	

Generalizing spatial metrics: P-Spaces

Minkowski distance :

$$MK(x, y) = \left(\sum_{i=1}^n |x_i - y_i|^p \right)^{\frac{1}{p}}$$

- Manhattan distance: $P=1$
- Euclidean distance: $P=2$
- Chebyshev distance: $\lim_{p \rightarrow \infty} MK(x, y)$
- Definition allows for arbitrary, even negative p

Cosine similarity/distance


- Angle between vectors
- Vector Product, normalized for length

$$C(x, y) = 1 - \frac{x \cdot y}{||x|| * ||y||} = 1 - \frac{\sum_{i=1}^n x_i y_i}{\sqrt{\sum_{i=1}^n x_i^2} \sqrt{\sum_{i=1}^n y_i^2}}$$

- Why cosine?:

$$x \cdot y = ||x|| * ||y|| * \cos \theta$$

- Cosine vs Euclidean
 - Cosine ignores length of vectors
 - Useful for weighted data (e.g., frequency counts in documents)



Edit / String Distance

- Measure distance between character sequences by the number of edit operations
- kitten -> sitting
- Levenshtein: Distance 3
 - Replace s for k
 - Replace i for e
 - Insert g at end
- LCS (only insert+remove): Distance 5
 - Remove k
 - Insert s at beginning
 - Remove e
 - Insert i before n
 - Insert g add end
- Could also apply hamming distance here: $A=\{k,e\}$, $B=\{s,i,g\}$

Representing Documents

- Naïve, extreme representations not really helpful
 - Full document (only for identity)
 - Individual symbols/characters (maybe for language detection)
- Split document - what is the right granularity?
 - Individual words
 - Groups of words or characters (n-grams)
- Does the order matter?
- Does the word frequency matter?
- Or is there any other unit of importance?
 - Structure of document?

Common Representation of Texts: Vector Space

- Each feature (word, shingle) of a document is assigned a dimension.
- The number represents the weight is the value in the respective dimension.
 - A value space of 0/1 may just denote **presence** or **absence** of a feature
 - Another common approach is the *number of occurrences* (**term frequency**)
 - Normalization: bring weights in same range for all documents
 - Feature Weighing over all documents (IDF)
- Useful orderings: lexicographical, weights, ...

Example (with lexicographical ordering)

[as, please, possible, soon, yes]

• D1: "yes as soon as possible" D1: [2, 0, 1, 1, 1]

• D2: "as soon as possible please" D2: [2, 1, 1, 1, 0]

- In real datasets:
 - Vectors are very long (many possible features), but also very sparse
 - Skewed frequency: few very common features, many rare features: stop words, "tail clipping"

Finding Similar Items – Exact Solution

First idea: Inverted Indexes

- Key idea: Similarity can be only greater than 0 if there are shared features
- Inverted Index: look up documents containing a term/features
- Documents (aka forward index)

w = [C,D, F]

z = [C,F, G]

y = [A,B, E]

x = [B, E, H]

Inverted index:

A : y

B : y, x

C : w, z

D : w

E : y, x

F : w, z

G : z

H: x

Find candidates:

for every term in source:

get other documents from index

Merge and traverse

Limitations of inverted indexes

- Document length and term skew:
 - Very common terms shared among many documents
 - Very long list for long words
- Many candidates, large set to store and test

Documents

- $w = [C, D, F]$
- $z = [A, B, E, F, G]$
- $y = [A, B, C, D, E]$
- $x = [B, C, D, E, F]$

Candidates:

- $w = \{y, x, z\}$
- $z = \{x, y, w\}$
- $y = \{x, z, w\}$
- $x = \{w, y, z\}$

Index :

A : z, y

B : z, y, x

C : w, y, x

D : w, y, x

E : z, y, x

F : w, z, x

G : z

Working on Prefixes

- Overlap: $O(x,y) = |x \cap y|$
(not a true distance function)

Given a threshold t , the following hold:

1. $J(x,y) \geq t \Leftrightarrow O(x,y) \geq \frac{t}{1+t} * (|x| + |y|)$
2. $J(x,y) \geq t \Rightarrow O(x,y) \geq t * |x|$
3. $J(x,y) \geq t \Rightarrow t * |x| < |y|$

- With 1 and/or 2, we can approximate J using O

Intuition for prefixes:

- for large data sets, we will aim for high thresholds
- high thresholds mean significant overlap
- With high overlap, we can eliminate candidates after seeing few non-matches when comparing (needs order)

Benefits:

- Only test prefix to find candidates
- Only index up to prefix length

Prefix with high similarity ($t=0.9$)

- $w = [\underline{C}, D, F]$
- $z = [\underline{A}, B, E, F, G]$
- $y = [\underline{A}, B, C, D, E]$
- $x = [\underline{B}, C, D, E, F]$
- Tokens in the prefixes are underlined.
- For each record, similarity has to be tested with respect to its candidates.
- candidates :
 - $z = \{y\}$
 - $y = \{z\}$
- We have to check only 1 pair which afterwards does not meet the threshold.

index :

A : z, y

B : x

C : w

Prefix with low similarity ($t=0.5$)

- $w = [\underline{C}, \underline{D}, F]$
- $z = [\underline{A}, \underline{B}, E, F, G]$
- $y = [\underline{A}, \underline{B}, \underline{C}, D, E]$
- $x = [\underline{B}, \underline{C}, \underline{D}, E, F]$

index :

A : z, y

B : z, y, x

C : w, y, x

D : w, x

E : z

- candidates :
 - $w = \{y, x\}$
 - $z = \{x, y\}$
 - $y = \{x, z, w\}$
 - $x = \{w, y, z\}$
- We have to check 5 pairs - one less than in Example 1 without prefix filtering
- w, x and x, y pass the threshold.

Formalizing prefixes

- Given an ordering over the tokens in all lists and $O(x,y) > a$, then the $(|x| - \alpha + 1)$ -prefix of x and the $(|y| - \alpha + 1)$ -prefix of y must share at least one token.
- Since information on both sides is needed, precomputation is not directly possible
- We can again approximate:
 $O(x,y) \geq t * |x|$ (out of 3 and 1)
- Applying the formulas before, we can determine the needed prefix length of a list u as

$$|u| - \lceil t * |u| \rceil + 1$$

Further improvement: positional filtering

- Consider $t = 0.8$ and
- $y = [\underline{A}, \underline{B}, C, D, E]$
- $x = [\underline{B}, \underline{C}, D, E, F]$
- x and y are candidates of each other for a final similarity check, however they will not pass the constraint $O(x, y) \geq 5$ resulting from $J(x, y) \geq 0.8$.
- Taking the position of the common token B into account the maximum possible overlap can be estimated with respect to the unseen tokens in x and y :
- $1 + \min(3, 4) = 4$ therefore x and y cannot pass the similarity check.
- Let $w = x[i]$ be the i -th token in x .
- w partitions x into a left $x_l(w)$ and right partition $x_r(w)$.
- If $O(x, y) \geq \alpha$, then for every token $w \in x \cap y$:
 $O(x_l(w), y_l(w)) + \min(|x_r(w)|, |y_r(w)|) \geq \alpha$.

Approximative Solution

Hashing for Approximation

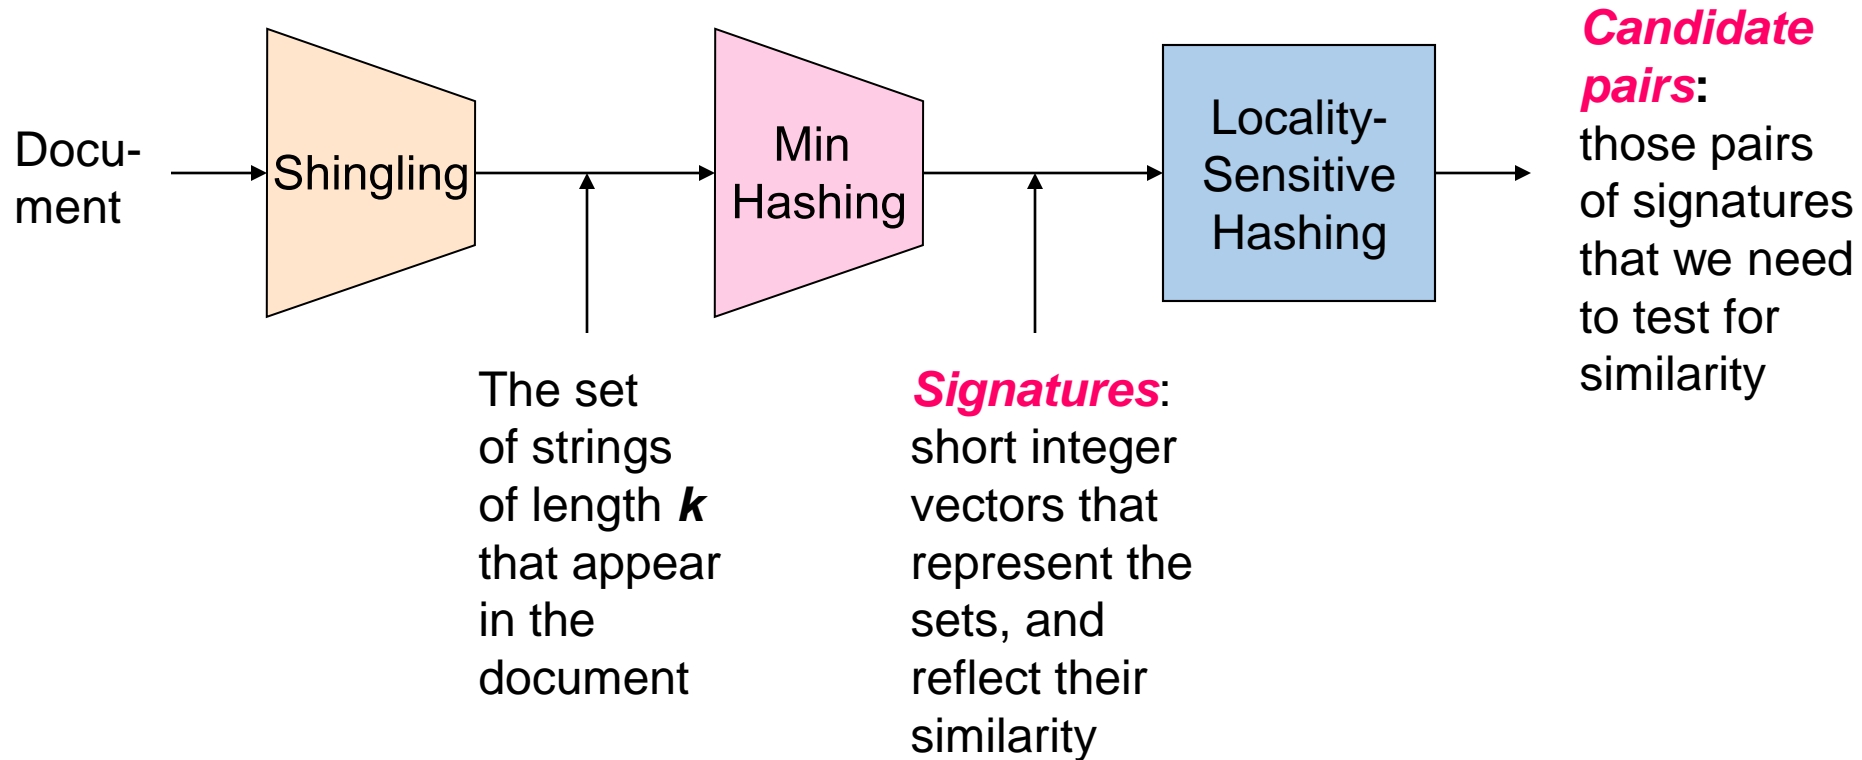
Key Ideas:

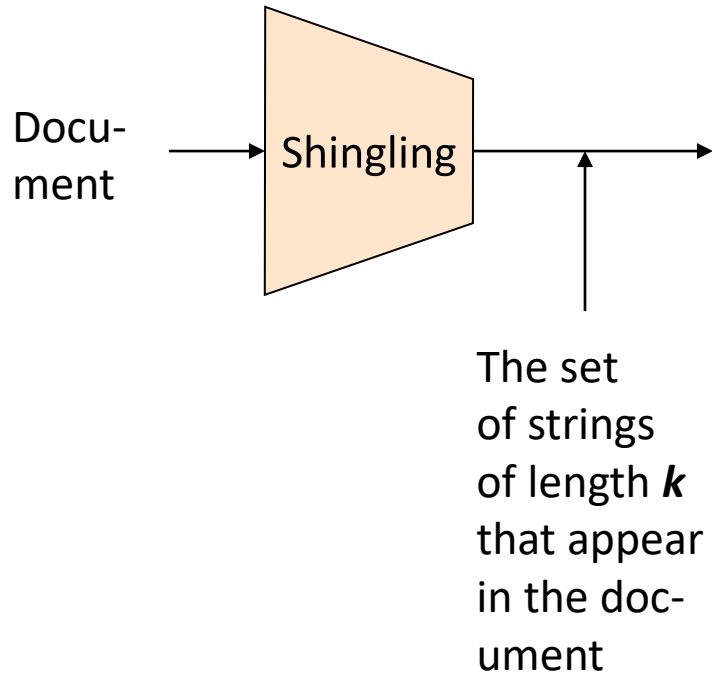
- Apply **multiple** different **hash functions** on the same data to express "more" than equality: **approximate similarity**
- Use **hash functions** that preserve of similarity (dependency on similarity function)
- Use **hashing** at **multiple stages** for **different purposes**

Strategy:

1. **Shingling**: Convert documents to sets
2. **Min-Hashing**: Convert large sets to short signatures, while preserving similarity
3. **Locality-Sensitive Hashing**: Focus on pairs of signatures likely to be from similar documents
 - **Reduce search space to generate candidate pairs**
 - **Careful about false negatives**

The Big Picture





Shingling

Step 1: *Shingling*: Convert documents to sets

Documents as High-Dim Data

- **Step 1: *Shingling*: Convert documents to sets**
- **Simple approaches:**
 - Document = set of words appearing in document
 - Document = set of “important” words
 - Don’t work well for this application. *Why?*
- **Need to account for ordering of words!**
- A different way: ***Shingles!***

Define: Shingles

- A ***k*-shingle** (or ***k*-gram**) is a sequence of k tokens that appears in a document
 - Tokens can be **characters**, **words** or something else, depending on the application
 - Assume tokens = characters for examples
 - Amount of shingles bigger than number of tokens or length of document
- **Example:** $k=2$; document $D_1 = \text{abcaab}$
Set of 2-shingles: $S(D_1) = \{\text{ab}, \text{bc}, \text{ca}\}$
 - **Option:** Shingles as a bag (multiset),
count ab twice: $S'(D_1) = \{\text{ab}, \text{bc}, \text{ca}, \text{ab}\}$

Assessment of Shingles

- **Benefit:** Shingles capture some order of the document
 - Full representation of order within a shingle
 - Overlapping shingles provide indication of overall ordering
 - Reordering the document invalidates only few shingles
- **Core Tuning question: How big to make k?**
 - Too short = most shingles in most documents
 - Too long = missing too many possible candidates

Compressing Shingles


- Motivation:
 - Shingles generate a large space: token space^{length}, e.g., 27^9
 - Shingles consume $O(\text{length})$ bytes
 - String operations are (relatively) inefficient
- To **compress long shingles**, we can **hash** them to (say) 4 bytes
- **Represent a document by the set of hash values of its k -shingles**
 - **Idea:** Two documents could (rarely) appear to have shingles in common, when in fact only the hash-values were shared
 - In practice, the space of 2^{32} is sufficient to cover all relevant shingles
- **Example: $k=2$** ; document $D_1 = \text{abcab}$
Set of 2-shingles: $S(D_1) = \{\text{ab}, \text{bc}, \text{ca}\}$
Hash the singles: $h(D_1) = \{1, 5, 7\}$

Similarity Metric for Shingles

- **Document D_1 is a set of its k -shingles $C_1=S(D_1)$**
- Equivalently, each document is a 0/1 vector in the space of k -shingles
 - Each unique shingle is a dimension
 - Vectors are very sparse
- **A natural similarity measure is the Jaccard similarity**

Assumption:

- **Documents that have lots of shingles in common have similar text, even if the text appears in different order**
- **Caveat:** You must pick k large enough, or most documents will have most shingles
 - $k = 5$ is OK for short documents
 - $k = 10$ is better for long documents



The set
of strings
of length k
that appear
in the doc-
ument


Signatures:
short integer
vectors that
represent the
sets, and
reflect their
similarity

MinHashing

Step 2: *Minhashing*: Convert large sets to short signatures, while preserving similarity

Encoding Sets as Bit Vectors

- Many similarity problems can be formalized as **finding subsets that have significant intersection**
- **Encode sets using 0/1 (bit, boolean) vectors**
 - One dimension per element in the universal set
- Interpret **set intersection as bitwise AND**, and **set union as bitwise OR**
- **Example:** $C_1 = 10111$; $C_2 = 10011$
 - Size of intersection = **3**; size of union = **4**,
 - **Jaccard similarity** (not distance) = **$3/4$**
 - **Distance:** $d(C_1, C_2) = 1 - (\text{Jaccard similarity}) = 1/4$



From Sets to Boolean Matrices

- **Rows** = elements (shingles)
- **Columns** = sets (documents)
 - 1 in row e and column s if and only if e is a member of s
 - Column similarity is the Jaccard similarity of the corresponding sets (rows with value 1)
 - **Typical matrix is sparse!**
- **Each document is a column:**
 - **Example:** $\text{sim}(C_1, C_2) = ?$
 - Size of intersection = 3; size of union = 6, Jaccard similarity (not distance) = $3/6$
 - $d(C_1, C_2) = 1 - (\text{Jaccard similarity}) = 3/6$

	Documents			
Shingles	1	1	1	0
	1	1	0	1
	0	1	0	1
	0	0	0	1
	1	0	0	1
	1	1	1	0
	1	0	1	0

Outline: Finding Similar Columns

- **So far:**
 - Documents → Sets of shingles
 - Represent sets as boolean vectors in a matrix
- **Next goal: Find similar columns while computing small signatures**
 - **Similarity of columns == similarity of signatures**

Outline: Finding Similar Columns

- **Next Goal: Find similar columns, Small signatures**
- **Naïve approach:**
 - **1) Signatures of columns:** small summaries of columns
 - **2) Examine pairs of signatures** to find similar columns
 - **Essential:** Similarities of signatures and columns are related
 - **3) Optional:** Check that columns with similar signatures are really similar
- **Warnings:**
 - Comparing all pairs may take too much time: **Job for LSH**
 - These methods can produce false negatives, and even false positives (if the optional check is not made)

Hashing Columns (Signatures)

- **Key idea:** “hash” each column C to a small **signature** $h(C)$, such that:
 - (1) $h(C)$ is small enough that the signature fits in RAM
 - (2) $\text{sim}(C_1, C_2)$ is the same as the “similarity” of signatures $h(C_1)$ and $h(C_2)$
 - **Goal:** Find a hash function $h(\cdot)$ such that:
 - If $\text{sim}(C_1, C_2)$ is high, then with high prob. $h(C_1) = h(C_2)$
 - If $\text{sim}(C_1, C_2)$ is low, then with high prob. $h(C_1) \neq h(C_2)$
- Hash docs into buckets. Expect that “most” pairs of near duplicate docs hash into the same bucket!

Min-Hashing

- **Goal: Find a hash function $h(\cdot)$ such that:**
 - if $\text{sim}(C_1, C_2)$ is high, then with high prob. $h(C_1) = h(C_2)$
 - if $\text{sim}(C_1, C_2)$ is low, then with high prob. $h(C_1) \neq h(C_2)$
- **Clearly, the hash function depends on the similarity metric:**
 - Not all similarity metrics have a suitable hash function
- **There is a suitable hash function for the Jaccard similarity: It is called Min-Hashing**

Min-Hashing

- Imagine the rows of the boolean matrix permuted under **random permutation** π
- Define a “**hash**” function $h_{\pi}(C)$ = the index of the **first** (in the permuted order π) row in which column C has value **1**:

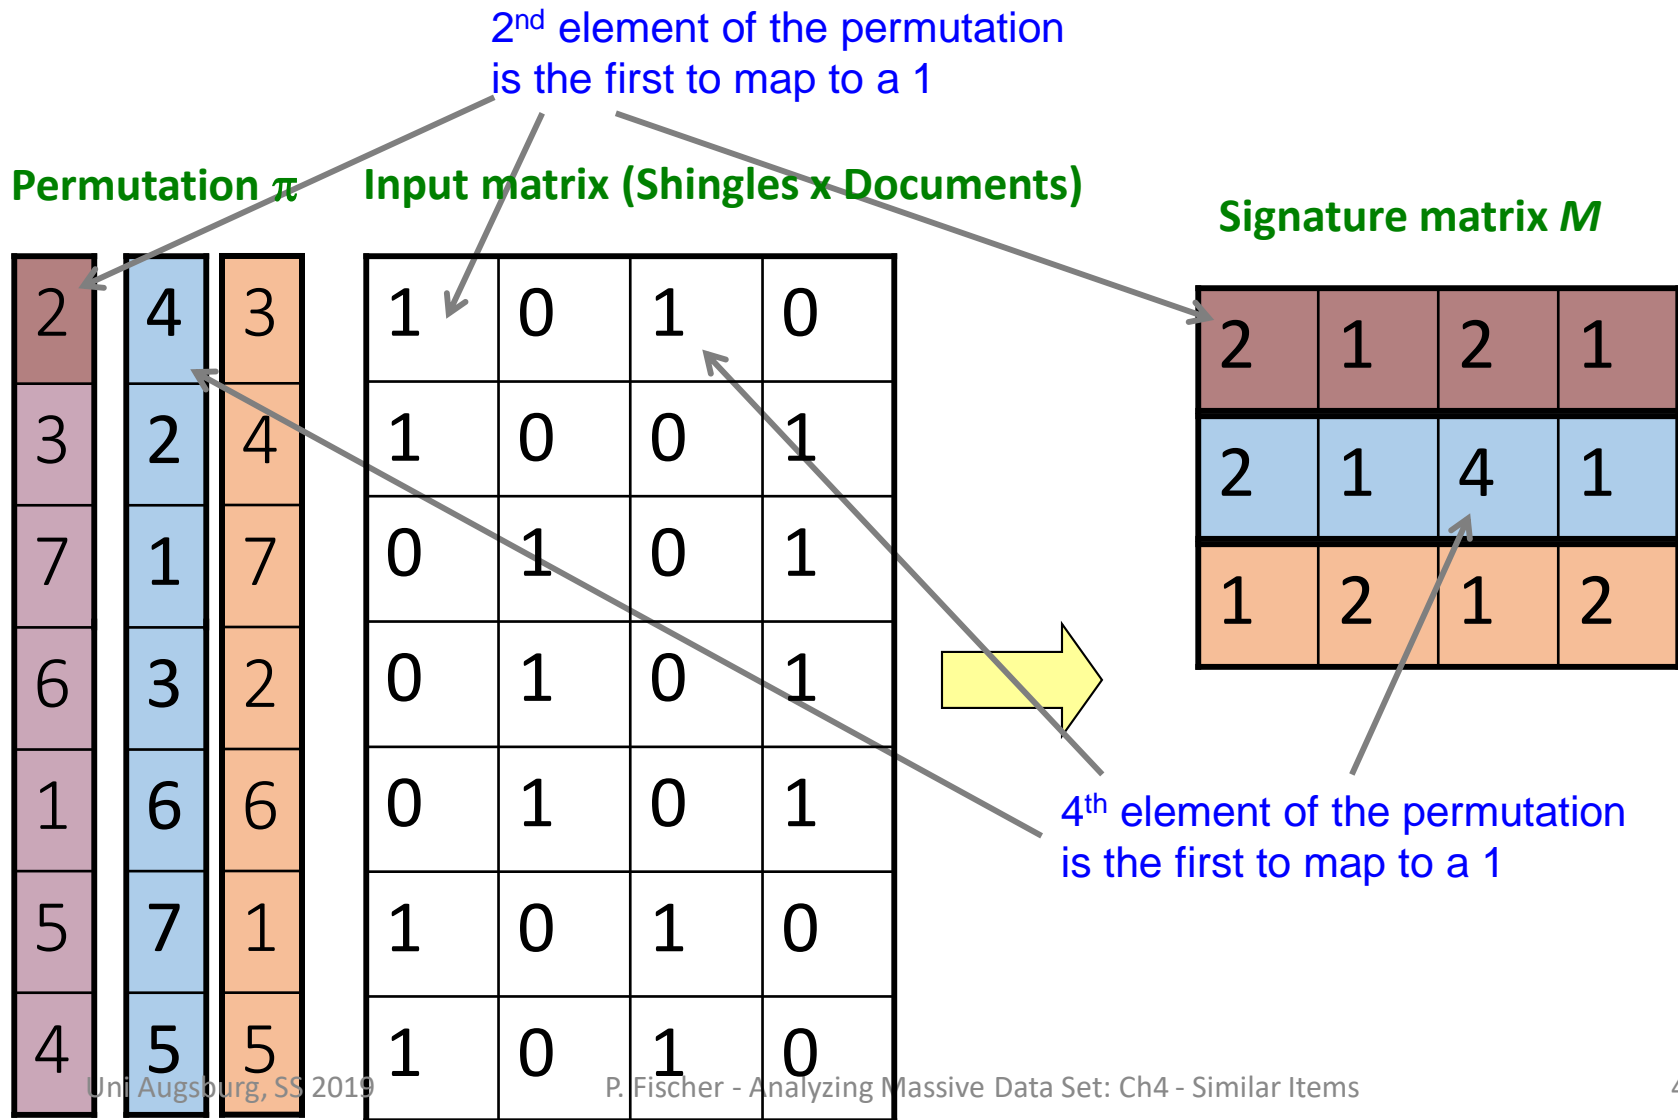
$$h_{\pi}(C) = \min_{\pi} \pi(C)$$

- Use several (e.g., 100) independent hash functions (that is, permutations) to create a signature of a column

Min-Hashing Example

Note: Another (equivalent) way is to store row indexes:

1	5	1	5
2	3	1	3
6	4	6	4



The Min-Hash Property

- Choose a random permutation π
- Claim: $\Pr[h_\pi(C_1) = h_\pi(C_2)] = \text{sim}(C_1, C_2)$
- Why?
 - Let X be a doc (set of shingles), $y \in X$ is a shingle
 - Then: $\Pr[\pi(y) = \min(\pi(X))] = 1/|X|$
 - It is equally likely that any $y \in X$ is mapped to the *min* element
 - Let y be s.t. $\pi(y) = \min(\pi(C_1 \cup C_2))$
 - Then either: $\pi(y) = \min(\pi(C_1))$ if $y \in C_1$, or $\pi(y) = \min(\pi(C_2))$ if $y \in C_2$
 - So the prob. that **both** are true is the prob. $y \in C_1 \cap C_2$
 - $\Pr[\min(\pi(C_1)) = \min(\pi(C_2))] = |C_1 \cap C_2| / |C_1 \cup C_2| = \text{sim}(C_1, C_2)$

0	0
0	0
1	1
0	0
0	1
1	0

One of the two
cols had to have
1 at position y

Min-Hash Signatures

- We know: $\Pr[h_{\pi}(C_1) = h_{\pi}(C_2)] = \text{sim}(C_1, C_2)$
- Now generalize to multiple hash functions
- The **similarity of two signatures** is the fraction of the hash functions in which they agree
- Pick **K=100 random permutations of the rows**
- Think of $\text{sig}(C)$ as a column vector
- $\text{sig}(C)[i]$ = according to the i -th permutation, the index of the first row that has a 1 in column C
$$\text{sig}(C)[i] = \min(\pi_i(C))$$
- **Note:** The sketch (signature) of document C is small **~100 bytes!**
- **We achieved our goal!** We “compressed” long bit vectors into short signatures
- **Real-Life Implementation:** Don't permute, but use random hash functions

Min-Hash Signature Example

Permutation π


2	4	3
3	2	4
7	1	7
6	3	2
1	6	6
5	7	1
4	5	5

Input matrix (Shingles x Documents)

1	0	1	0
1	0	0	1
0	1	0	1
0	1	0	1
0	1	0	1
1	0	1	0
1	0	1	0

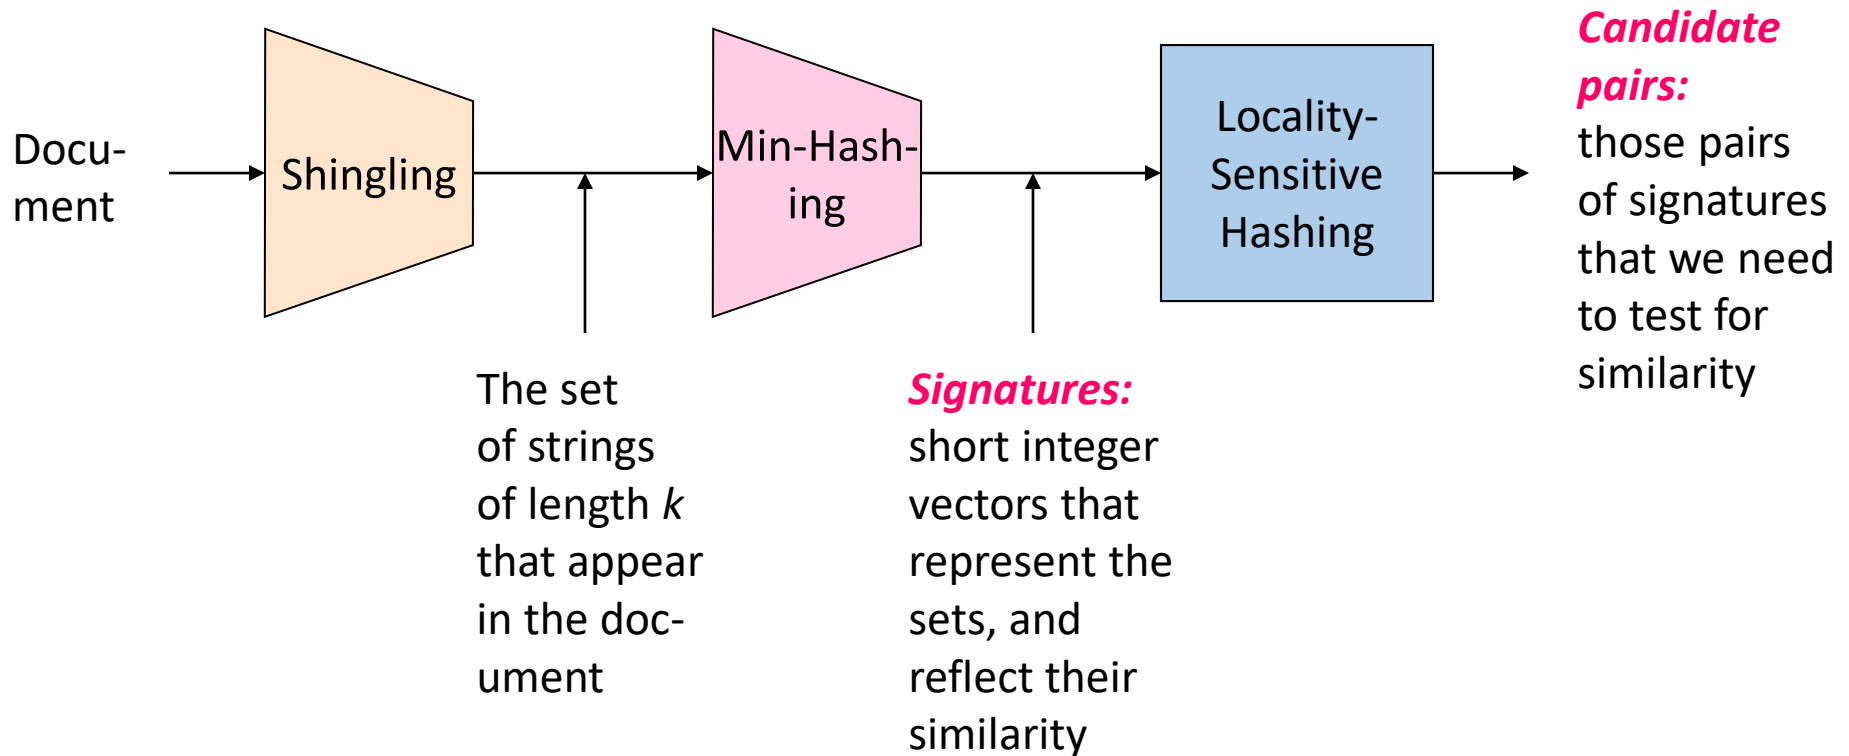
Signature matrix M

2	1	2	1
2	1	4	1
1	2	1	2



Similarities:

	1-3	2-4	1-2	3-4
Col/Col	0.75	0.75	0	0
Sig/Sig	0.67	1.00	0	0



Locality Sensitive Hashing

Step 3: *Locality-Sensitive Hashing:*

Focus on pairs of signatures likely to be from similar documents

2	1	4	1
1	2	1	2
2	1	2	1

LSH: First Cut

- **Goal:** Find documents with Jaccard similarity at least s (for some similarity threshold, e.g., $s=0.8$)
- **LSH – General idea:** Use a function $f(x,y)$ that tells whether x and y is a **candidate pair**: a pair of elements whose similarity must be evaluated
- **For Min-Hash matrices:**
 - Hash columns of **signature matrix M** to many buckets
 - Each pair of documents that hashes into the same bucket is a **candidate pair**
- **Caveat: This approach can generate**
 - **False positives:** Same bucket but too little overlap -> check afterwards
 - **False negatives:** Never same bucket, but enough overlap

Candidates from Min-Hash

2	1	4	1
1	2	1	2
2	1	2	1

- Pick a similarity threshold s ($0 < s < 1$)
- Columns x and y of M are a **candidate pair** if their signatures agree on at least fraction s of their rows:
 $M(i, x) = M(i, y)$ for at least frac. s values of i
 - We expect documents x and y to have the same (Jaccard) similarity as their signatures


LSH for Min-Hash

2	1	4	1
1	2	1	2
2	1	2	1

- **Big idea:** Hash columns of signature matrix M several times
- Arrange that (only) **similar columns** are likely to **hash to the same bucket**, with high probability
- **Candidate pairs are those that hash to the same bucket**

2	1	4	1
1	2	1	2
2	1	2	1

Partition M into b Bands




Signature matrix M

Partition M into Bands

- Divide matrix M into b bands of r rows
- For each band, hash its portion of each column to a hash table with k buckets
 - Make k as large as possible
- **Candidate** column pairs are those that hash to the same bucket for ≥ 1 band
- Tune b and r to catch most similar pairs, but few non-similar pairs

Hashing Bands



2	1	4	1
1	2	1	2
2	1	2	1

Example of Bands

Assume the following case:

- Suppose 100,000 columns of \mathbf{M} (100k docs)
- Signatures of 100 integers (rows)
- Therefore, signatures take 40Mb
- Choose $b = 20$ bands of $r = 5$ integers/band
- **Goal:** Find pairs of documents that are at least $s = 0.8$ similar

2	1	4	1
1	2	1	2
2	1	2	1

C_1, C_2 are 80% Similar

- Find pairs of $\geq s=0.8$ similarity, set $b=20, r=5$
- **Assume:** $\text{sim}(C_1, C_2) = 0.8$
 - Since $\text{sim}(C_1, C_2) \geq s$, we want C_1, C_2 to be a **candidate pair**: We want them to hash to at **least 1 common bucket** (at least one band is identical)
- **Probability C_1, C_2 identical in one particular band:** $(0.8)^5 = 0.328$
- Probability C_1, C_2 are **not** similar in all of the 20 bands: $(1-0.328)^{20} = 0.00035$
 - i.e., about 1/3000th of the 80%-similar column pairs are **false negatives** (we miss them)
 - **We would find 99.965% pairs of truly similar documents**

2	1	4	1
1	2	1	2
2	1	2	1

C_1, C_2 are 30% Similar

- Find pairs of $\geq s=0.8$ similarity, set $b=20, r=5$
- **Assume:** $\text{sim}(C_1, C_2) = 0.3$
 - Since $\text{sim}(C_1, C_2) < s$ we want C_1, C_2 to hash to **NO common buckets** (all bands should be different)
- **Probability C_1, C_2 identical in one particular band:** $(0.3)^5 = 0.00243$
- Probability C_1, C_2 identical in at least 1 of 20 bands: $1 - (1 - 0.00243)^{20} = 0.0474$
 - In other words, approximately 4.74% pairs of docs with similarity 0.3% end up becoming **candidate pairs**
 - They are **false positives** since we will have to examine them (they are candidate pairs) but then it will turn out their similarity is below threshold s

2	1	4	1
1	2	1	2
2	1	2	1

LSH Involves a Tradeoff


- **Pick:**

- The number of Min-Hashes (rows of M)
- The number of bands b , and
- The number of rows r per band


to balance false positives/negatives

- **Example:** If we had only 15 bands of 5 rows, the number of false positives would go down, but the number of false negatives would go up

Analysis of LSH – What We Want




What 1 Band of 1 Row Gives You



b bands, r rows/band

- Columns C_1 and C_2 have similarity t
- Pick any band (r rows)
 - Prob. that all rows in band equal = t^r
 - Prob. that some row in band unequal = $1 - t^r$
- Prob. that no band identical = $(1 - t^r)^b$
- Prob. that at least 1 band identical = $1 - (1 - t^r)^b$

What b Bands of r Rows Gives You



Example: $b = 20$; $r = 5$


- Similarity threshold s
- Prob. that at least 1 band is identical:

s	$1-(1-s^r)^b$
.2	.006
.3	.047
.4	.186
.5	.470
.6	.802
.7	.975
.8	.9996

Picking r and b : The S-curve

- **Picking r and b to get the best S-curve**

- 50 hash-functions ($r=5$, $b=10$)



Blue area: False Negative rate
Yellow area: False Positive rate

LSH Summary

- Tune M , b , r to get almost all pairs with similar signatures, but eliminate most pairs that do not have similar signatures
- Check in main memory that **candidate pairs** really do have **similar signatures**
- **Optional:** In another pass through data, check that the remaining candidate pairs really represent similar documents

Summary: 3 Steps

- **Shingling:** Convert documents to sets
 - We used hashing to assign each shingle an ID
- **Min-Hashing:** Convert large sets to short signatures, while preserving similarity
 - We used **similarity preserving hashing** to generate signatures with property $\Pr[h_{\pi}(C_1) = h_{\pi}(C_2)] = \text{sim}(C_1, C_2)$
 - We used hashing to get around generating random permutations
- **Locality-Sensitive Hashing:** Focus on pairs of signatures likely to be from similar documents
 - We used hashing to find **candidate pairs** of similarity $\geq s$