



# Decision Tree Learning

A method for approximating discrete-valued functions which is

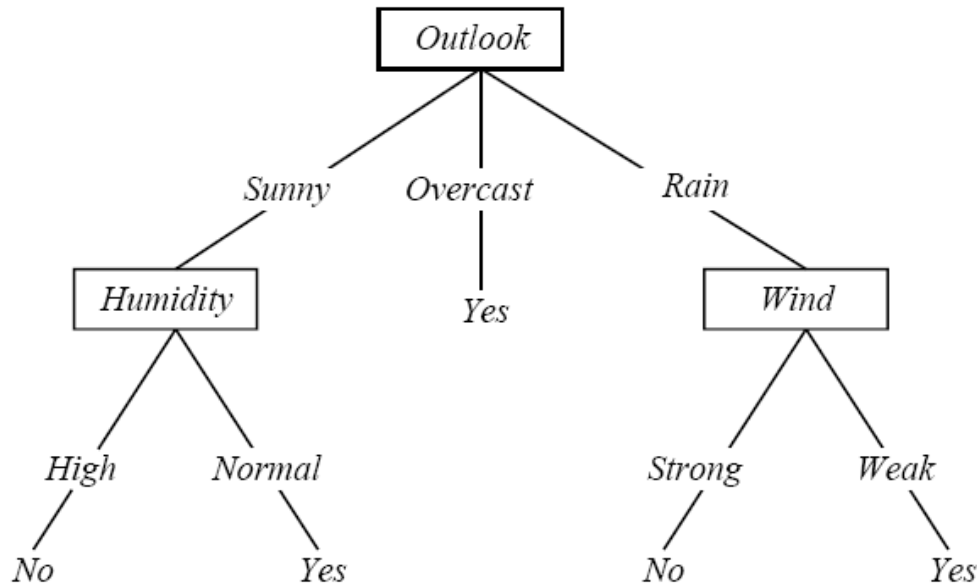
1. robust to noisy data
2. capable of learning disjunctive expressions
  - ➔ Searches a completely expressive hypothesis space
3. easy to understand by humans (especially its learned results!)

Inductive Bias: “Prefer small trees over large trees”

- Decision tree representation
- ID3 learning algorithm
- Entropy, Information gain
- Overfitting

| Day | Outlook  | Temperature | Humidity | Wind   | PlayTennis |
|-----|----------|-------------|----------|--------|------------|
| 1   | Sunny    | Hot         | High     | Weak   | No         |
| 2   | Sunny    | Hot         | High     | Strong | No         |
| 3   | Overcast | Hot         | High     | Weak   | Yes        |
| 4   | Rain     | Mild        | High     | Weak   | Yes        |
| 5   | Rain     | Cool        | Normal   | Weak   | Yes        |
| 6   | Rain     | Cool        | Normal   | Strong | No         |
| 7   | Overcast | Cool        | Normal   | Strong | Yes        |
| 8   | Sunny    | Mild        | High     | Weak   | No         |
| 9   | Sunny    | Cool        | Normal   | Weak   | Yes        |
| 10  | Rain     | Mild        | Normal   | Weak   | Yes        |
| 11  | Sunny    | Mild        | Normal   | Strong | Yes        |
| 12  | Overcast | Mild        | High     | Strong | Yes        |
| 13  | Overcast | Hot         | Normal   | Weak   | Yes        |
| 14  | Rain     | Mild        | High     | Strong | No         |

# Decision Tree for *PlayTennis*



How to convert tree into disjunctions of conjunctions of constraints on attribute values of instances?

→ Convert into set of if-then rules

→ Each path from the root node to a leaf corresponds to a conjunctions of attribute tests

→ Tree is a disjunctions of conjunctions of constraints on attribute values

1. How Would you classify <outlook=Sunny, Temperature=Hot, Humidity=High, Wind=Strong>?

# Typical Datamining Task

Given:

- 9714 patient records each describing a pregnancy and birth
- Each patient record contains 215 features

From lecture 1-1

Learn to predict:

- Classes of future patients at high risk for Emergency Cesarean Section

---

Data:

| <i>Patient103</i> time=1   | <i>Patient103</i> time=2   | ... | <i>Patient103</i> time=n        |
|----------------------------|----------------------------|-----|---------------------------------|
| Age: 23                    | Age: 23                    |     | Age: 23                         |
| FirstPregnancy: no         | FirstPregnancy: no         |     | FirstPregnancy: no              |
| Anemia: no                 | Anemia: no                 |     | Anemia: no                      |
| Diabetes: no               | Diabetes: YES              |     | Diabetes: no                    |
| PreviousPrematureBirth: no | PreviousPrematureBirth: no |     | PreviousPrematureBirth: no      |
| Ultrasound: ?              | Ultrasound: abnormal       |     | Ultrasound: ?                   |
| Elective C-Section: ?      | Elective C-Section: no     |     | Elective C-Section: no          |
| Emergency C-Section: ?     | Emergency C-Section: ?     |     | <b>Emergency C-Section: Yes</b> |
| ...                        | ...                        |     | ...                             |

Learned from medical records of 1000 women, neg. examples are C-sections:

[833+,167-] .83+ .17-

- Fetal\_Presentation = 1: [822+,116-] .88+ .12-
  - Previous\_Csection = 0: [767+,81-] .90+ .10-
    - Primiparous = 0: [399+,13-] .97+ .03-
    - Primiparous = 1: [368+,68-] .84+ .16-
      - Fetal\_Distress = 0: [334+,47-] .88+ .12-
        - » Birth\_Weight < 3349: [201+,10.6-] .95+ .05-
        - » Birth\_Weight >= 3349: [133+,36.4-] .78+ .22-
      - Fetal\_Distress = 1: [34+,21-] .62+ .38-
    - Previous\_Csection = 1: [55+,35-] .61+ .39-
  - Fetal\_Presentation = 2: [3+,29-] .11+ .89-
  - Fetal\_Presentation = 3: [8+,22-] .27+ .73-

- Decision tree representation:
  - Each internal node tests an attribute
  - Each branch corresponds to an attribute value
  - Each leaf node assigns a classification
- How would we represent:
  - AND, OR, XOR
  - $(A \wedge B) \vee (C \wedge \neg D \wedge E)$
  - $M$  of  $N$





$$A \wedge B$$



$A \vee B$



$A \text{ XOR } B$





$$(A \wedge B) \vee (C \wedge \neg D \wedge E)$$

- Instances described by attribute values
- Target function is discrete valued
- Disjunctive hypotheses may be required
- Possibly noisy training data
- Training data may contain missing attribute values
- Examples:
  - Equipment or medical diagnosis
  - Credit risk analysis
  - Modeling calendar scheduling preferences


- Main loop:
  - $A \leftarrow$  the “best” decision attribute for next *node*
  - Assign  $A$  as decision attribute for *node*
  - For each value of  $A$ , create new descendant of *node*
  - Sort training examples to leaf nodes
  - If training examples perfectly classified, then STOP, else iterate over new leaf nodes

- Which attribute is best?



# (Shannon) Entropy (1)

- $S$  is a sample of training examples
- $p_+$  is the fraction of positive examples in  $S$
- $p_-$  is the fraction of negative examples in  $S$
- Entropy measures the impurity of  $S$



$$\text{Entropy}(S) \equiv -p_+ \log_2 p_+ - p_- \log_2 p_-$$



## (Shannon) Entropy (2)

*Entropy*( $S$ ) = expected number of bits needed to encode class + or - of randomly drawn member of  $S$  (under the optimal, shortest-length code)

Why?

*Information theory*: Optimal length code assigns  $-\log_2 p$  bits to message having probability  $p$ .

So, expected number of bits to encode + or - of random member of  $S$ :

$$p_+(-\log_2 p_+) + p_-(-\log_2 p_-)$$

$$\text{Entropy}(S) = -p_+ \log_2 p_+ - p_- \log_2 p_-$$

$$\textit{Entropy}(S) \equiv \sum_{i=1}^c -p_i \log_2 p_i$$

- $p_i$  = proportion of  $S$  belonging to class  $i$
- Maximum value =  $\log_2 c$
- Measures the expected encoding length measured in *bits*

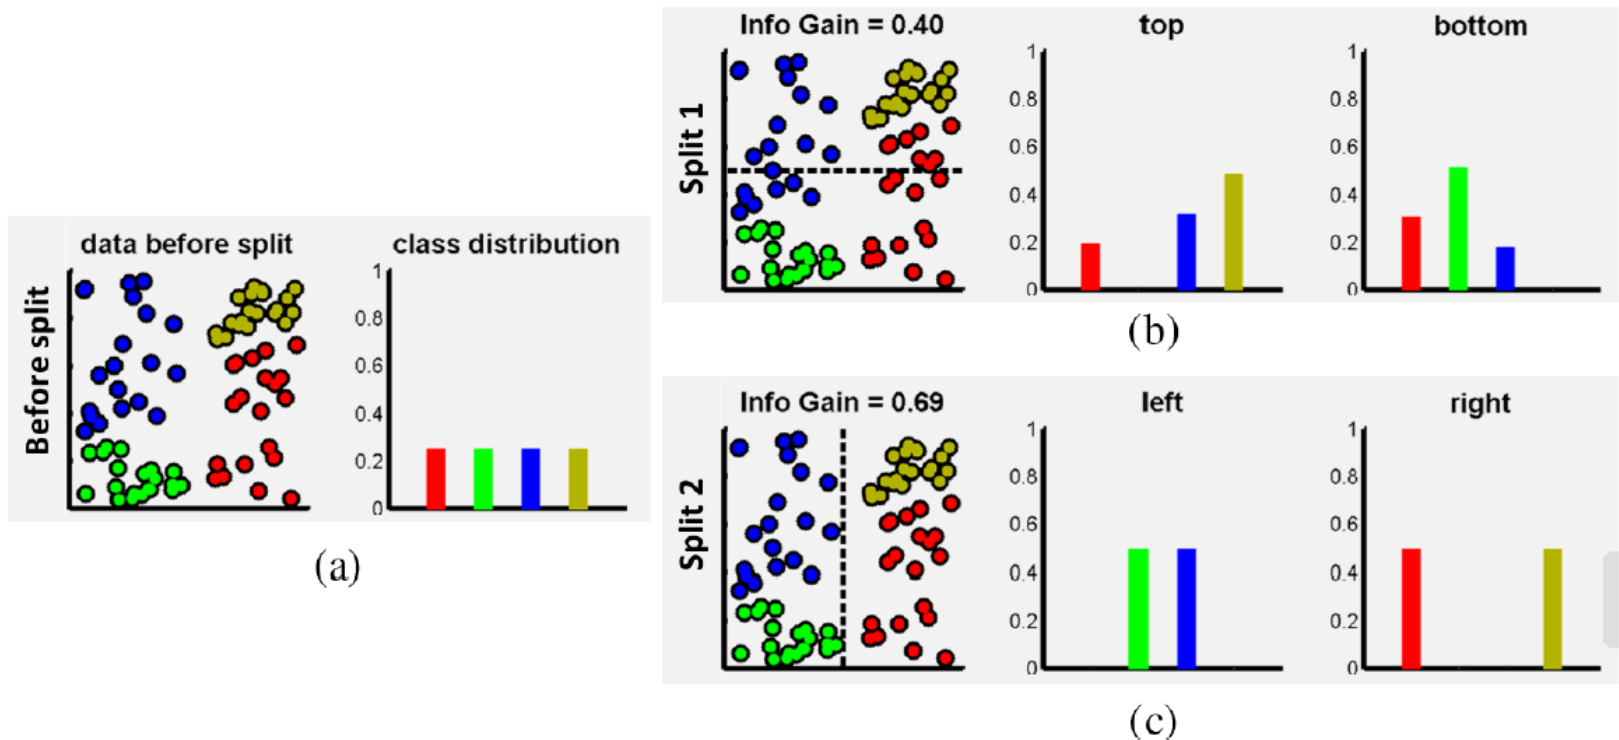




Fig. 2.5 Information gain for discrete, non-parametric distributions. (a) Dataset  $\mathcal{S}$  before a split. (b) After a horizontal split. (c) After a vertical split. In this example the vertical split produces purer class distributions in the child nodes. Classes are colour coded.

- $Gain(S, A)$  = expected reduction in entropy due to sorting on  $A$

Original entropy

$$Gain(S, A) \equiv Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

Entropy after partitioning




→ Let's compute all Entropy and Gain values


| Day | Outlook  | Temperature | Humidity | Wind   | PlayTennis |
|-----|----------|-------------|----------|--------|------------|
| 1   | Sunny    | Hot         | High     | Weak   | No         |
| 2   | Sunny    | Hot         | High     | Strong | No         |
| 3   | Overcast | Hot         | High     | Weak   | Yes        |
| 4   | Rain     | Mild        | High     | Weak   | Yes        |
| 5   | Rain     | Cool        | Normal   | Weak   | Yes        |
| 6   | Rain     | Cool        | Normal   | Strong | No         |
| 7   | Overcast | Cool        | Normal   | Strong | Yes        |
| 8   | Sunny    | Mild        | High     | Weak   | No         |
| 9   | Sunny    | Cool        | Normal   | Weak   | Yes        |
| 10  | Rain     | Mild        | Normal   | Weak   | Yes        |
| 11  | Sunny    | Mild        | Normal   | Strong | Yes        |
| 12  | Overcast | Mild        | High     | Strong | Yes        |
| 13  | Overcast | Hot         | Normal   | Weak   | Yes        |
| 14  | Rain     | Mild        | High     | Strong | No         |

# Selecting the Next Attribute

- Which attribute is the best classifier?

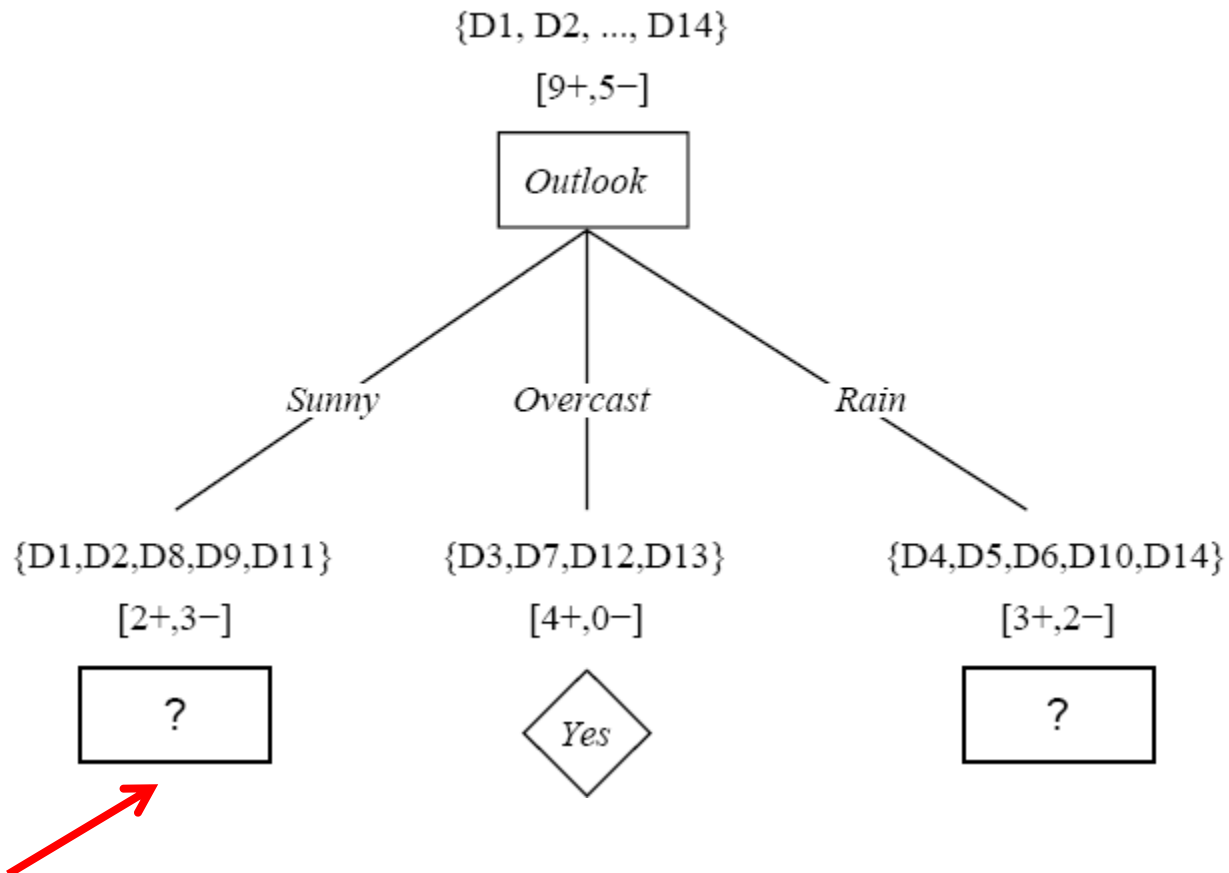


$$\begin{aligned}
 \text{Gain}(S, \text{Humidity}) &= .940 - (7/14).985 - (7/14).592 \\
 &= .151
 \end{aligned}$$



$$\begin{aligned}
 \text{Gain}(S, \text{Wind}) &= .940 - (8/14).811 - (6/14)1.0 \\
 &= .048
 \end{aligned}$$

# Selecting the Next Attribute



- Which attribute should be tested here?

- $S_{\text{sunny}} = \{D1, D2, D8, D9, D11\}$

$$\text{Gain}(S_{\text{sunny}}, \text{Humidity}) = .970 - (3/5) 0.0 - (2/5) 0.0 = .970$$


$$\text{Gain}(S_{\text{sunny}}, \text{Temperature}) = .970 - (2/5) 0.0 - (2/5) 1.0 - (1/5) 0.0 = .570$$

$$\text{Gain}(S_{\text{sunny}}, \text{Wind}) = .970 - (2/5) 1.0 - (3/5) .918 = .019$$

- Humidity has highest information gain



# Hypothesis Space Search by ID3 (1)



- Hypothesis space is complete!
  - Target function surely is in there...
- Outputs a single hypothesis (which one?)
  - Can't play 20 questions...
- No back tracking
  - Local minima...
- Statistically-based search choices
  - Robust to noisy data...
- Inductive bias:
  - Approximately: „prefer shortest tree“

Note  $H$  is the power set of instances  $X$

- Unbiased?

Not really...

- Preference for short trees, and for those with high information gain attributes near the root
- Bias is a *preference* for some hypotheses, rather than a *restriction* of hypothesis space  $H$
- Occam's razor: prefer the shortest hypothesis that fits the data

# Occam's Razor: Why Prefer Short Hypotheses?

- In favor:
  - There are fewer short hypotheses than long hypotheses
  - A short hypothesis that fits the data is unlikely to be coincidence
  - A long hypothesis that fits the data might be coincidence
- Opposed:
  - There are many ways to define small sets of hypotheses  
e.g., all trees with a prime number of nodes that use attributes beginning with "Z"
  - What is short depends on language

# Issues with Decision Trees:

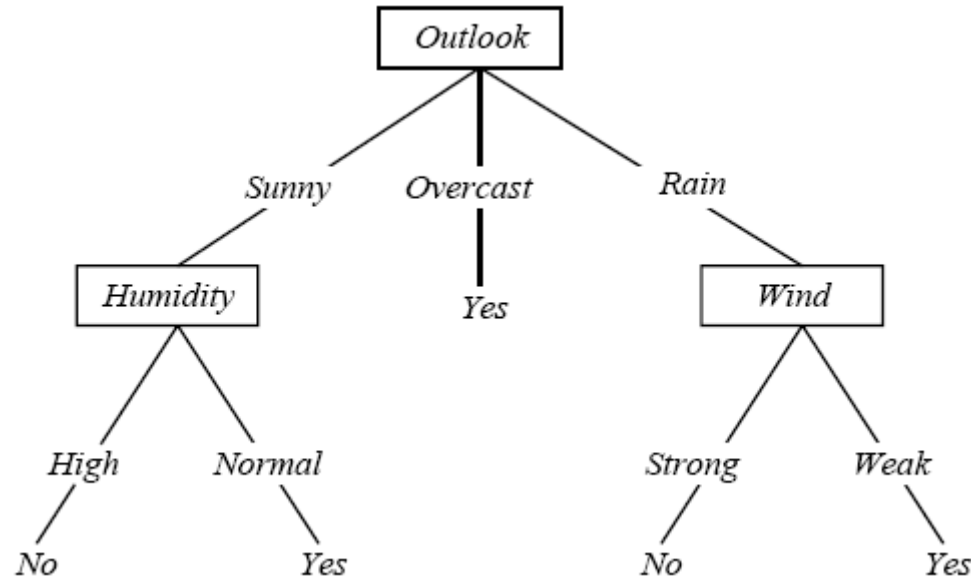
## 1. Overfitting

Consider adding noisy training example #15:

(Outlook, Temperature, Humidity, Wind)

*Sunny, Hot, Normal, Strong, PlayTennis=No*

What is the effect on earlier tree?




- Consider error of hypothesis  $h$  over
  - Training data:  $error_{train}(h)$
  - Entire distribution of data:  $error_D(h)$
- Hypothesis  $h \in H$  overfits training data if there is an alternative hypothesis  $h' \in H$  such that

$$error_{train}(h) < error_{train}(h')$$

and

$$error_D(h) > error_D(h').$$

# Overfitting in Decision Tree Learning




- How can we avoid overfitting?
  - stop growing when data split is not statistically significant
  - grow full tree, then post-prune
- How to select „best“ tree:
  - Measure performance over training data
  - Measure performance over separate validation data set
  - Minimum Description Length: minimize  $size(tree) + size(misclassifications(tree))$

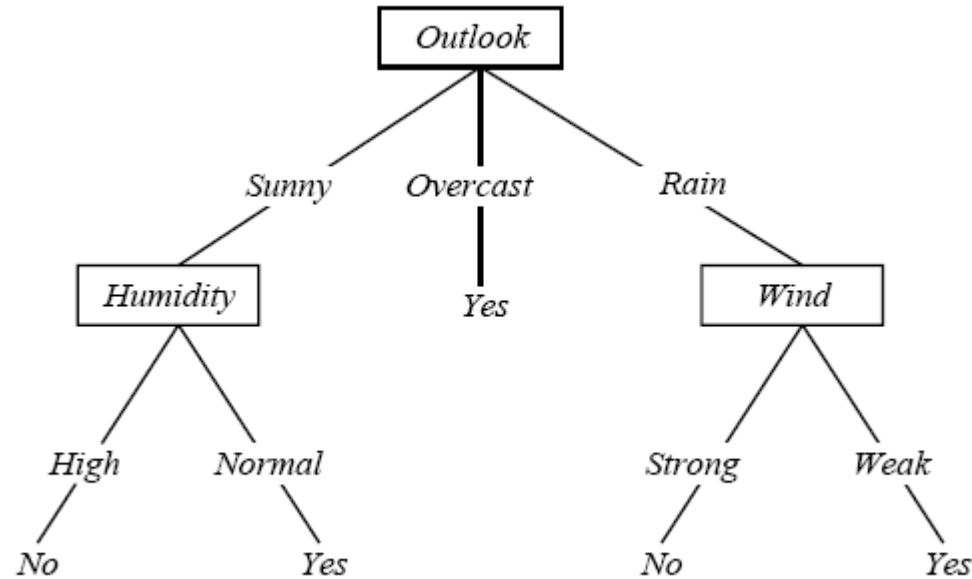


- Split data into *training* and *validation* set
- Do until further pruning is harmful:
  1. Evaluate impact on *validation* set of pruning each possible node (plus those below it)
  2. Greedily remove the one that most improves *validation* set accuracy

# Effect of Reduced Error-Pruning



# Converting a Tree to Rules



- IF (*Outlook=Sunny*) and (*Humidity=High*)  
THEN *PlayTennis=No*
- IF (*Outlook=Sunny*) and (*Humidity=Normal*)  
THEN *PlayTennis=Yes*
- ...

## Issues with Decision Trees: 2. Continuous Valued Attributes

- Create a discrete attribute to test continuous
  - Temperature = 82.5
  - (Temperature > 54) = true ELSE false

|                     |    |    |     |     |     |    |
|---------------------|----|----|-----|-----|-----|----|
| <i>Temperature:</i> | 40 | 48 | 60  | 72  | 80  | 90 |
| <i>PlayTennis:</i>  | No | No | Yes | Yes | Yes | No |

# Issues with Decision Trees:

## 3. Attributes With Many Values

- Problem: Gain measure has a natural bias towards attributes with many values
  - if attribute has many values, *Gain* will select it
  - Imagine using Date = Jun 3 1996 as attribute
- Use alternative method for selecting attributes
- One approach: use GainRatio instead:

$$\text{GainRatio}(S, A) \equiv \frac{\text{Gain}(S, A)}{\text{SplitInformation}(S, A)}$$

$$\text{SplitInformation}(S, A) \equiv - \sum_{i=1}^c \frac{|S_i|}{|S|} \log_2 \frac{|S_i|}{|S|}$$

where  $S_i$  is subset of  $S$  for which  $A$  has value  $v_i$

$\log_2 c$

Discourages the selection of attributes with many uniformly distributed values

E.g. two-class problem in 2D ( $\mathbf{x} \in \mathbb{R}^2$ ):

$$f(\mathbf{x}) = I_{x_2 > x_1}$$

This linear, but not axis aligned decision function needs to be approximated by **many** nodes. It cannot be efficiently learned. It also needs many training examples in order to be approximately learned.

What if some examples have missing values of  $A$ ?

Use training example anyway, sort through tree

1. If node  $n$  tests  $A$ , assign most common value of  $A$  among other examples sorted to node  $n$
2. Assign most common value of  $A$  among other examples with same target value
3. Assign probability  $p_i$  to each possible value  $v_i$  of  $A$ 
  - assign fraction  $p_i$  of example to each descendant in tree

Classify new examples in same fashion.

Model  $n$ -dimensional data  $\{(x, y)\}$ ,  $x \in \mathbb{R}^n$  by  $n$ -variate Gaussian:

$$G = ..$$

In the continuous case we define the differential entropy

$$H(S) = - \int_y p(y) \log(p(y)) dy$$

which becomes for the multi-variate Gaussian

$$H(S) = \frac{1}{2} \log((2\pi e)^n) \cdot \det(\Sigma)$$



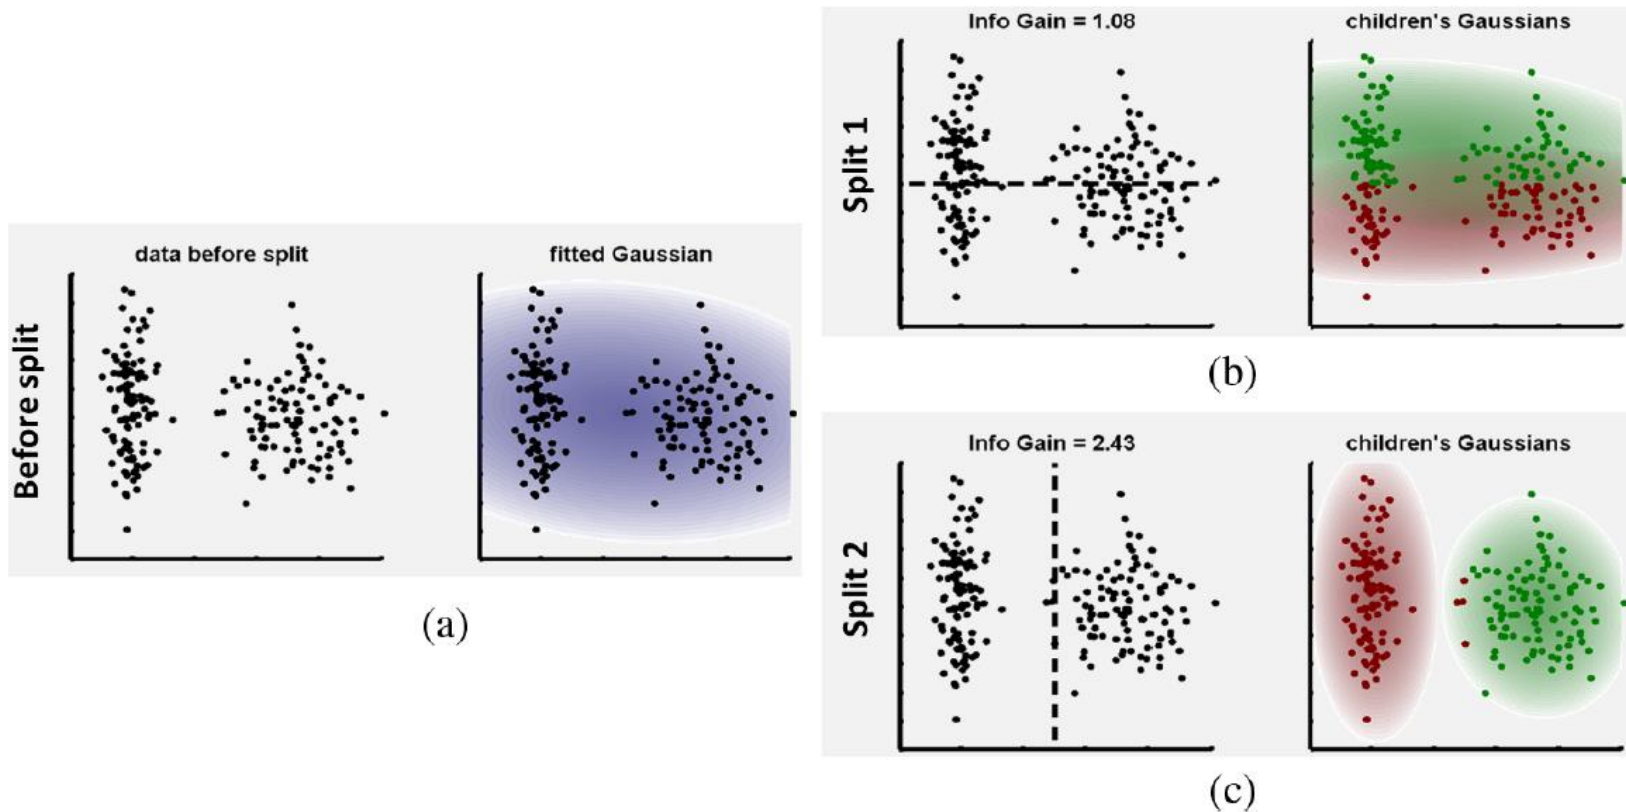


Fig. 2.6 Information gain for continuous, parametric densities. (a) Dataset  $\mathcal{S}$  before a split. (b) After a horizontal split. (c) After a vertical split. A vertical split produces better separation and a correspondingly higher information gain.