



# Deep Learning

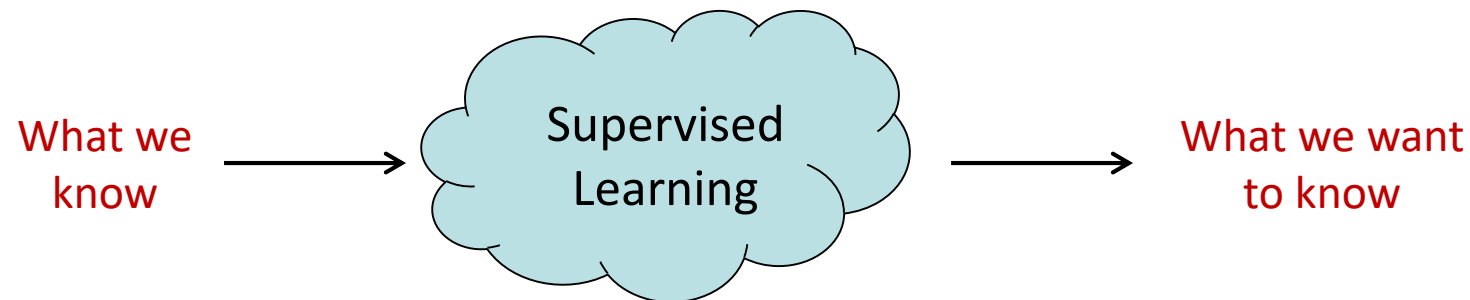
## Introduction to Gradient Descent

Tuesday 29<sup>th</sup> November

Dr. Nicholas Cummins

- **Supervised Learning**


- Transforming one dataset into another
- Taking what you know as input and transforming it into what you want to know at the output
- Input: observable, recordable, and knowable data
- Output: data for logical analysis



- **Supervised parametric learning**
  - A learning model that summarizes data with a set of parameters of fixed size
    - Independent of the number of training examples
  - Such algorithms involve two steps:
    1. Select a form for the function
    2. Learn the coefficients for the function from the training data
- Example simple linear regression



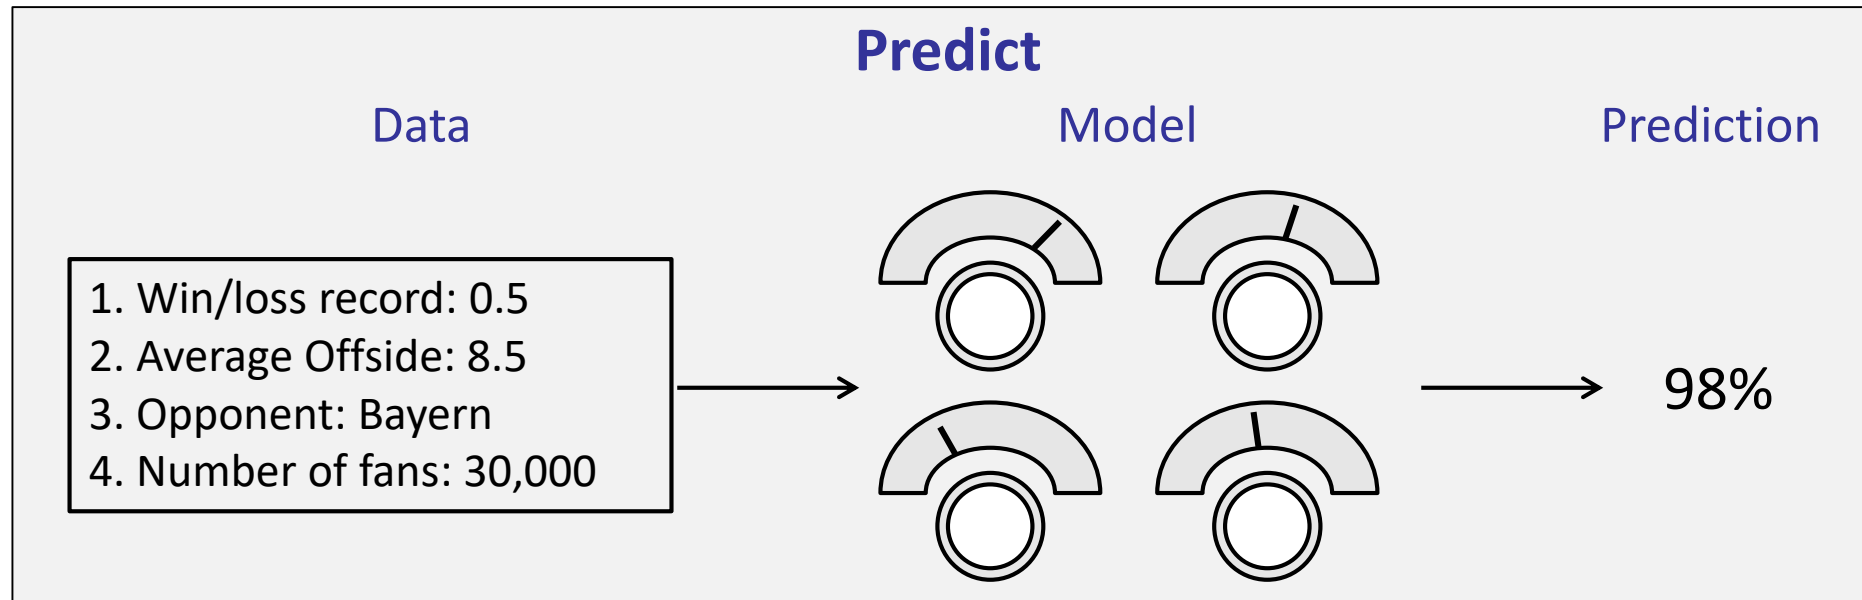
$$b_0 + b_1x_1 + b_2x_2 = y$$

- **Supervised parametric learning analogy:**
  - Machine with a fixed number of knobs
  - Position of knobs indicates how to process the data
  - Processing transforms input data into an output prediction
  - Learning is accomplished by tuning the knobs



- Key steps in supervised parametric learning
  - Step 1: Predict
    - Gather data, send through machine, make a prediction
  - Step 2: Compare with truth
    - Compare the prediction with the actual score


**Prediction: 98% > Truth: 0%**
  - Step 3: Learn the pattern
    - Adjust the knobs to make a more accurate prediction
      - Considers the input data, and how much the models prediction missed by
      - Each knob represents the prediction's sensitivity to the different types of input data



# Prediction with a single network

- **Simple predictions**

- One input data point, one output prediction
- Build a network with one single knob (*the weight*), to learn a mapping to one single output




# Prediction with a single network

- **Key steps**

- 1) Define network
- 2) Feed input value into network
- 3) Multiply input value by weight
- 4) Output Prediction

```
weight = 0.1  
  
def neural_network(input, weight):  
    prediction = input * weight  
    return prediction  
  
number_of_offsides = [8.5, 9.5, 10, 9]  
input = number_of_offsides[0]  
pred = neural_network(input, weight)  
print(pred)
```





- **What is a prediction?**
  - The output of the network, *given the current input data*
- **Is this prediction always right?**
  - Of course not, neural networks make mistakes
  - During training, the network learns from these mistakes
    - Prediction is too high, adjust the weights lower
    - Prediction is too low, adjust the weights higher
- **How does the network learn?**
  - Trail and error: making predictions and learn from them

## How do we set the weights?

- **Compare**


- Evaluate how well the network performed
- Measure of how much a prediction 'missed' by
- Mean Squared Error (MSE) metric

- **Learn**

- Adjusting each weight to reduce the error
- Gradient Descent Algorithm

- **Measuring predictive performance**


- Calculate error by squaring the difference between the networks prediction and its goal
- Squaring the error forces it to be positive



- **Squaring the error**

- Big errors become bigger
- Small errors become smaller
- These effects help the network learn
  - Pays more attention to the big errors
  - Pays less attention to smaller ones
- Mean Square Error Equation:

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$



- **Why measure error?**

- Aim of network training is to make correct predictions
- This can be achieved two ways:

- Adjust weights such that prediction equal target

$$\text{pred} = \text{goal\_pred}$$

- Adjust weights such that error equals zero

$$(\text{pred} - \text{goal\_pred}) = 0$$

- Both essentially say the same thing
- Tuning weights to predict the target is actually a more complicated task than tuning weights to set error to zero
- Therefore tune network such that `Error == 0`


- **Why positive errors?**

- Large networks can have millions of connections
- Therefore millions of `pred`  $\leftrightarrow$  `goal_pred` pairs
- In these circumstances we need take the *average* error down to zero
  - This presents a problem if the error can be positive and negative
  - Consider network predicting two data points with errors:
    - $(y_1 - \hat{y}_1) = 1000$  and  $(y_2 - \hat{y}_2) = -1000$
    - The average error is zero!
- We want positive errors so they don't cancel each other out when they are average

# Updating the weights

- **Measuring error and finding the direction and amount!**
  - Represents how you want to change weight
    - Pure Error
    - Scaling, negative reversal and stopping
  - This can be achieved in a single line of code!

## Basic Gradient Descent Algorithm



```
weight = 0.5
goal_pred = 0.8
input = 0.5
for iteration in range(20):
    pred = input * weight
    error = (pred - goal_pred) ** 2
    direction_and_amount = (pred - goal_pred) * input
    weight = weight - direction_and_amount
```

- What is direction and amount?

```
direction_and_amount = (pred - goal_pred) * input
```

- How much to change weight by to reduce the error
- Two core parts:
  1. The Pure Error: **pred - goal\_pred**
  2. Multiplication by the input: **\*input**
    - Modifying the pure error so it's ready to update the weight.
    - Performs scaling, negative reversal and stopping



- **What is Pure Error?**

$$\text{pred} - \text{goal\_pred}$$

- An indication of the *raw direction* the current prediction missed by
  - If this is a *positive* number, the prediction is too *high*
  - If this is a *negative* number, the prediction is too *low*
- Also an indication of the *amount* the current prediction missed by:
  - If this is a *big* number, the prediction has missed by a *big* amount
  - If this is a *small* number, the prediction has missed by a *small* amount

- **What are scaling, negative reversal, and stopping?**
  - Have the combined effect of translating the pure error into the absolute amount you want to change weight.
  - **Stopping**
    - Do not adjust weights when input is zero
  - **Negative Reversal**
    - Ensuring that weight moves in the correct direction even when input is negative
  - **Scaling**
    - Weight changes are proportional to input size

- The golden method for neural learning

```
pred = input * weight
error = (pred - goal_pred) ** 2
delta = pred - goal_pred
weight_delta = delta * input
weight = weight - weight_delta
```

- This approach adjusts each weight in the correct direction and by the correct amount so that error reduces to 0
  - Secret lies in the `pred` and `error` calculations

- Measuring error and finding the direction and amount
  - Gradient Descent in Action
    - Albeit in a bit of an oversimplified environment!

## Basic Gradient Descent Algorithm

```
weight = 0.5
goal_pred = 0.8
input = 0.5
for iteration in range(20):
    pred = input * weight
    error = (pred - goal_pred) ** 2
    direction_and_amount = (pred - goal_pred) * input
    weight = weight - direction_and_amount
    print("Error:" + str(error) + " Prediction:" + str(pred))
```

```
Error:0.3025 Prediction:0.25
Error:0.170 Prediction:0.388
Error:0.096 Prediction:0.491
...
Error:1.709e-05 Prediction:0.796
Error:9.615e-06 Prediction:0.797
Error:5.408e-06 Prediction:0.798
```

- Combining the `pred` and `error` calculations

```
error = ((input * weight) - goal_pred) ** 2
```

- For any input/prediction pair, an exact relationship can be defined between error and weight
- This is found by combining the prediction and error formulas
  - E.g when `goal_pred = 0.8` and `input = 0.5`

```
error = ((0.5 * weight) - 0.8) ** 2
```


# Error/Weight Relationship

- The relationship between error and weight

$$\text{error} = ((0.5 * \text{weight}) - 0.8) ** 2$$

## Graph

- Black dot: the current point of both error and weight
- The dotted circle is where we want to be ( $\text{error} == 0$ ).



## Key Points

- No matter where you are, the slope also points to the *minimum* point in the function.
- You can use this to find the minimum

# Error/Weight Relationship

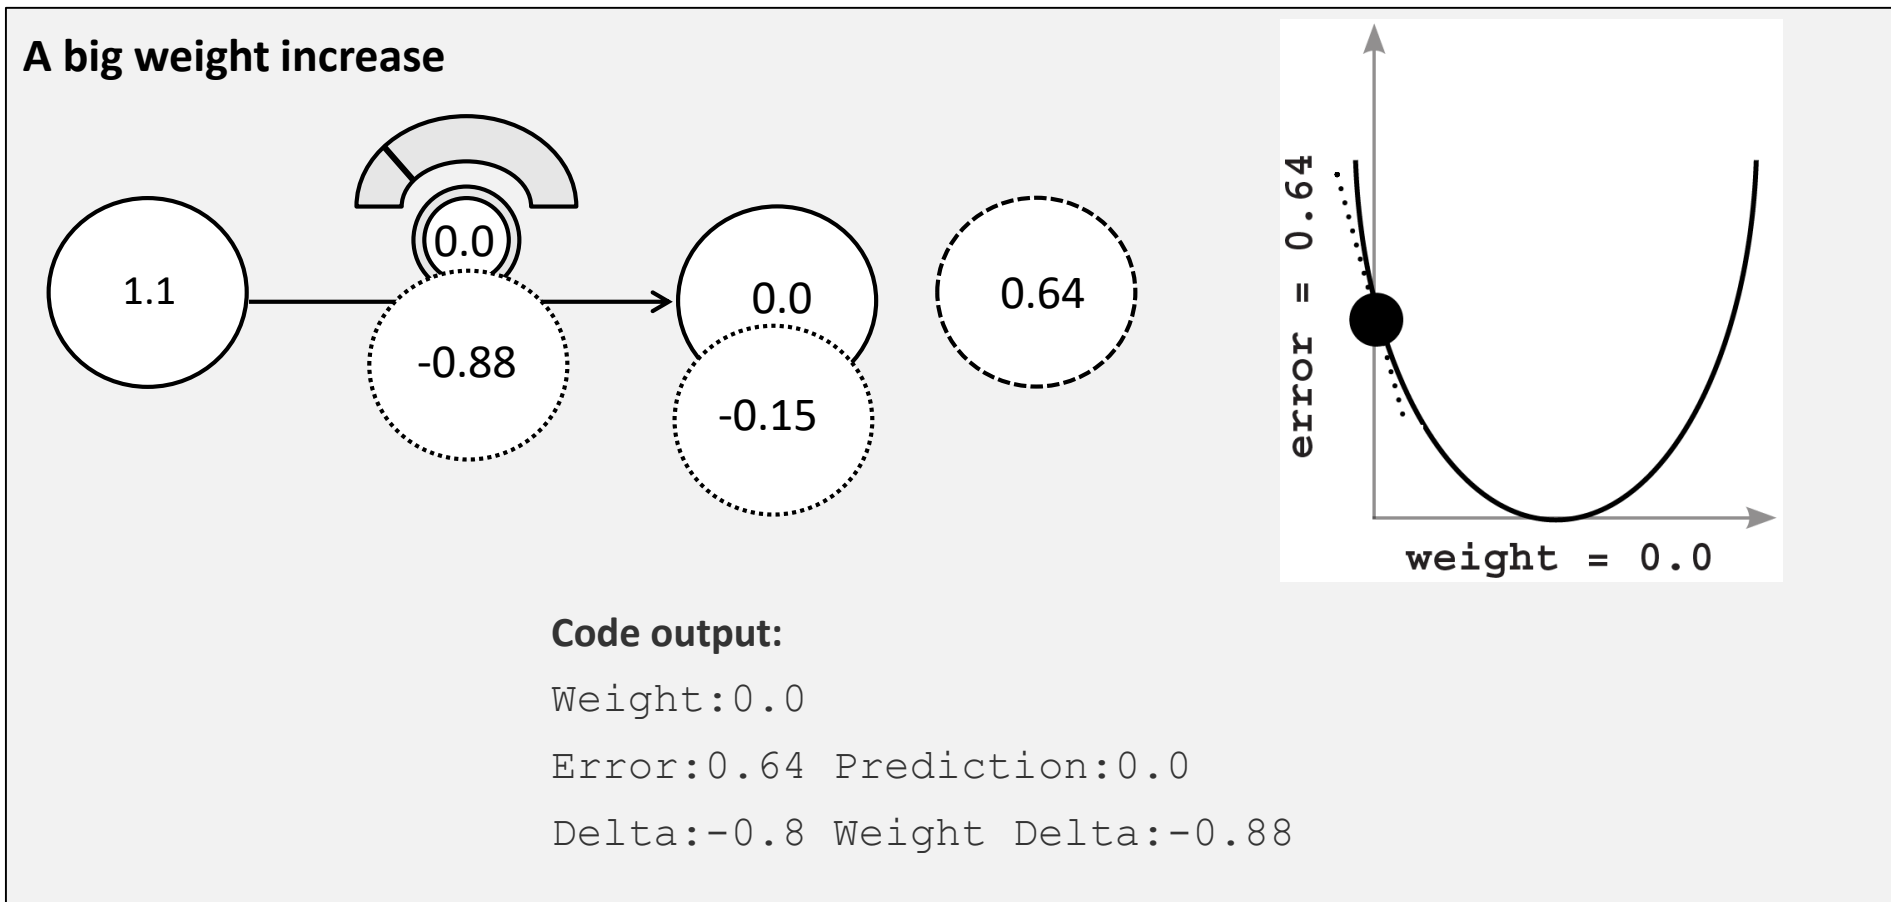
- Can we find the minima?
  - Consider the following function

```
weight, goal_pred, input = (0.0, 0.8, 1.1)

for iteration in range(4):
    print("-----\nWeight:" + str(weight))
    pred = input * weight
    error = (pred - goal_pred) ** 2
    delta = pred - goal_pred
    weight_delta = delta * input
    weight = weight - weight_delta
    print("Error:" + str(error) + " Prediction:" + str(pred))
    print("Delta:" + str(delta) + " Weight Delta:" + str(weight_delta))
```


Note:  
direction\_and\_amount  
has been split over two lines

- First update:







- Second update:



- Third update:

Small overshoot, head back, but only a little



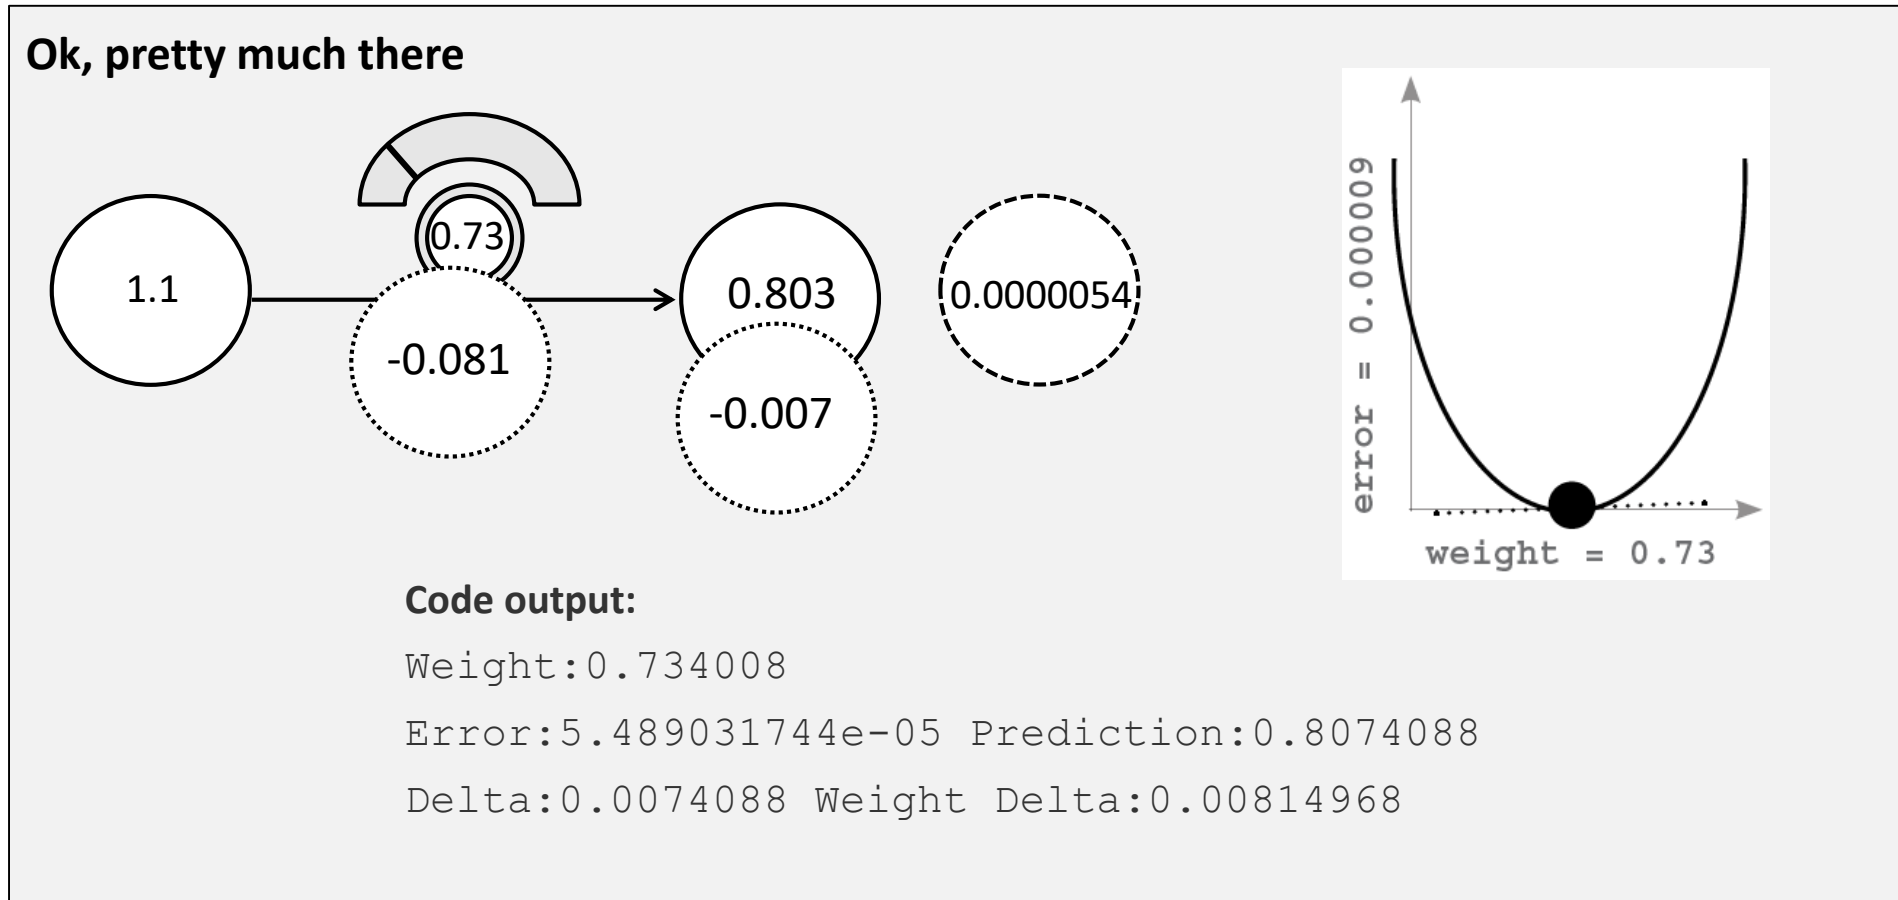
**Code output:**

Weight:0.6952

Error:0.0012446784 Prediction:0.76472

Delta:-0.03528 Weight Delta:-0.038808

- Last update:



- **What is happening?**

- Consider a function

- A function defines some sort of relationship between the input number(s) and the output number(s).

- Every function has what you might call *moving parts*

- Pieces we can tweak or change to make the output different.

```
error = ((input * weight) - goal_pred) ** 2
```

- What's controlling the relationship between `input` and the output (`error`)?

- **What is happening?**

- What's controlling the relationship between `input` and the output (`error`)?

```
error = ((input * weight) - goal_pred) ** 2
```

- We could change `goal_pred` to reduce error
  - Essentially denying we missed
- We could change `input` to reduce error
  - This would not work in the real world!
- We could change the squaring, or the mathematical operators
  - This is just changing how you calculate error in the first place.

- **What is happening?**

- What's controlling the relationship between `input` and the output (`error`)?

```
error = ((input * weight) - goal_pred) ** 2
```

- The only thing left we can change is `weight`
  - Adjusting this doesn't change your perception of the world, doesn't change your goal, and doesn't destroy your error measure.
- Changing weight means the function conforms to the patterns in the data.

- **Key Message:**
  - Learning is adjusting the weight to reduce the error to 0
  - Knowing how to do this is all about understanding the relationship between weight and error
    - How does changing one variable effect the other?
  - This is the sensitivity between the two variables
  - **Goal:** know the direction and the amount that error changes when you change weight
    - This relationship is defined through the derivative of the error function

- **Derivatives**
  - With derivatives, you can pick any two variables in any formula, and know how they interact
- **Use in training a neural network**
  - A neural network is essentially a bunch of weights used to compute an error function
  - For any error function we can compute the relationship between any `weight` and the final `error` of the network.
    - With this information, we change each `weight` in the network to reduce `error` to 0

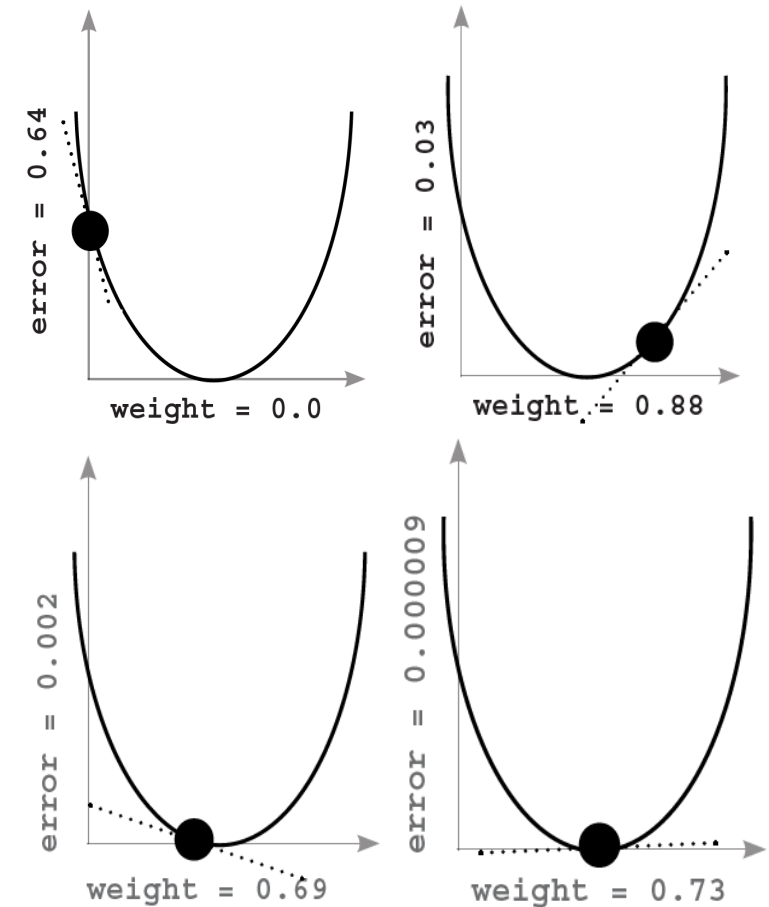


- **Gradient Descent for neural learning**

```
pred = input * weight
error = (pred - goal_pred) ** 2
delta = pred - goal_pred
weight_delta = delta * input
weight = weight - weight_delta
```

- `error` is a measure of how much the network missed by
  - We define `error` to be always positive
- `weighted_delta` is the derivate of `weight` and `error`
  - Defines the relationship between the weights and the total error

- **weighted\_delta** is the derivate
    - For each point of our error function the derivative tells you how much `error` changes when we change `weight`
    - The derivative is always pointed in the opposite direction to the minimum point
    - To reduces `error` to zero, move the `weight` value opposite the gradient value
- $$\text{weight} = \text{weight} - \text{weight\_delta}$$
- It is the minus sign in the weight update step that enables us to go in the opposite direction




# Gradient Descent: Divergence

- **Divergence in Gradient Descent**
  - Consider the following function


```
weight, goal_pred, input = (0.5, 0.8, 2.0)

for iteration in range(20):
    print("-----\nWeight:" + str(weight))
    pred = input * weight
    error = (pred - goal_pred) ** 2
    delta = pred - goal_pred
    weight_delta = delta * input
    weight = weight - weight_delta
    print("Error:" + str(error) + " Prediction:" + str(pred))
```

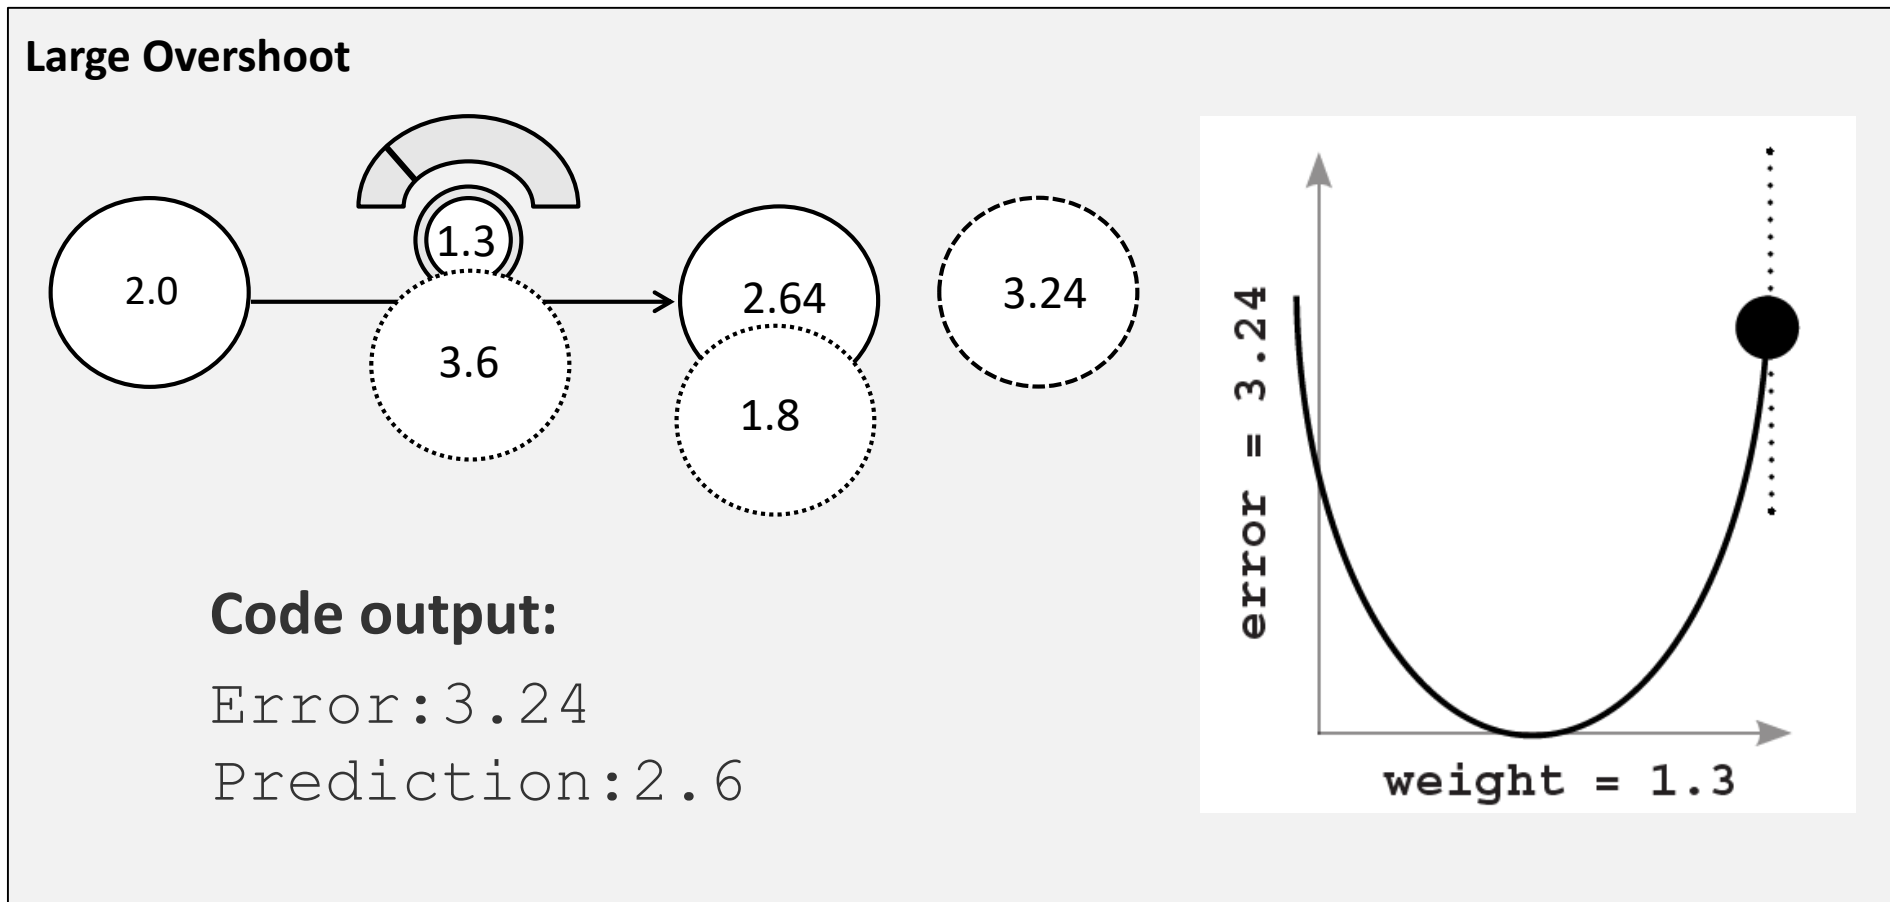
- First update:



- Second update:




- Third update:



- **Observation: The predictions are exploding**
  - At every update the network overcorrects the weights
    - They alternate from negative to positive and negative to positive, getting farther away from the true answer at every step. In

```
Error:0.04 Prediction:1.0  
Error:0.36 Prediction:0.2  
Error:3.24 Prediction:2.6  
...  
Error:6.67e+14 Prediction:-25828031.8  
Error:6.00e+15 Prediction:77484098.6  
Error:5.40e+16 Prediction:-232452292.6
```



- **Why are the predictions exploding?**

- Consider how the weights are updated:

```
weight = weight - (input * (pred - goal_pred))
```

- If the input is sufficiently large, the weight update will also be large, even when the error is small.
- When you have a large weight update and a small error, the network overcorrects
- The bigger the error, the more the network overcorrects



- **How to prevent divergence:**

- Introduce a new variable,  $\alpha$ , to scale the weight updates

```
weight = weight - derivative  
weight = weight - (alpha*derivative)
```

- In most cases, this involves multiplying the weight update by a single real-valued number between 0 and 1

- **How to set Alpha?**

$$\text{weight} = \text{weight} - (\text{alpha} * \text{derivative})$$

- Empirically, watching errors over time
  - If it starts diverging (going up), then the alpha is too high and needs to be decreased
  - If learning is happening too slowly, then the alpha is too low and should be increased

# Gradient Descent: Learning Rate

- **Divergence Example**

- Effect of introducing `alpha` into previous example

```
weight = 0.5
goal_pred = 0.8
input = 2
alpha = 0.1

for iteration in range(20):
    pred = input * weight
    error = (pred - goal_pred) ** 2
    derivative = input * (pred - goal_pred)
    weight = weight - (alpha * derivative)
    print("Error:" + str(error) + " Prediction:" + str(pred))
```

- **Divergence Example**

- Effect of introducing `alpha` into previous example
  - The neural network can now make good predictions again

Using:

`weight = weight - derivative`

Error:0.04 Prediction:1.0

Error:0.36 Prediction:0.2

Error:3.24 Prediction:2.6

...

Error:6.67e+14 Prediction:-25828031.8

Error:6.00e+15 Prediction:77484098.6

**Error:5.40e+16 Prediction:-232452292.6**

Using:

`weight = weight - (alpha*derivative)`

Error:0.04 Prediction:1.0

Error:0.0144 Prediction:0.92

Error:0.005184 Prediction:0.872

...

Error:1.146e-09 Prediction:0.800033853319

Error:4.126e-10 Prediction:0.800020311991

**Error:1.485e-10 Prediction:0.800012187195**

# Basic Gradient Descent for Neural Learning

```
pred = input * weight
error = (pred - goal_pred) ** 2
derivative = input * (pred - goal_pred)
weight = weight - (alpha * derivative)
```

- `Pred` is the output of the network given the current input
  - We want the network to learn to make correct predictions
- `error` is a measure of how much the network missed by
  - We define `error` to be always positive, we learn by reducing the `error` to zero
- `derivative` is the derivate of `weight` and `error`
  - Predicts both direction and amount to adjust the weights
- `Alpha` scales the weight update
  - Helps minimise divergence effects when the input is large