

## Exercise - DL Tutorial 1

Please complete the following notebook and submit your solutions to manuel.milling@informatik.uni-augsburg.de

**student name:**

Solutions from exercise sheet 3 (class methods below).

```
In [1]: #Equations from handout 3 and for are referred to as (3.X) and (4.X)

import numpy as np
#numpy random seed
np.random.seed(42)

trainx, trainy, testx, testy = np.load('mnist.npy', allow_pickle=True)
print("Trainx shape: {}".format(trainx.shape))
print("Trainy shape: {}".format(trainy.shape))
print("Testx shape: {}".format(testx.shape))
print("Testy shape: {}".format(testy.shape))

def sigmoid(X):
    return 1/(1 + np.exp(-X))

def softmax(X):
    #more stable
    eps = X.max()
    return np.exp(X + eps) / (np.sum(np.exp(X + eps), axis=1).reshape((X.shape[0], 1)))

def fcc_one_layer(X, W, b, activation):
    return activation(np.matmul(X, W) + b)

def cross_entropy(pred_logits, y):
    num_data_points = pred_logits.shape[0]
    correct_logits = pred_logits[np.arange(num_data_points), y]
    return np.mean(-np.log(correct_logits))

def accuracy(logits, labels):
    class_predictions = np.argmax(logits, axis=1)
    return np.mean(class_predictions == labels)

Trainx shape: (60000, 784)
Trainy shape: (60000,)
Testx shape: (10000, 784)
Testy shape: (10000,)
```

1. Implement the error of the last layer.

```
In [2]: def delta_last_layer(h, y):
        """
        :param h: softmax activations of shape (num_examples, num_classes)
        :param y: correct labels of shape num_classes
        :return: delta of softmax
        """
        num_data_points = h.shape[0]
        # get  $H^n_{mi}$  for (4.31)
        correct_logits = h[np.arange(num_data_points), y]
        # (4.31)  $i \neq j$ 
        h_i_neq_j = - np.reshape(correct_logits, (correct_logits.shape[0], 1)) * h
        # (4.31)  $i=j$ 
        h_i_eq_j = correct_logits*(1- correct_logits)
        #replace the  $i=j$  terms in  $i \neq j$  matrix
        h_i_neq_j[np.arange(num_data_points),y] = h_i_eq_j
        # (4.30)
        h_i_neq_j = h_i_neq_j / np.reshape(correct_logits, (correct_logits.shape[0],
1))
        #transpose h --> delta shape
        return - np.transpose(h_i_neq_j)
```

```
In [3]: def delta_last_layer_easy_approach(h, y):
        # create one hot vectors for every row
        one_hots = np.zeros((h.shape[0], h.shape[1]))
        one_hots[np.arange(h.shape[0]), y] = 1.0
        # (4.30) and (4.31) can be reshaped
        return (h - one_hots).T
```

1. Implement the derivative of the sigmoid function in terms of the sigmoid function.

```
In [4]: def del_sigmoid(h):
        """
        :param h: output of sigmoid function, i.e.  $h = \sigma(x)$ 
        """
        # $dh/dx = d\sigma(x)/dx = \sigma(x)(1-\sigma(x)) = h(1-h)$ 
        return h * (1 - h)
```

1. Implement the backpropagation as a class method.
2. Implement the the optimisation step as a class method.

```

In [11]: class fcc:
    def __init__(self, n_input, n_hidden1, n_hidden2, n_out):
        # Initialisation and Declaration of class variables
        self.W_i_h1 = np.random.randn(n_input, n_hidden1)
        self.b_h1 = np.random.randn(n_hidden1)
        self.W_h1_h2 = np.random.randn(n_hidden1, n_hidden2)
        self.b_h2 = np.random.randn(n_hidden2)
        self.W_h2_o = np.random.randn(n_hidden2, n_out)
        self.b_out = np.random.randn(n_out)
        # not necessary, but for better overview
        self.X = None
        self.h1 = None
        self.h2 = None
        self.out = None
        self.dW_i_h1 = None
        self.db_h1 = None
        self.dW_h1_h2 = None
        self.db_h2 = None
        self.dW_h2_o = None
        self.db_out = None

        # calculation of network parameters
        n_trainable_bias = self.b_h1.shape[0] + self.b_h2.shape[0] + self.b_out.sh
pe[0]
        n_trainable_weights = self.W_i_h1.shape[0] * self.W_i_h1.shape[1] + self.W_
h1_h2.shape[0] * self.W_h1_h2.shape[1] + self.W_h2_o.shape[0] * self.W_h2_o.shape
[1]
        print("Number of parameters: {}".format(n_trainable_bias + n_trainable_weig
hts))

    def forward_propagation(self, X):
        self.X = X
        # (3.4)
        self.h1 = fcc_one_layer(X, self.W_i_h1, self.b_h1, sigmoid)
        # (3.4)
        self.h2 = fcc_one_layer(self.h1, self.W_h1_h2, self.b_h2, sigmoid)
        # (3.4)
        self.out = fcc_one_layer(self.h2, self.W_h2_o, self.b_out, softmax)
        return self.out

    def backprop(self, y):
        self.num_train_ex = y.shape[0]
        self.delta_out = delta_last_layer_easy_approach(self.out, y)
        # (4.27)
        self.dW_h2_o = np.transpose(np.matmul(self.delta_out, self.h2))/self.num_tr
ain_ex
        # (4.28)
        self.db_out = np.mean(self.delta_out, axis=1)
        # (4.26)
        self.delta_h2 = np.matmul(self.W_h2_o, self.delta_out) * np.transpose(del_s
igmoid(self.h2))
        # (4.27)
        self.dW_h1_h2 = np.transpose(np.matmul(self.delta_h2, self.h1)) / self.num_
train_ex
        # (4.28)
        self.db_h2 = np.mean(self.delta_h2, axis=1)
        # (4.26)
        self.delta_h1 = np.matmul(self.W_h1_h2, self.delta_h2) * np.transpose(del_s
igmoid(self.h1))
        # (4.27)
        self.dW_i_h1 = np.transpose(np.matmul(self.delta_h1, self.X)) / self.num_tr
ain_ex
        # (4.28)
        self.db_h1 = np.mean(self.delta_h1, axis=1)

```

1. Implement the training routine.

```
In [15]: learning_rate = 0.1  
         neural_net = fcc(784, 400, 400, 10)  
  
         Number of parameters: 478410
```

Normal Gradient Descent

```
In [13]: #1000 trainingssteps
num_iterations = 1000
for i in range(num_iterations):
    # evaluate after each 100 steps
    if i % 100 == 0:
        print
        (
            "-----"
            "-----"
            print("Iteration:\t\t{}".format(i))
            logits = neural_net.forward_propagation(testx)
            print("Test Loss:\t\t{}".format(cross_entropy(logits, testy)))
            print("Test Accuracy:\t\t{}".format(accuracy(logits, testy)))
            logits = neural_net.forward_propagation(trainx)
            if i%100 == 0:
                print("Train Loss:\t\t{}".format(cross_entropy(logits, trainy)))
                print("Train Accuracy:\t\t{}".format(accuracy(logits, trainy)))
                print
                (
                    "-----"
                    "-----"
                    neural_net.backprop(trainy)
                    neural_net.gradient_step(learning_rate)
                )
            print
            (
                "-----"
                "-----"
            )
            print("Iteration:\t\t{}".format(i))
            logits = neural_net.forward_propagation(testx)
            print("Test Loss:\t\t{}".format(cross_entropy(logits, testy)))
            print("Test Accuracy:\t\t{}".format(accuracy(logits, testy)))
            logits = neural_net.forward_propagation(trainx)
            print("Train Loss:\t\t{}".format(cross_entropy(logits, trainy)))
            print("Train Accuracy:\t\t{}".format(accuracy(logits, trainy)))
```

```
-----  
-----  
Iteration:           0  
Test Loss:           19.479279548133853  
Test Accuracy:       0.1353  
Train Loss:          19.653523403657122  
Train Accuracy:      0.13226666666666667  
-----  
-----  
-----
```

```
-----  
-----  
Iteration:           100  
Test Loss:           2.282769668031916  
Test Accuracy:       0.5941  
Train Loss:          2.4319569058890083  
Train Accuracy:      0.57965  
-----  
-----  
-----
```

```
-----  
-----  
Iteration:           200  
Test Loss:           1.573850466163723  
Test Accuracy:       0.7007  
Train Loss:          1.6520382064719452  
Train Accuracy:      0.6910666666666667  
-----  
-----  
-----
```

```
-----  
-----  
Iteration:           300  
Test Loss:           1.296911010188398  
Test Accuracy:       0.7445  
Train Loss:          1.3371711957554957  
Train Accuracy:      0.7404666666666667  
-----  
-----  
-----
```

```
-----  
-----  
Iteration:           400  
Test Loss:           1.142345153274019  
Test Accuracy:       0.7731  
Train Loss:          1.1521656169086645  
Train Accuracy:      0.7698666666666667  
-----  
-----  
-----
```

```
-----  
-----  
Iteration:           500  
Test Loss:           1.0402612586580389  
Test Accuracy:       0.7876  
Train Loss:          1.0261024142379838  
Train Accuracy:      0.7897  
-----  
-----  
-----
```

```
-----  
-----  
Iteration:           600  
Test Loss:           0.9656370013641344  
Test Accuracy:       0.8009  
Train Loss:          0.9324061314836917  
Train Accuracy:      0.8044166666666667  
-----  
-----  
-----
```

## Stochastic Gradient Descent

```
In [16]: #rerun initialisation of neural_network before executing

batch_size = 64
permutation = np.arange(trainx.shape[0])
epochs = 10
for i in range(epochs):
    print
    ( "-----"
    -----")
    print("Epoch:\t\t{}".format(i))
    logits = neural_net.forward_propagation(testx)
    print("Test Loss:\t\t{}".format(cross_entropy(logits, testy)))
    print("Test Accuracy:\t\t{}".format(accuracy(logits, testy)))
    print
    ( "-----"
    -----")
    #create new permutation of trainings examples
    np.random.shuffle(permutation)
    #loop over epoch (= one permutation of all data)
    for j in range(int(trainx.shape[0]/batch_size)):
        #take one minibatch
        batch = permutation[j*batch_size:(j+1) * batch_size]
        trainx_batch = trainx[batch]
        trainy_batch = trainy[batch]
        logits = neural_net.forward_propagation(trainx_batch)
        # print every 100 training steps
        if j%100 == 0:
            print("Train Loss:\t\t{}".format(cross_entropy(logits, trainy_batch)))
            print("Train Accuracy:\t\t{}".format(accuracy(logits, trainy_batch)))
            neural_net.backprop(trainy_batch)
            neural_net.gradient_step(learning_rate)
        logits = neural_net.forward_propagation(testx)
    print
    ( "-----"
    -----")
    print("Final:\t\t")
    print("Test Loss:\t\t{}".format(cross_entropy(logits, testy)))
    print("Train Accuracy:\t\t{}".format(accuracy(logits, testy)))
```



```
-----  
-----  
Epoch:          0  
Test Loss:              20.258685506517264  
Test Accuracy:          0.0926  
-----
```

```
-----  
Train Loss:              19.821905152851123  
Train Accuracy:          0.109375  
Train Loss:              1.8949693999235606  
Train Accuracy:          0.5625  
Train Loss:              2.0527701690070286  
Train Accuracy:          0.671875  
Train Loss:              1.6121803157171737  
Train Accuracy:          0.65625  
Train Loss:              2.077289663492683  
Train Accuracy:          0.71875  
Train Loss:              1.453578726376255  
Train Accuracy:          0.71875  
Train Loss:              1.3055405866063823  
Train Accuracy:          0.640625  
Train Loss:              1.0969084533218267  
Train Accuracy:          0.828125  
Train Loss:              0.7339847306076936  
Train Accuracy:          0.796875  
Train Loss:              0.717445025227063  
Train Accuracy:          0.828125  
-----
```

```
-----  
-----  
Epoch:          1  
Test Loss:              0.829246421718839  
Test Accuracy:          0.8072  
-----
```

```
-----  
Train Loss:              0.7272308116995612  
Train Accuracy:          0.859375  
Train Loss:              0.4546252790733857  
Train Accuracy:          0.875  
Train Loss:              0.41340715768929903  
Train Accuracy:          0.890625  
Train Loss:              0.5671620821361496  
Train Accuracy:          0.875  
Train Loss:              0.607704523731818  
Train Accuracy:          0.859375  
Train Loss:              0.9237634688270586  
Train Accuracy:          0.828125  
Train Loss:              0.6420787375359898  
Train Accuracy:          0.828125  
Train Loss:              0.6920089752589389  
Train Accuracy:          0.796875  
Train Loss:              0.5326557310728961  
Train Accuracy:          0.875  
Train Loss:              0.5589444050035213  
Train Accuracy:          0.859375  
-----
```

```
-----  
-----  
Epoch:          2  
Test Loss:              0.6006060028525041  
Test Accuracy:          0.8547  
-----
```

```
-----  
Train Loss:              0.21403026054300162  
Train Accuracy:          0.921875  
Train Loss:              0.5835133616491388  
-----
```

In [ ]: