

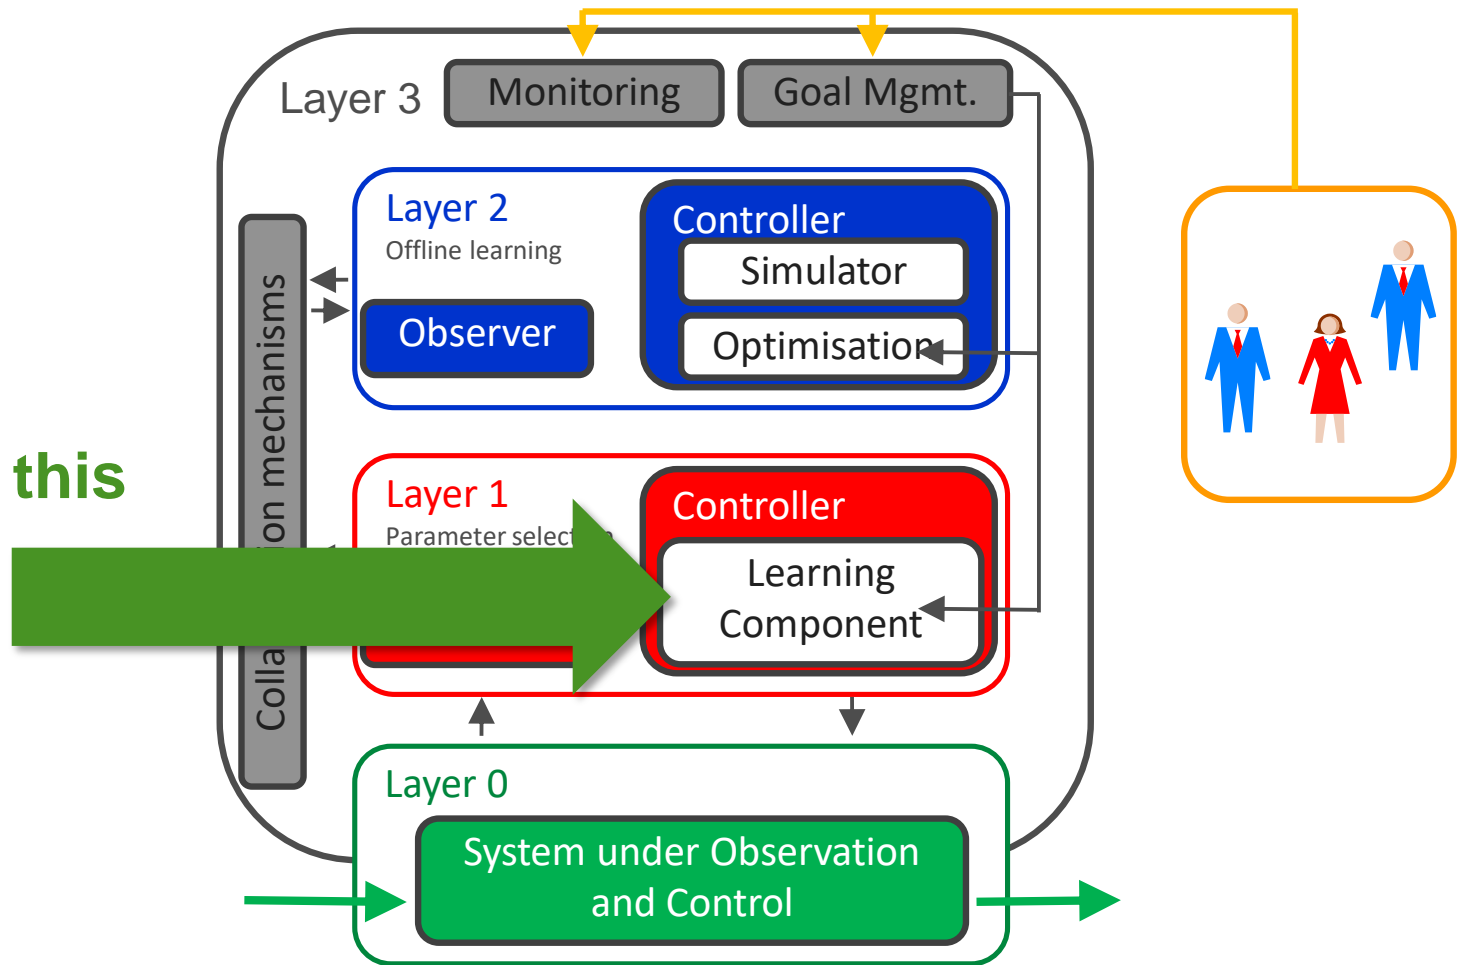

Organic Computing

Lecture
Organic Computing II
Summer term 2020

Chapter 6: Learning

Lecturer: Jörg Hähner

Focus of this
Lecture!



Content

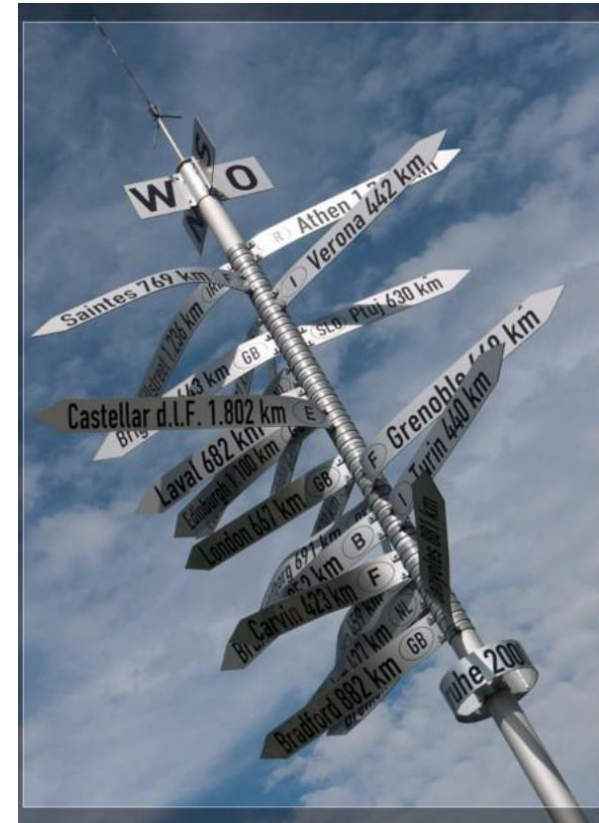
- Motivation
- Extended Classifier System
- XCS-O/C
- Artificial Neural Networks
- Conclusion and further readings

Goals

Students should be able to:

- Explain what machine learning is and why it is needed in organic systems.
- Compare basic concepts such as supervised vs. reinforcement learning or exploration vs. exploitation.
- Outline an XCS and explain the main loop with all components.
- Discuss the necessary modifications to XCS for OC.
- Explain basic Feed-Forward ANNs and how backpropagation works

- **Motivation**
- Extended Classifier System
- XCS-O/C
- Artificial Neural Networks
- Conclusion and further readings



Complex learning tasks:

- **Sparse and imbalanced data**
 - E.g. due to non-uniform distributions and class imbalances
- **Non-stationary environments**
 - May exhibit severe changes in the target concepts.
 - Also called *concept drift*.
- **Necessity of exploration boundaries**
 - Unrestricted or unknown feature spaces
 - Continuous or large discrete action spaces
 - Legal constraints
 - Trial-and-error must be avoided!
- **Complexity of underlying problem space**
 - Functions mapping inputs to certain outputs regarding are complex.
 - E.g. due to their dimensionality, continuity, obliqueness and curvature.
- Knowledge and expected behaviour must be represented in a **human comprehensible** manner (e.g. as rules).

- Machine learning techniques seem to be a **promising approach for continuous self-improvement** in Organic Computing systems.
→ How can computers be programmed so that **problem solving** capabilities are built up by specifying “**what is to be done**” rather than “how to do it”? (Holland, 1975).
- Major issues:
 - How can the system react to unforeseen situations?
 - How can the system automatically improve its performance (if possible) at runtime?
 - How can knowledge (and expected behaviour) be encoded in a human comprehensible manner?
 - Overall: flexible and **autonomous reaction to changes of the environments and/or the system itself are desirable.**

- Machine learning is used when:
 - human expertise does not exist (navigating on Mars)
 - humans are unable to explain their expertise (speech recognition)
 - solution changes in time (routing on a computer network)
 - solution needs to be adapted to particular cases (user biometrics)
 - human learning is not feasible (map protein sequences to secondary structures)
 - ...
- The slides for the introduction part to Machine Learning are mainly based on the book:

- Common definition:

*“A computer program is said to **learn** from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .”*

- *Example:*

- consider a program that learns playing checker
- task T : play checker
- performance measure P : percentage of the games won
- experience E : play against itself



- consider a given set of examples (**training set**)
- make accurate predictions about future examples (**validation set**)
- with the goal of predicting concrete values for unknown feature sets, the **learning problem** can be mapped to learning a **function**

$f(X) = Y$, whereas X is an input example and Y the desired output

Supervised Learning

training = desired (target) output



error = (target output – actual system output)

- **Supervised learning** means that the **training set** of (X, Y) is provided (by a ``supervisor'' or ``teacher'')
- **Unsupervised learning** means only the input data X and feedback on the prediction performance are provided
- Example supervised learning: Face recognition
 - x : Bitmap picture of person's face
 - $f(x)$: Name of the person
- Example unsupervised learning: Customer segmentation
 - x : Customer's buying habits, address, age,...
 - $f(x)$: Customer class

Supervised learning

1. **Prediction** of future cases:
Use the learned model (or *hypothesis*) to predict the output for future input
2. **Knowledge extraction**:
The hypothesis is easy to understand
3. **Compression**:
The hypothesis is simpler than the data it explains
4. **Outlier detection**:
Exceptions that are not covered by the hypothesis (e.g. fraud)

Unsupervised learning

1. **Clustering** of data
2. **Density estimation** of data
→ gain insights into the organisation of data

Reinforcement Learning

training information = evaluation (“rewards” / “penalties”)



Goal: achieve as much **reward** as possible!

- Act “successfully” in the environment
- Implication: maximise the sequence of rewards R_t

- What is Reinforcement Learning?
 - German: “Bestärkendes Lernen“
 - Learning from interaction
 - Goal-oriented learning
 - Learning **by/from/during** interaction with an external environment
 - Learning “what to do“ (how to map situations to actions) to maximise a numeric reward

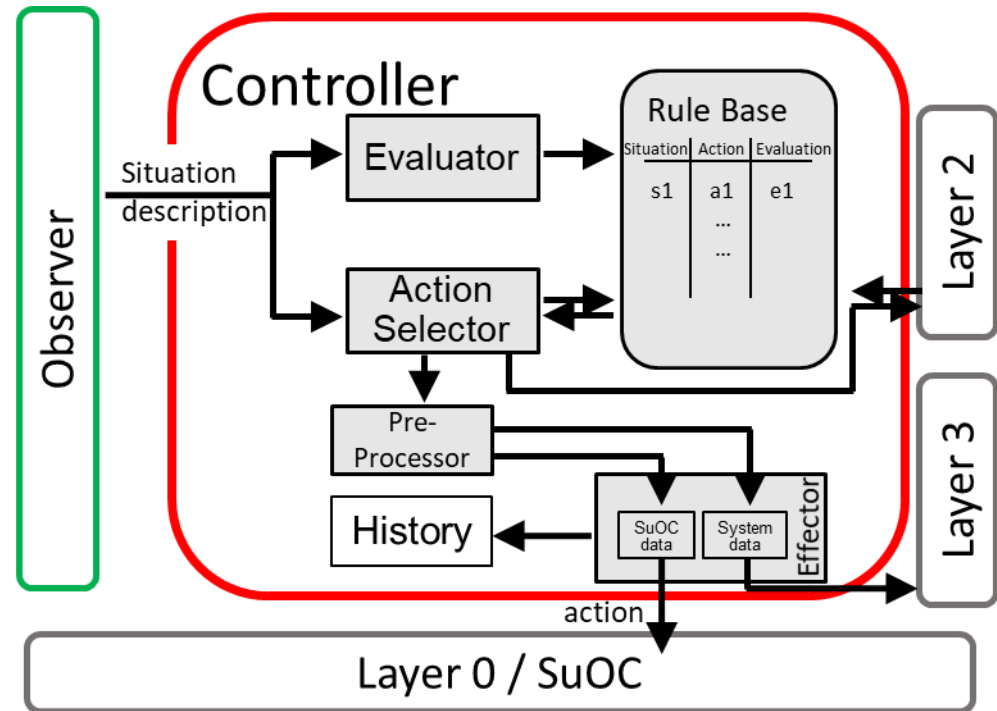
- Learning based on feedback
 - feedback tells the system how well it performs (**not** what it should be doing)
 - no supervised output; but a delayed “reward”
- Aspects:
 - credit assignment problem
 - game playing
 - robot in a maze
 - multiple agents, partial observability, ...

This is exactly what we will have to cope with in most OC applications!

Recap: Layer 1 Controller

Observer/Controller

- Controller has to learn from feedback.
- Basic concept: rule-based system
- Learning is done by „book-keeping“ attributes, i.e. evaluation parameters.
- These are modified depending on the observed success.



Example: 'Woods' scenario / Maze

- Example of an Animat problem
- Basis: rectangular toroidal regular (n x m)-grid
- Each grid cell may contain a tree (t), food (F), or it may be empty.
- Food and trees fixed per instance
- Animat/agent/robot is initially randomly placed on empty cell.
- Walks around, looking for food
- In each step, agent can go to one of the eight neighboring cells (empty and food cells only).

t	t	t	t	t	t	t	t	t	t	t	t	t
t	t				t	t	t	t		t	t	
t		t	t	t		t	t		t		t	
t		t	t	t		t		t	t	t		t
t	F	t	t	t		t	t		t	t	t	t
t	t	t	t	t	t			t	t	t	t	t

Woods14

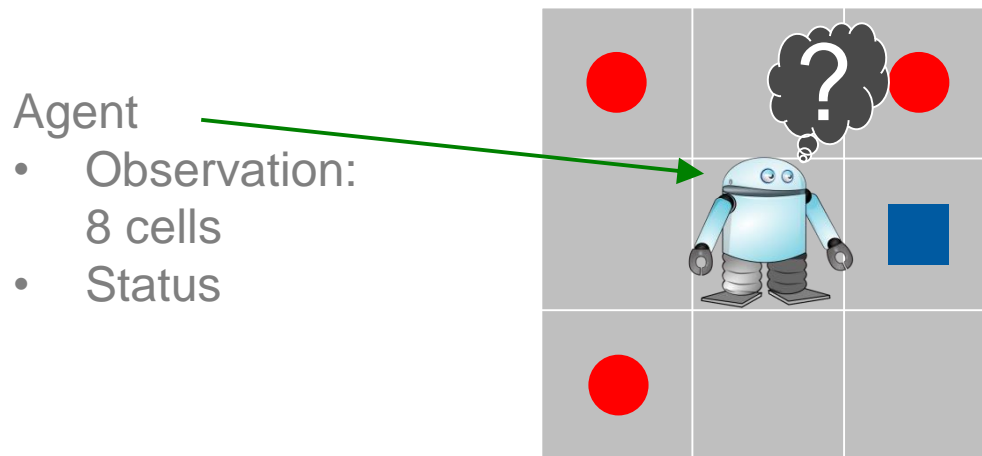
t	t	F		
t	t	t		
t	t	t		

Woods1

Woods1: optimal average number of steps to reach food: 1.7 steps (Bull & Hurst, "ZCS redux")

Example: 'Woods' and the underlying problem

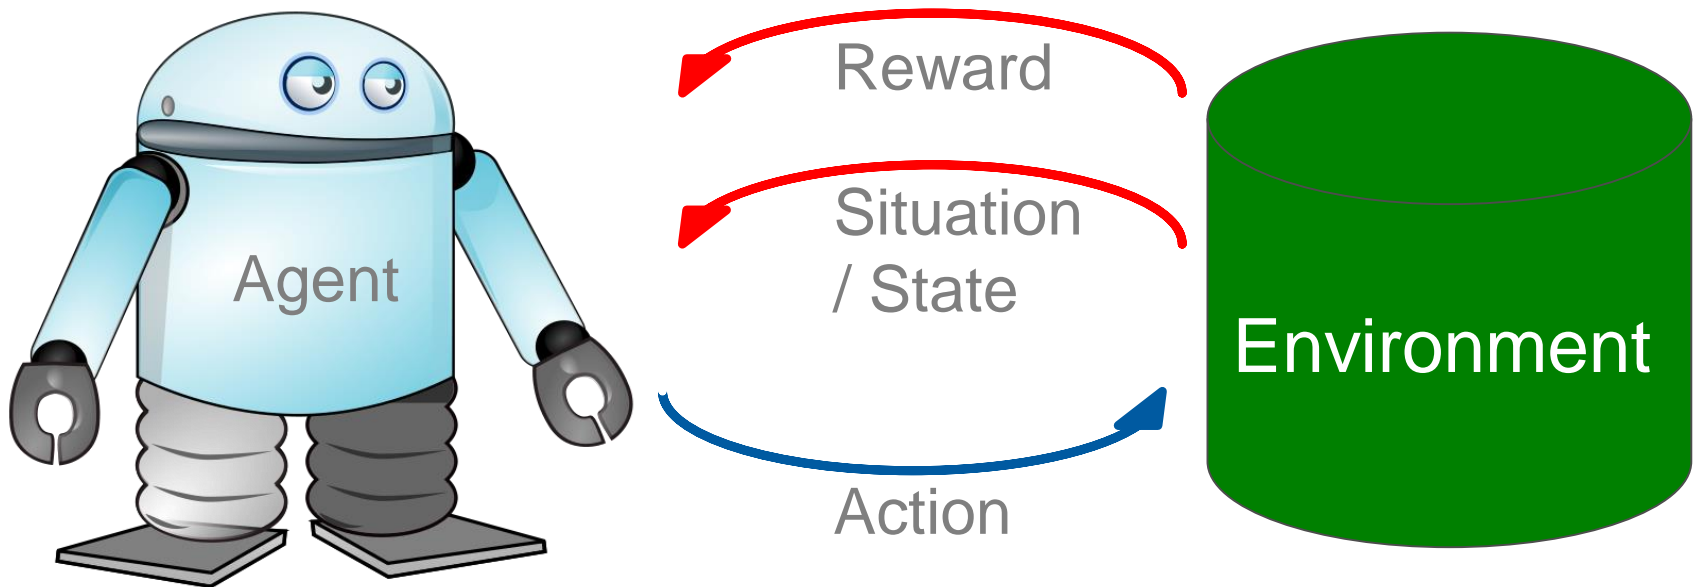
- Question: Can we build an agent that can efficiently find food “in the Woods” **without global knowledge**?
- One idea to build such an agent:
 - Suppose the agent can “see” the eight surrounding cells.
 - Based upon this perception, it has to decide where to go next.
 - Reward is paid once the food is found.



- Obstacle
- Previous cell

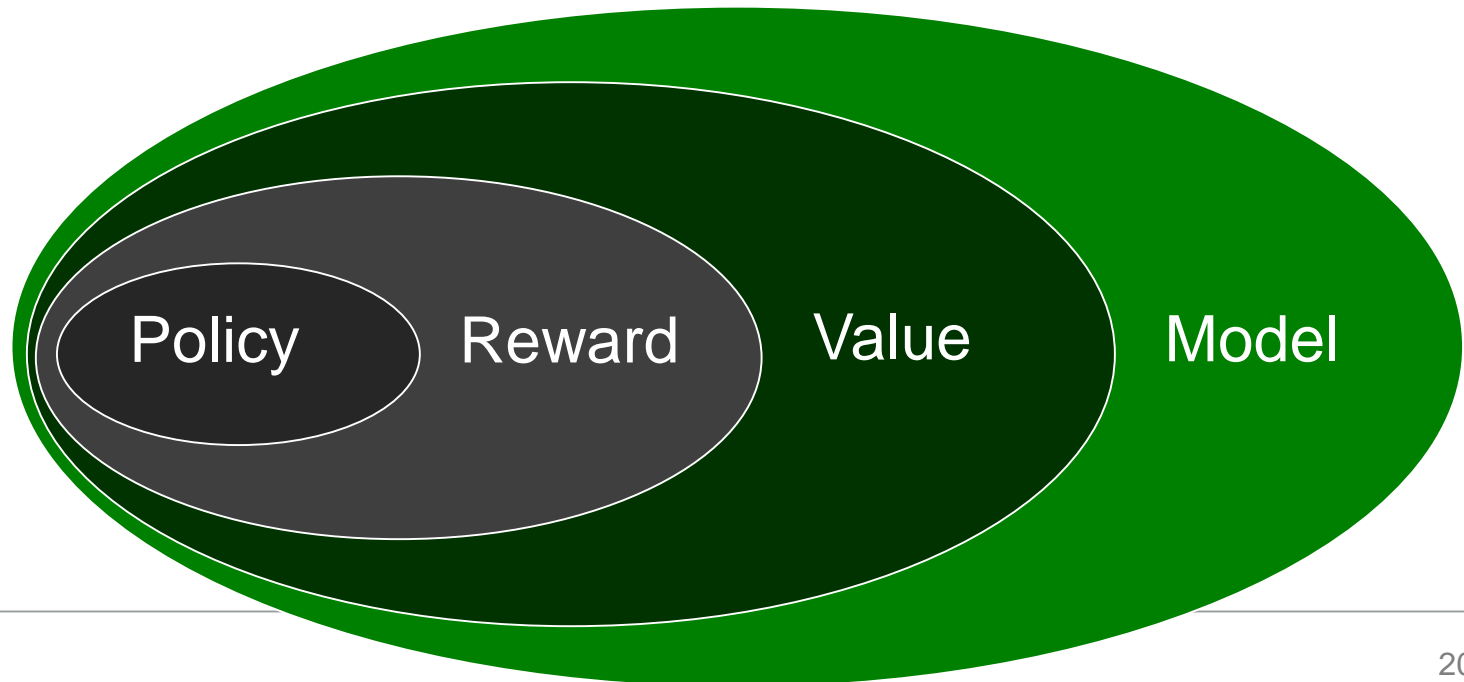
Where to go?

- The complete agent
 - Chronologically situated
 - Constant learning and planning
 - Affects the environment
 - Environment is stochastic and uncertain



Elements

- Policy: What to do in a particular situation?
- Reward: What is good or bad behaviour (experience)?
- Value: What is a good action due to the expected reward?
- Model: What follows from the actions? What is the impact?



Exploration:

A process of visiting entirely new regions of a search space.

VS.

Exploitation:

A process of visiting regions of a search space based on previously visited points (neighbourhood).

To be successful, a search algorithm needs to find a good balance between exploration and exploitation.

- **Exploration** is important in early stages:
 - seek good patterns
 - spread out through the search space
 - avoid local optima
- **Exploitation** is important in later stages:
 - exploit good patterns
 - focus on good areas of the search space
 - refine to global optimum

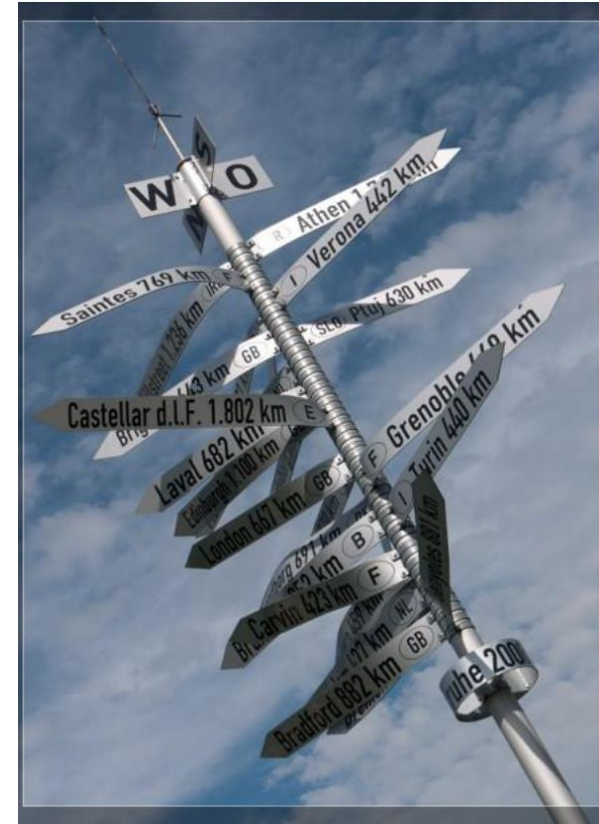
The Exploration / Exploitation Problem: Formalisation

- Suppose values are estimated:
 $Q_t(a) \approx Q^*(a)$; **estimation of action values**
- The greedy-action for time t is:

$$\begin{aligned} a_t^* &= \arg \max_a Q_t(a) \\ a_t &= a_t^* \Rightarrow \text{exploitation} \\ a_t &\neq a_t^* \Rightarrow \text{exploration} \end{aligned}$$

- Insights:
 - You cannot explore all the time, but also not exploit all the time.
 - Exploration should never be stopped, but it should be reduced.

- Motivation
- **Extended Classifier System**
- XCS-O/C
- Artificial Neural Networks
- Conclusion and further readings



Names to remember in LCS research

- Initial Learning Classifier System (LCS) was introduced by **John H. Holland** in 1975.
- He was (and still is) interested in complex adaptive systems.
- How can computers be programmed so that **problem-solving capabilities are built up by specifying “what is to be done”** rather than “how to do it”? (Holland, 1975)
- An important development in LCS was done by **Stewart W. Wilson** in 1995.
- Based on the initial approach by Holland, Wilson proposed a simplified and more efficient classifier system called Extended Classifier System (XCS).
- XCS is today one of the most studied classifier systems.
- Many extensions have been proposed.



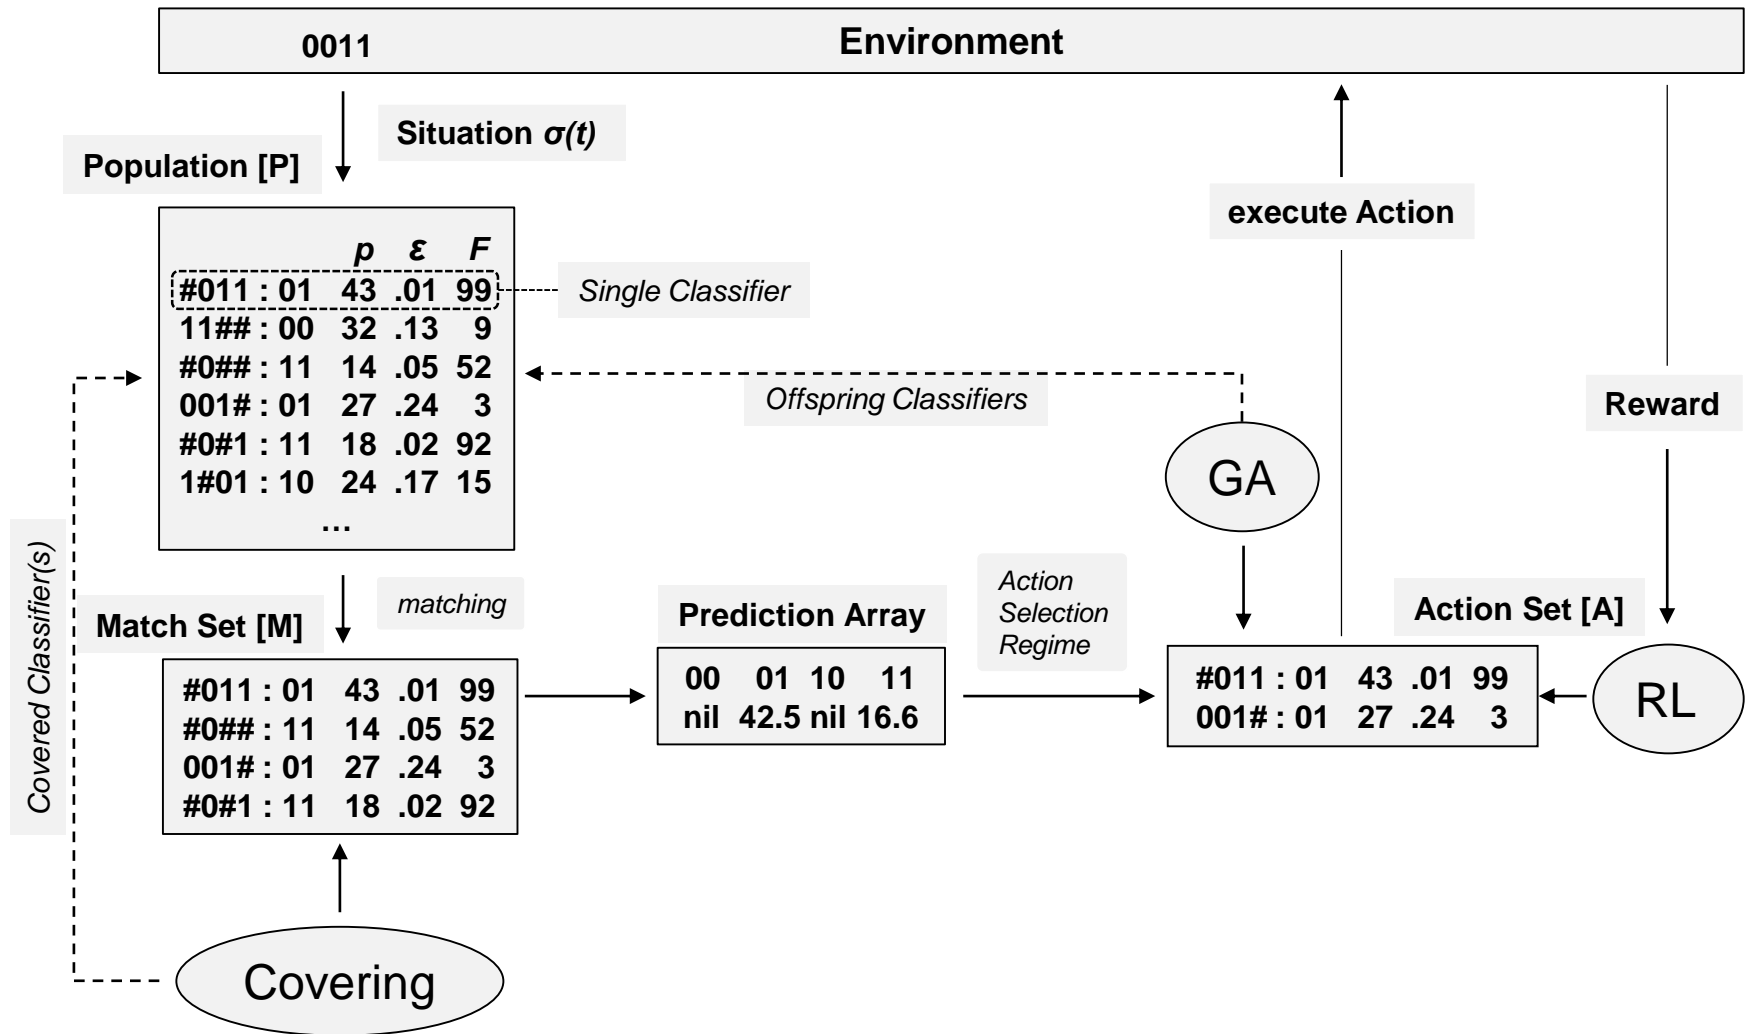
- Initially system
 - Holland designed a first system in 1978.
 - System is called CS1.
- System contains
 - **Set of classifiers** (condition/action)-pairs
 - Not called “rule” since they compete (classifier is a “may rule”)!
 - **Input interface** to receive state from the environment
 - **Output interface** to apply actions to the environment
 - Internal **message list** as an internal “workspace” for I/O
 - **Evolutionary process** (genetic algorithm) to generate new classifiers

The Extended Classifier System (XCS) by Wilson

- XCS is a rule-based (online) learning system.
- It can be used for pure classification as well as for regression problems.
- It is a derivative of the overall class of *Learning Classifier Systems (LCS)*, initially proposed by Holland in 1978 (CS-1)
- Wilson in 1994 simplified Hollands CS-1 to the so-called *Zeroth-Classifier System (ZCS)*.
- In 1995 Wilson presented the *Extended Classifier System (XCS)*.
- Initially designed for binary problems, Wilson further extended XCS toward the ability to cope with real-valued inputs (XCSR) in 2001.

- XCS stores rules (termed 'classifiers') in a **limited set** of max. N classifiers called **population** $[P]$.
- A single classifier cl is comprised of:
 - A **condition** C that defines a subspace of the input space X
 - An **action** a that determines a reaction executed on the environment (e.g. '0' and '1' for 'turn left' or 'turn right')
 - A **predicted payoff scalar** p which is an estimate of the expected reward when the action a of this classifier is selected for execution
 - An **absolute error of the payoff prediction** ϵ
 - A **measure of accuracy termed fitness** F which is some sort of inverse function of ϵ
 - Some more so-called 'book-keeping' parameters (e.g. experience)

One cycle through XCS



A single iteration through the main loop

1. At each **timestep** t , XCS retrieves a **situation** $\sigma(t)$ from the observed environment.
2. XCS scans $[P]$ for matching classifiers and builds a so-called **match set** $[M]$.
3. Among all matching classifiers, the '**prediction array**' PA calculates the most promising action a .
4. All classifiers from $[M]$ with the selected action a , from another subset $[A]$ called the **action set**.
5. The **selected action** a_{exec} is actualised on the environment which in turn delivers a so-called **payoff** or **reward** r .
6. r is used to update and refine all classifiers in $[A]$, since these particular classifiers advocated the same action as the one executed.

Why such a triple ranking by p , ε , and F ?

- What is the difference of a classifier's strength and its accuracy?
 - **Strength** = predicted payoff p
 - **Accuracy** = Fitness (inverse of prediction error ϵ)
- Is a classifier predicting a high payoff also an accurate one?
 - When a classifier predicts a high payoff, this does not necessarily mean that its prediction is correct!
- Is it beneficial to know low performing (regarding p) but highly accurate (F) classifiers?
 - Yes, indeed!
 - The system has an indicator which action delivers low payoff, and thus will decide more likely against this action.

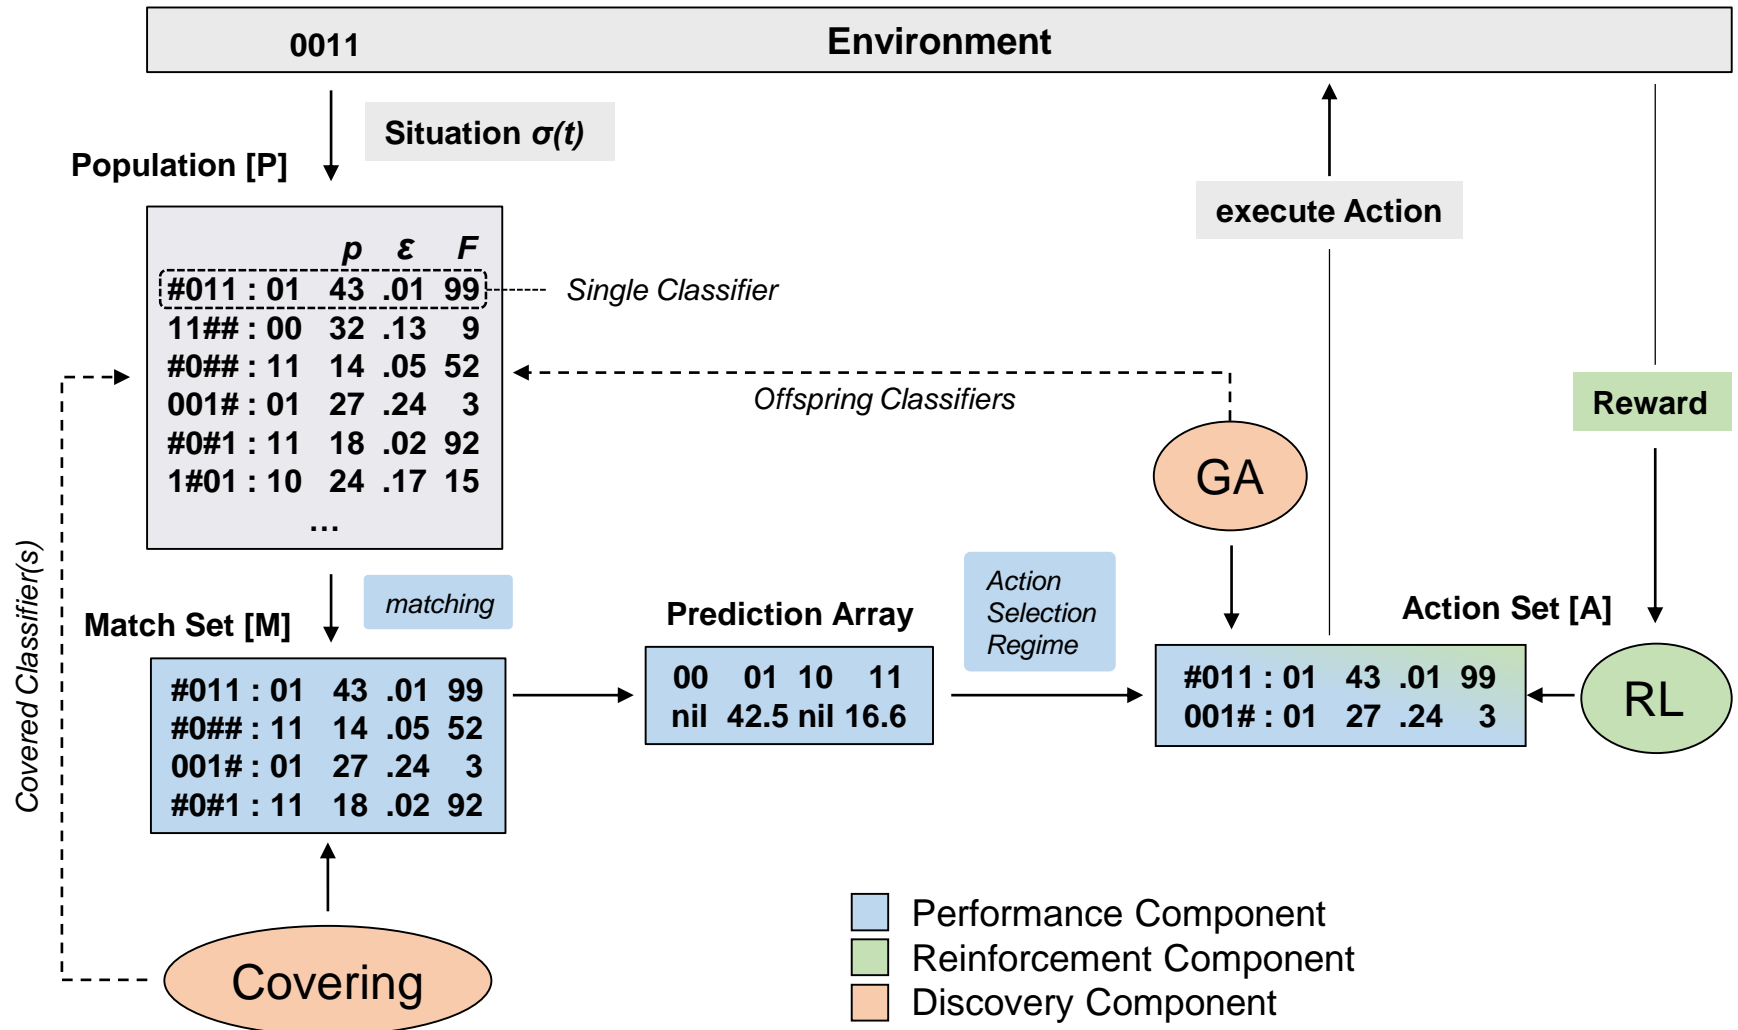
- Wilson hypothesised that XCS constructs classifiers that are **maximally general and accurate at the same time**.
- Thus, XCS attempts to construct a map/approximation of the underlying payoff-landscape, that is $X \times A \rightarrow P$, by means of single classifiers:
 - X is the input space (possible input)
 - A is the action space (possible outputs)
 - P is the payoff space (possible rewards)
- This map/approximation shall be:
 - **Complete**, in the sense that the entire payoff landscape is covered.
 - **Compact**, in terms of the # physical classifiers (macro-classifiers).
 - **Accurate**, since the system error shall be as minimal as possible (of course).
 - **Maximally general**, since the shape of a classifier (determined by its condition) shall be large enough to cover the environmental niche within X but specific enough to remain accurate.

- The separation of strength and accuracy combined with the incorporated 'niche' genetic algorithm exerts **evolutionary pressure** toward the aforementioned properties.
- The GA favours accurate (high fitness) classifiers within the environmental niche.
- Thus, accurate classifiers are more likely to reproduce and will eventually take over the environmental niche.

XCS' three main components

- Performance component
 - Matching, Payoff Prediction, Action Selection
- Reinforcement component
 - Attribute update, deferred credit assignment
- Discovery component
 - Covering of non-explored niches, refinement of poorly explored niches

XCS' algorithmic structure (2)



Matching

- At each time step t XCS retrieves a binary string on length n
- This string is denoted as $\sigma(t) \in \{0,1\}^n$
- Example for $n = 6$ and $t = 1$: $\sigma(1) = 011001$
- Each classifier maintains a condition or schema C .
- The conditions are encoded ternary, i.e. $C \in \{0,1,\#\}^n$.
- The $\#$ symbol serves as wildcard or 'don't care' operator.
- Examples of conditions: (is matching $\sigma(1)$?)
 - 0#1001 (yes)
 - #01001 (no)
 - 011##1 (yes)

Matching is the process of scanning the entire population [P] for classifiers with a condition that fits the situation $\sigma(t)$

The system prediction

- The system prediction $P(a)$ is a **fitness-weighted sum of predictions** of all classifiers advocating action a

$$P(a) = \frac{\sum_{cl \in [M] | cl.a=a} cl.F * cl.p}{\sum_{cl \in [M] | cl.a=a} cl.F}$$

- Especially at this place, the separation of strength and accuracy plays a major role!
- For each **possible action** $a \in A$ there exists one entry within the PA.
→ There may be several classifiers supporting the same action.

Update rules:

- $\epsilon_j = \epsilon_j + \beta(|P - p_j| - \epsilon_j)$
- $p_j = p_j + \beta(P - p_j)$
- $F_j = F_j + \beta(k'_j - F_j), \quad k'_j = \frac{k_j}{\sum_{cl_i \in [A]} cl_i.k}, \quad k_j = \alpha \left(\frac{\epsilon_j}{\epsilon_0} \right)^{-\nu}$
- β is the **learning rate** (typically set to 0.2)
- α (often set to 0.1) and ν (usually set to 5) control **how strong accuracy decreases** when the error is higher than ϵ_0
- ϵ_0 defines the **targeted error level** of the system
- In single-step problems: P is set to the reward r_{imm}
- Classifier attributes are updated by means of the **modified delta rule** (Widrow-Hoff delta rule) in combination with the moyenne adaptiv modifee (MAM) technique.

Covering

- Covering is the process of generating a novel classifier that matches the current input whenever:
 - Match set $[M]$ is empty (i.e. no matching cl in $[P]$).
 - $[M]$ is poor, i.e. average fitness below a certain threshold.
 - $[M]$ contains less than θ_{mna} distinct actions.
- The condition of the covered classifier cl_{cov} is set to the current input.
- Additionally, each bit is replaced by a # (for generalisation purposes) with probability $P_{\#}$.
- Values for p , ϵ and F are set to predefined initial values (typically 10.0, 0.0 and 0.01).

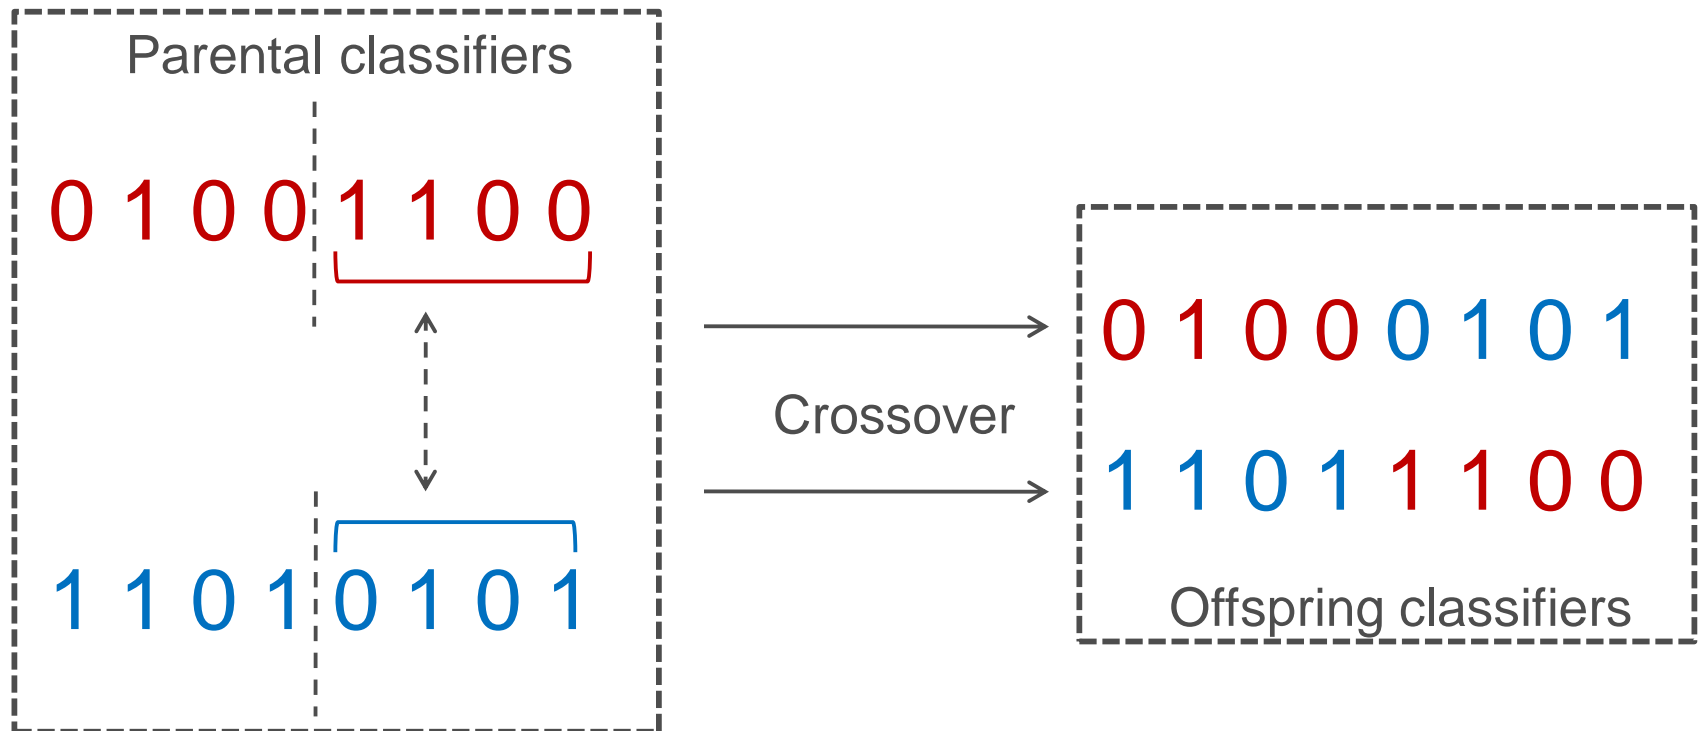
Genetic Algorithm:

- One of the most essential part of XCS is the incorporated niche Genetic Algorithm (GA).
- It is triggered when the average time of all classifiers in $[A]$ since the last GA invocation is greater than θ_{GA} (often set to 50).
- The GA selects two parents from $[A]$ with a probability proportional to their fitness values (roulette-wheel selection).
 - The higher a classifier's fitness, the higher the selection chance.
- The selected parents are copied to generate two offspring classifiers cl_{off} .

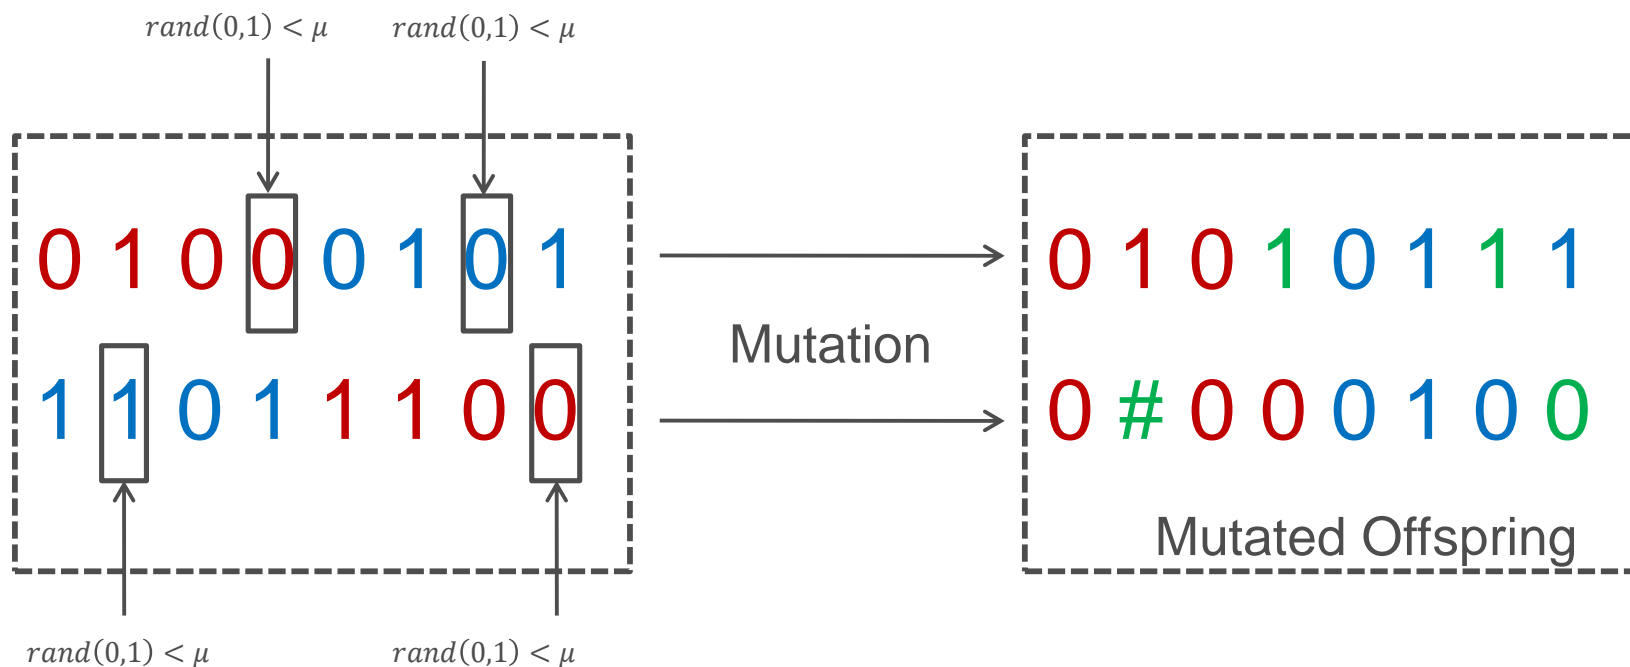
Genetic operators

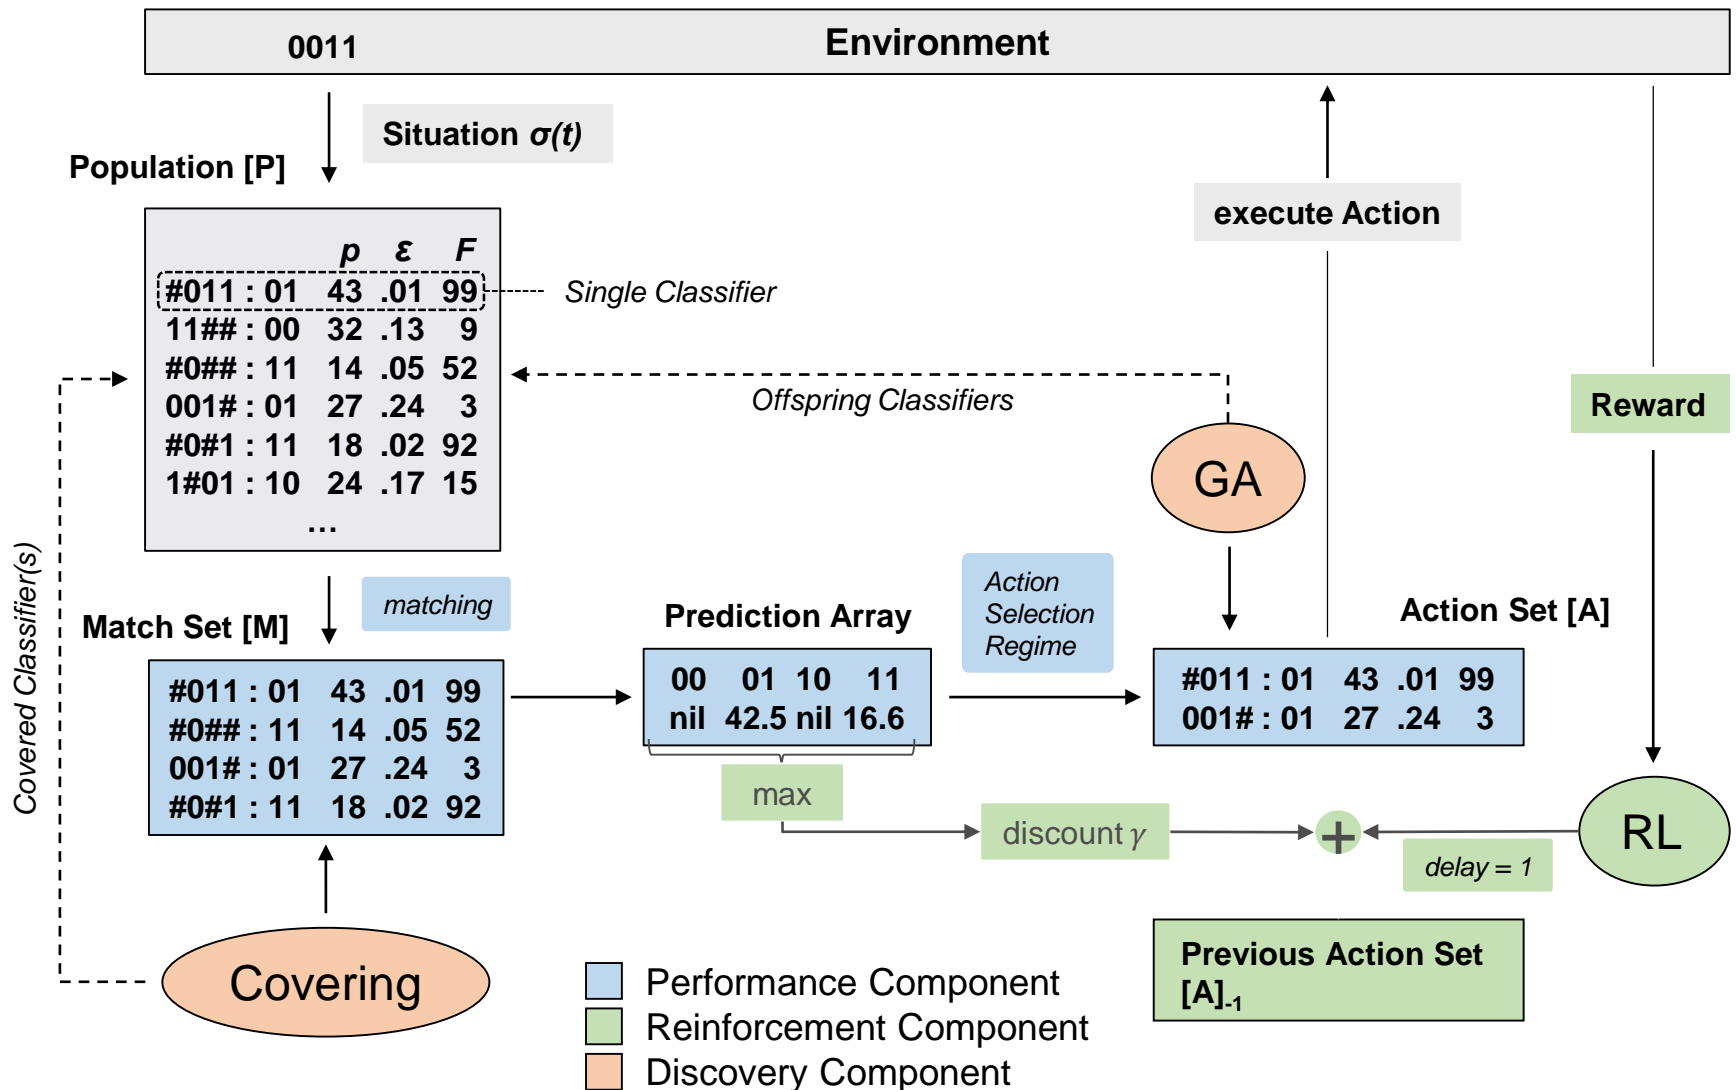
- The conditions of both cl_{off} are crossed (**crossover operator**):
 - One-point crossover: Each offspring classifier's condition is split at a certain point and switched with the other offspring classifier.
 - n-point crossover: more than one point is determined for switching.
 - Uniform crossover: Each value is switched with probability $P_{\chi} = 0.8$.
- Afterward, each bit is flipped with probability $P_{\mu} = 0.04$ to one of the other allowed alleles, that is $\{0, 1, \#\}$.

One-point crossover:



Mutation:





Credit assignment

- r may or may not be retrieved in each step.
- Update of classifier attributes is performed on the action set of the previous time step $t - 1$ ($[A]_{-1}$).
- The maximum **system prediction** $P(a)$ from the PA is discounted by a factor γ (usually $\gamma = 0.95$).
- Additionally, the reward from the previous time-step is added in (may be 0).
- This delay allows to retrieve “**information from the future**”.

Distinguish:

- In **single-step** environments $P = r_{imm}$.
- In **multi-step** problems $P = r_{t-1} + \gamma * \max P(a)$.

Real-world problems

- E.g.: traffic control
- There is no 'end' of the process!
- Hence: there is no reward!
- However, we can handle the control problem as single-step problem.
 - Activate XCS in discrete cycles.
 - Perform observation and adaptation loop.
 - Use utility function: (i) to estimate success, (ii) to analyse conditions.
- For the remainder of this lecture, we only consider single-step problems with immediate reward.

- Wilson proposed changes to conventional XCS to allow for real-valued input (cf. [Wilson2000]).
- To accomplish this, some changes were necessary regarding the internal calculations:
 - Representation of the condition
 - Matching
 - Covering
 - GA (Crossover, Mutation)
- The extended system is called **XCSR** in literature.

The condition part \mathcal{C} :

- The situation $\sigma(t)$ is now represented as an **n -dimensional input/feature vector**.
- The input vector is defined as $\vec{x}_t = (x_1 \dots x_n) \in X \subseteq \mathbb{R}^n$.
- Thus: a binary string is not appropriate anymore!
- For each dimension x_i a so-called **interval predicate** has to be defined.
- An **interval predicate is a tuple (l_i, u_i)** representing:
 - a lower l_i and
 - an upper bound u_i .
- The geometric interpretation of a condition for real-valued inputs is that of hyper-rectangles.
- Accordingly, this condition representation is called **hyper-rectangular representation**.

- In the following, we assume that the input space X is normalised to the standard interval: $[0,1]$.
- Thus, for an $n = 2$ dimensional problem space a classifier's condition C may look like:

$$cl.C = [(0.30, 0.70), (0.55, 0.95)]$$

- E.g.: This condition would match the input
 $\sigma(t) = (0.4, 0.75)$
- In general, a classifier matches the current input if and only if

$$\forall i: l_i \leq x_i < u_i \quad i = 1 \dots n$$

- When covering occurs the newly generated classifier is initialised with predefined initial values as before.

- The condition is set to the current situation

$$\sigma(t) = (x_1 \dots x_n)$$

- Additionally, to provide interval predicates

$$(l_i, u_i), i = 1 \dots n$$

- $l_i = x_i - rand(r_0)$
- $u_i = x_i + rand(r_0)$
- $rand(r_0)$ delivers a uniformly distributed random number between 0 and r_0 .
- r_0 is a predefined default spread.

- **Crossover** is actualised as in standard XCS
 - It can be distinguished whether it is allowed to cross in-between a certain interval predicate or only between the interval predicates
- When an allele is selected for **mutation** its current value is updated according to the following rule:
 - $l_i \pm rand(m_0)$
 - $u_i \pm rand(m_0)$
 - m_0 is a predefined mutation value to extend or shrink the current interval.
- The alleles to mutate are selected probabilistically as in standard XCS.

A 'real toy problem': the checkerboard problem

Checkerboard Problem - CBP(2,8)

cf. [Stone2003]

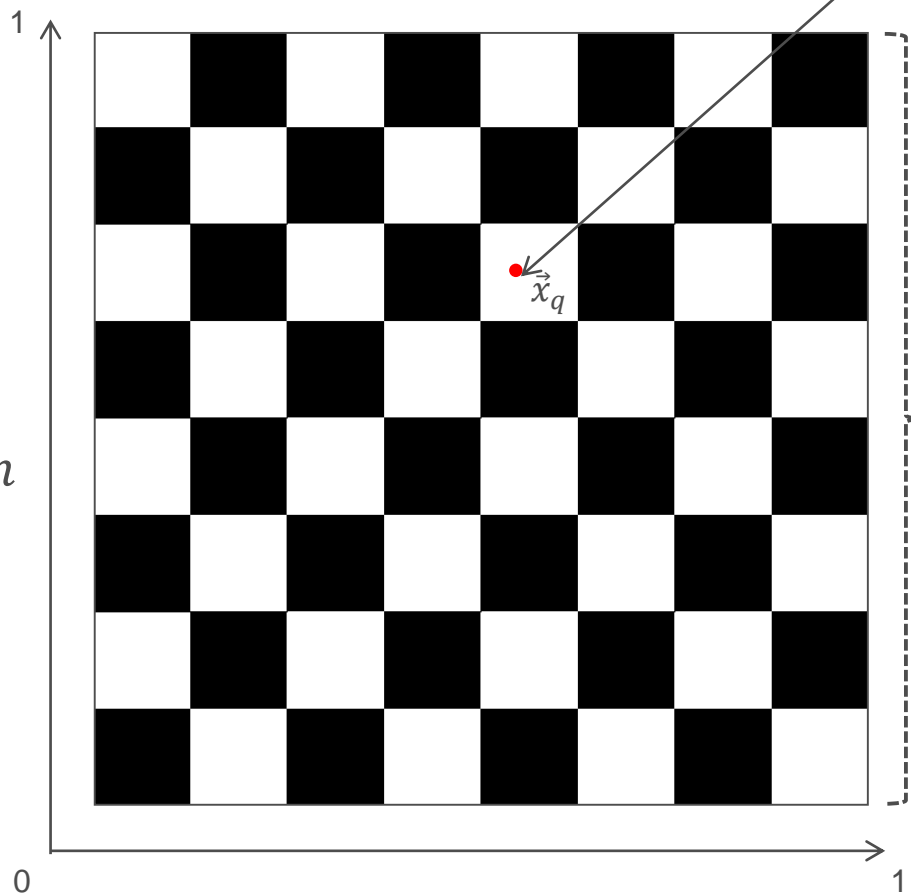
2 dimensions

$n = 2$

each within $[0,1]$

$\vec{x}_q \in \mathbb{R}^n$

$x_i \in [0,1], i = 1 \dots n$



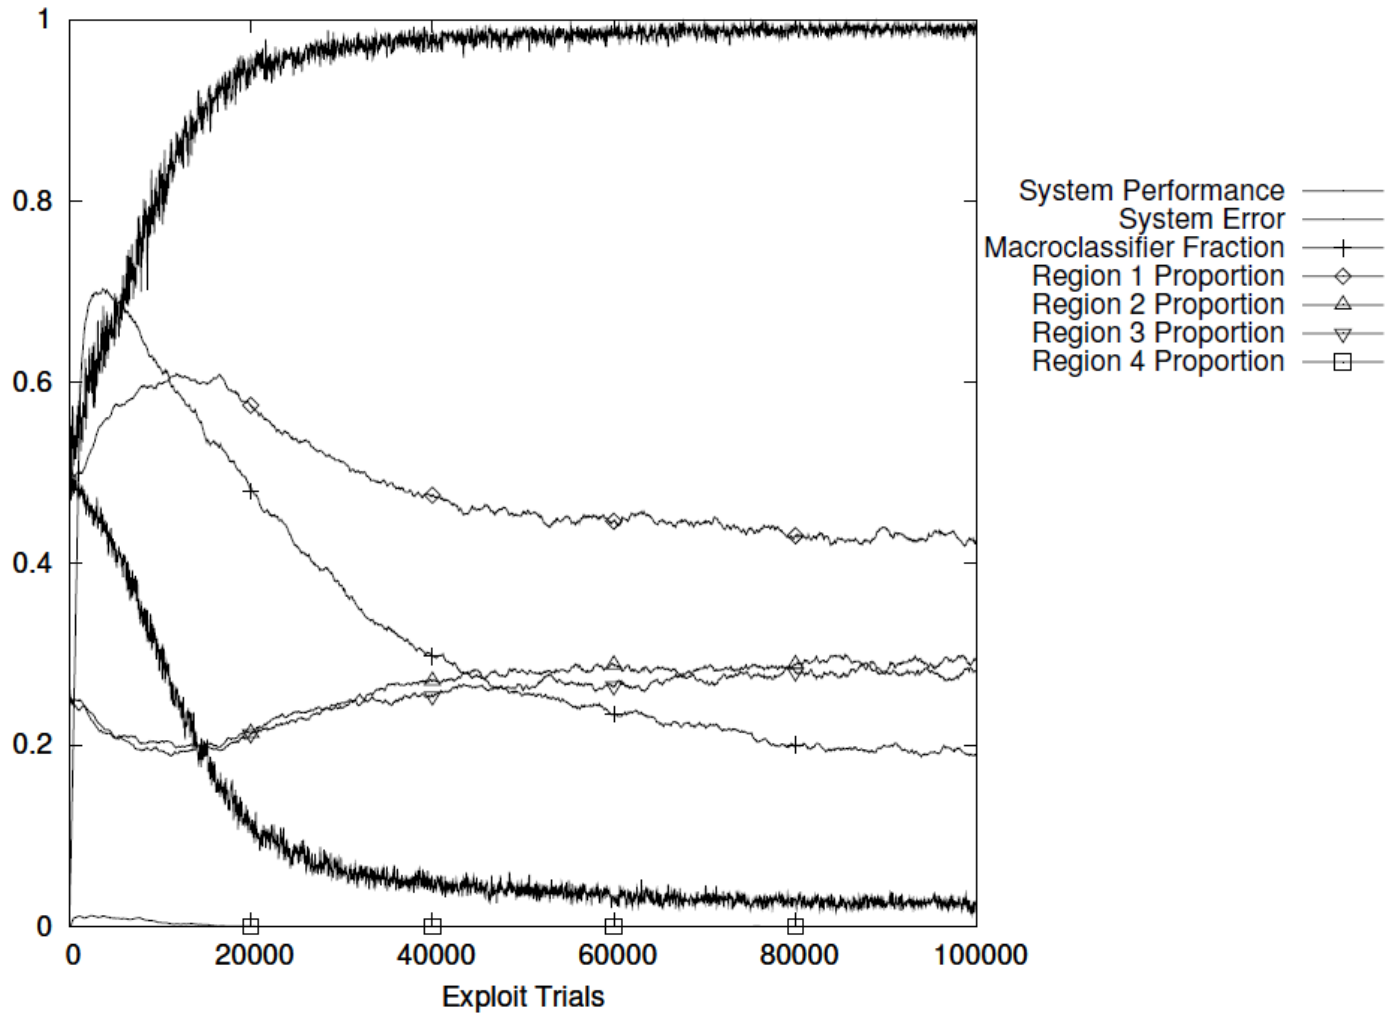
Task of XCS-R:

Of which colour is the field encompassing the query point \vec{x}_q ?

8 divisions $n_d = 8$
for each dimension
with alternate field
colours (black/white)

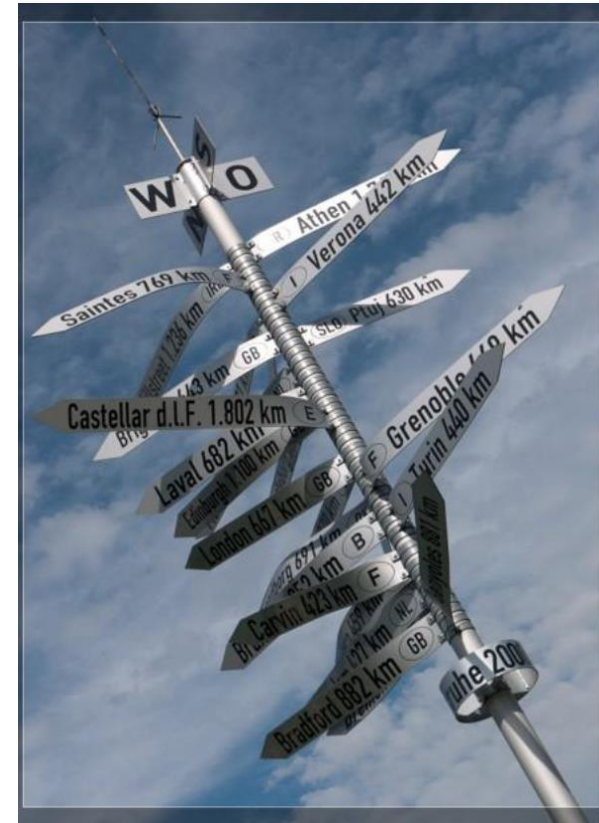
- A CBP(2,8) situation $\sigma(t)$ represents coordinates.
 - E.g.: $\vec{x}_q = (0.25, 0.79)$
- Possible actions are 'black' and 'white', respectively '0' and '1', thus $A := \{0,1\}$
- Reward is 1000 for correct guess and 0 for wrong guess

Single-step or multi-step problem?



cf. [Stone2003]

- Motivation
- Extended Classifier System
- **XCS-O/C**
- Artificial Neural Networks
- Conclusion and further readings



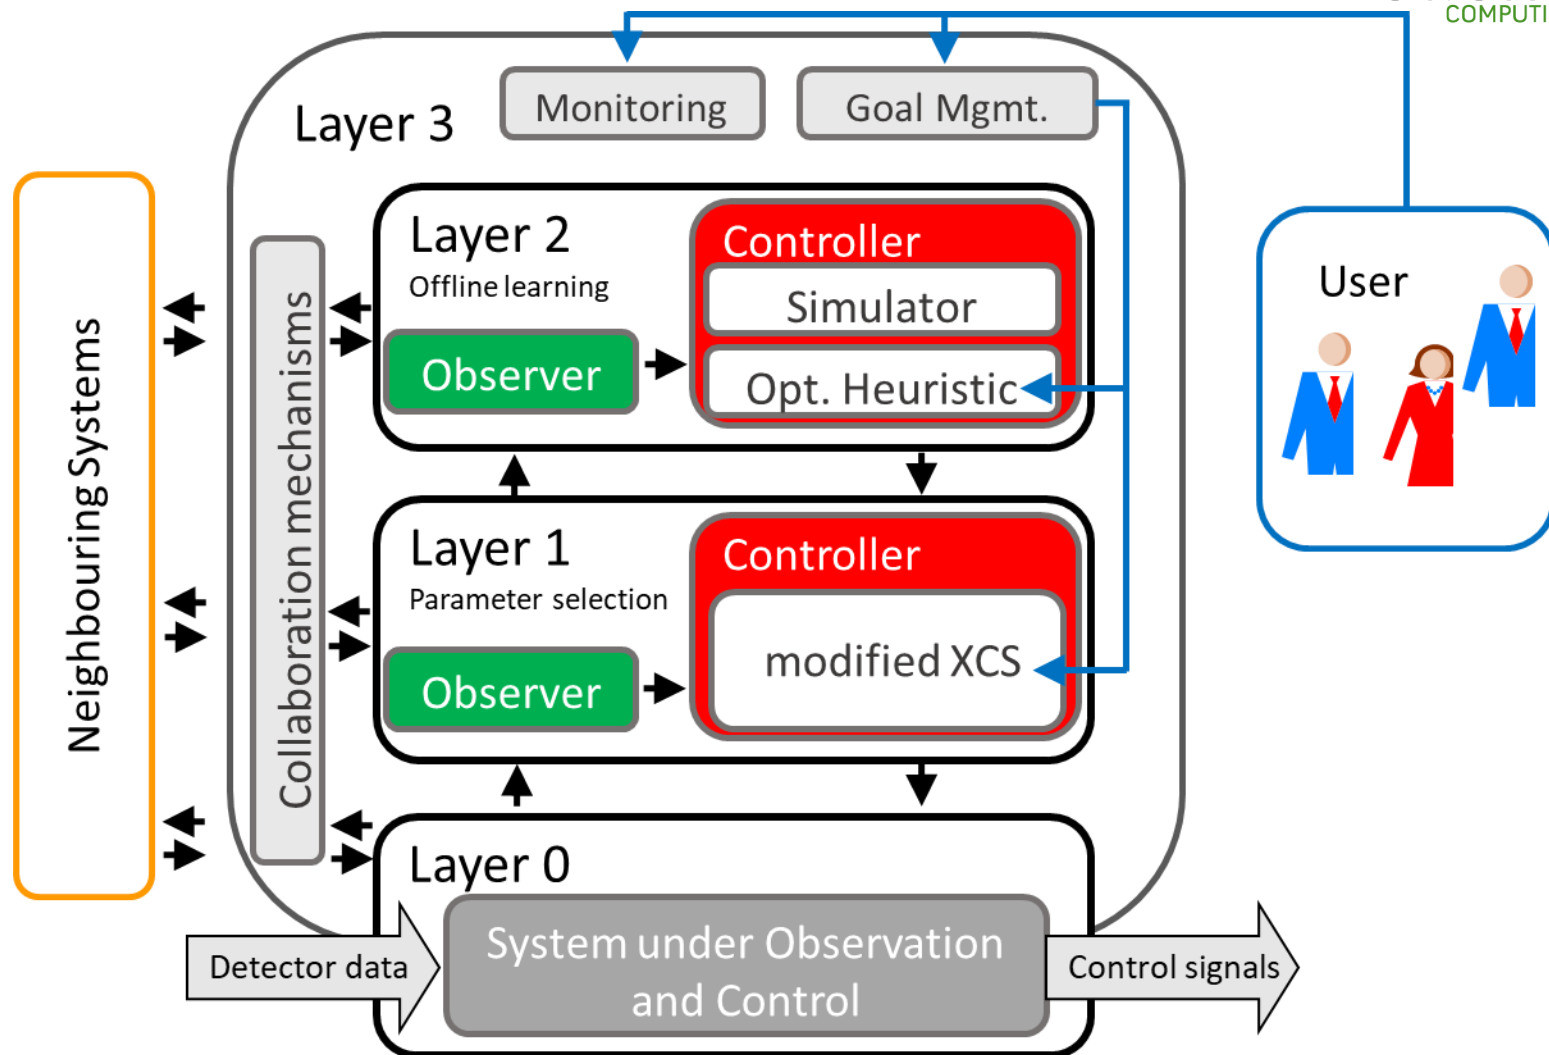
Modifications needed

- Exploring configuration space **online** using GA **can be dangerous!**
 - System will try suboptimal (or even bad) solutions.
 - Example: system could set all traffic lights to green.
 - We cannot allow failures!
- Learning requires experience!
 - What to do in case of missing knowledge?
 - How to **avoid failures** if the action is unknown (i.e. bad/good)?
- Approach:
 1. Provide '**sandbox**' for trying novel behaviour.
 2. Use action that works under '**similar**' conditions.

Multi-level observer/controller framework

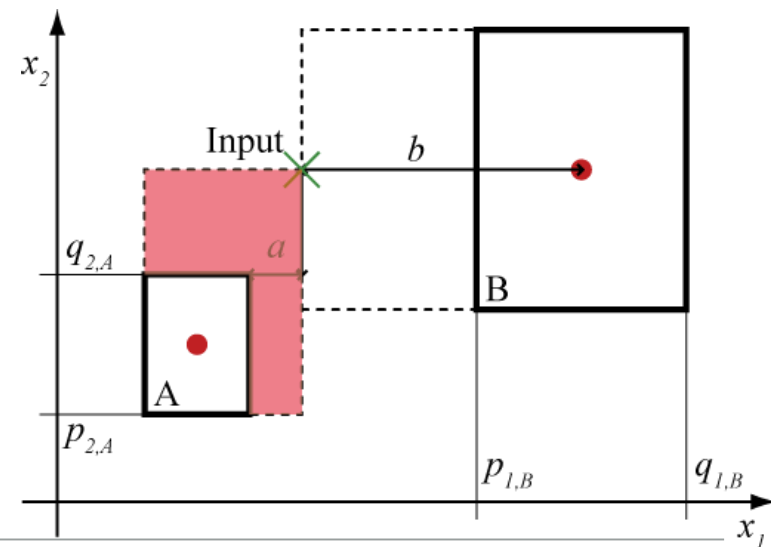
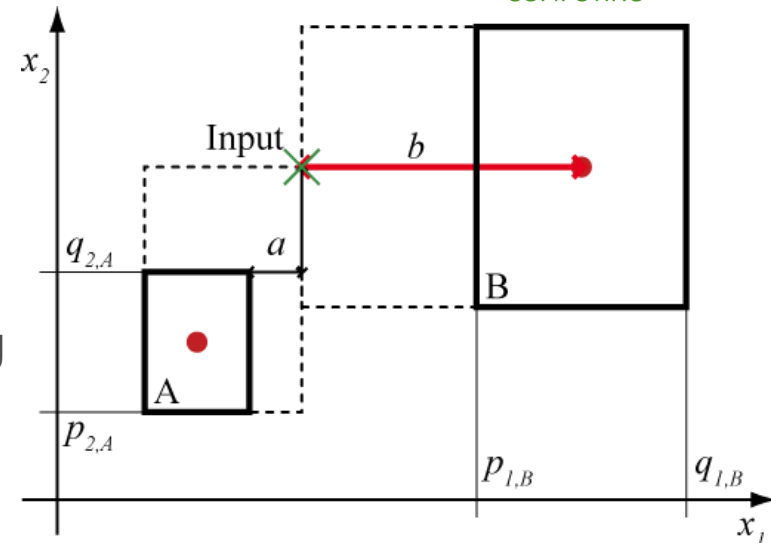


ORGANIC
COMPUTING



Covering

- Original XCS: covering creates new classifier for the current situation randomly.
- OC: application specific widening of existing classifiers.
 - Select “closest” classifier.
 - Copy classifier.
 - Widen condition until matching.
- Trade-off between “use only tested solutions” and quick reaction time.
- Additionally: threshold used to trigger rule generation at Layer 2.



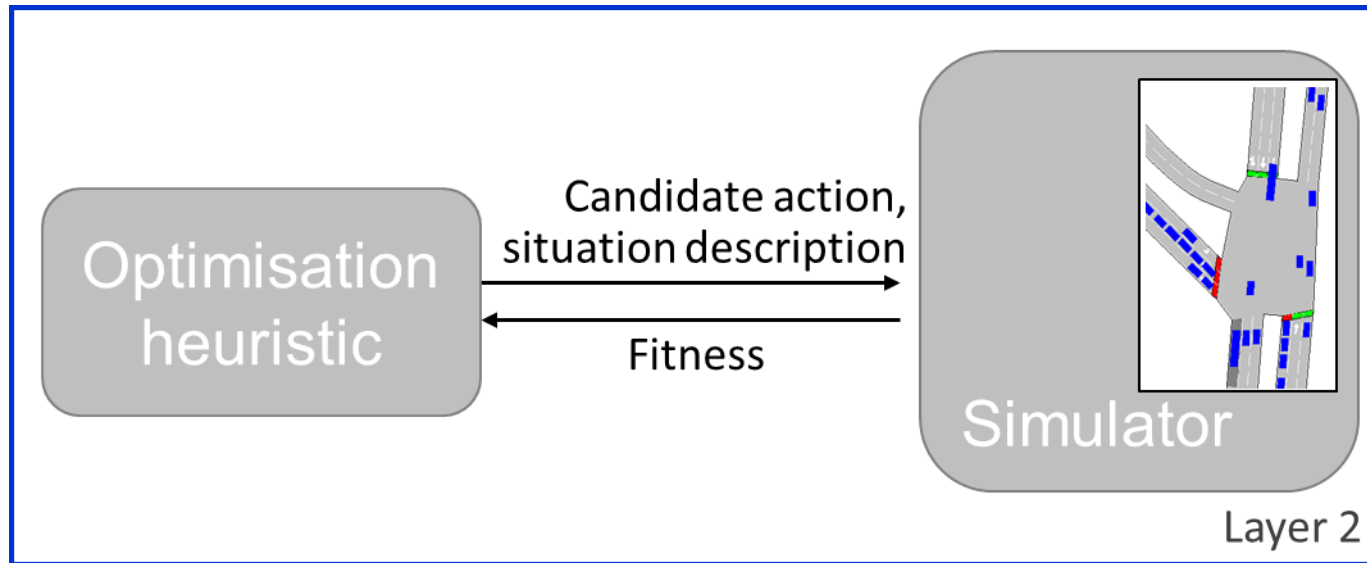
Idea:

- Covering at Layer 1 solves only a part of the problem.
- Building of match set: trade-off between “use only matching solutions” and competition needed for learning.
- What to do with an empty population?

Approach:

- Generation of new rules is done in a separate component.
- Learning is done offline in a simulator (Layer 2).
- Offline means: takes some time...
 - In the meantime, Layer 1 reacts with covering.

- Generation of candidate actions.
- Quality of action is tested using simulations.
- Simulator is configured using observed conditions (situation).



- Fitness is measured.
- Process is repeated until stop criterion is reached.

Novel rule

- Build as follows:
 - **Condition**: situation description, widened using standard interval.
 - **Action**: parameter set as result of optimisation process.
 - **Prediction**: measured in simulation.
 - I.e. utility as observed for best action.
 - **Fitness**: average of all classifiers in population
 - **Error**: zero
- Added to rule base of Layer 1

Process is activated:

- If match set is empty.
- If fitness of rules in match set is below certain threshold.
- Periodically.

Modifications:

- Representation
 - Use classifiers as basis for interpolation
 - Avoids knowledge gaps
- Generalisation
 - More sophisticated rule combination concepts
 - Recombination of partly matching classifiers
- Involve user
 - Combine, e.g., with active learning concept.
 - Proactive knowledge generation
- Second-order optimisation
 - XCS comes with several parameters
 - Adapt them at runtime (i.e. customisation)

XCS makes use of a Genetic Algorithm

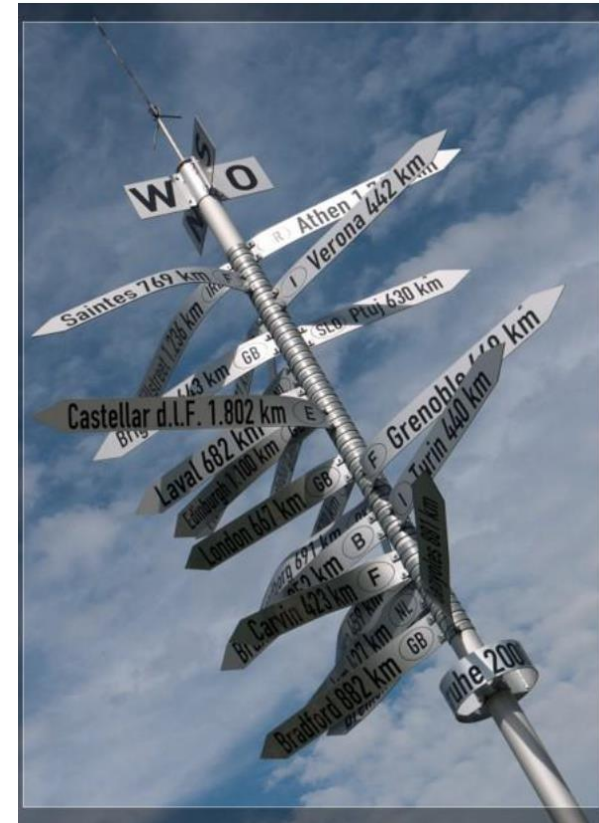
- Is part of the research domain of Evolutionary Computing
- There are alternatives!
- Requirements:
 - Find “good enough” solutions
 - Come up with preliminary result quite fast
- Possible methods and techniques:
 - If available: approximation functions or mathematical functions
 - Swarm-based optimisation heuristics, e.g. Particle Swarm Optimisation
 - OC-based, e.g. Role-based imitation algorithm by Cakar et al.
 - Mimicking physical processes, e.g. simulated annealing
 - An many more...

Everything in software...?

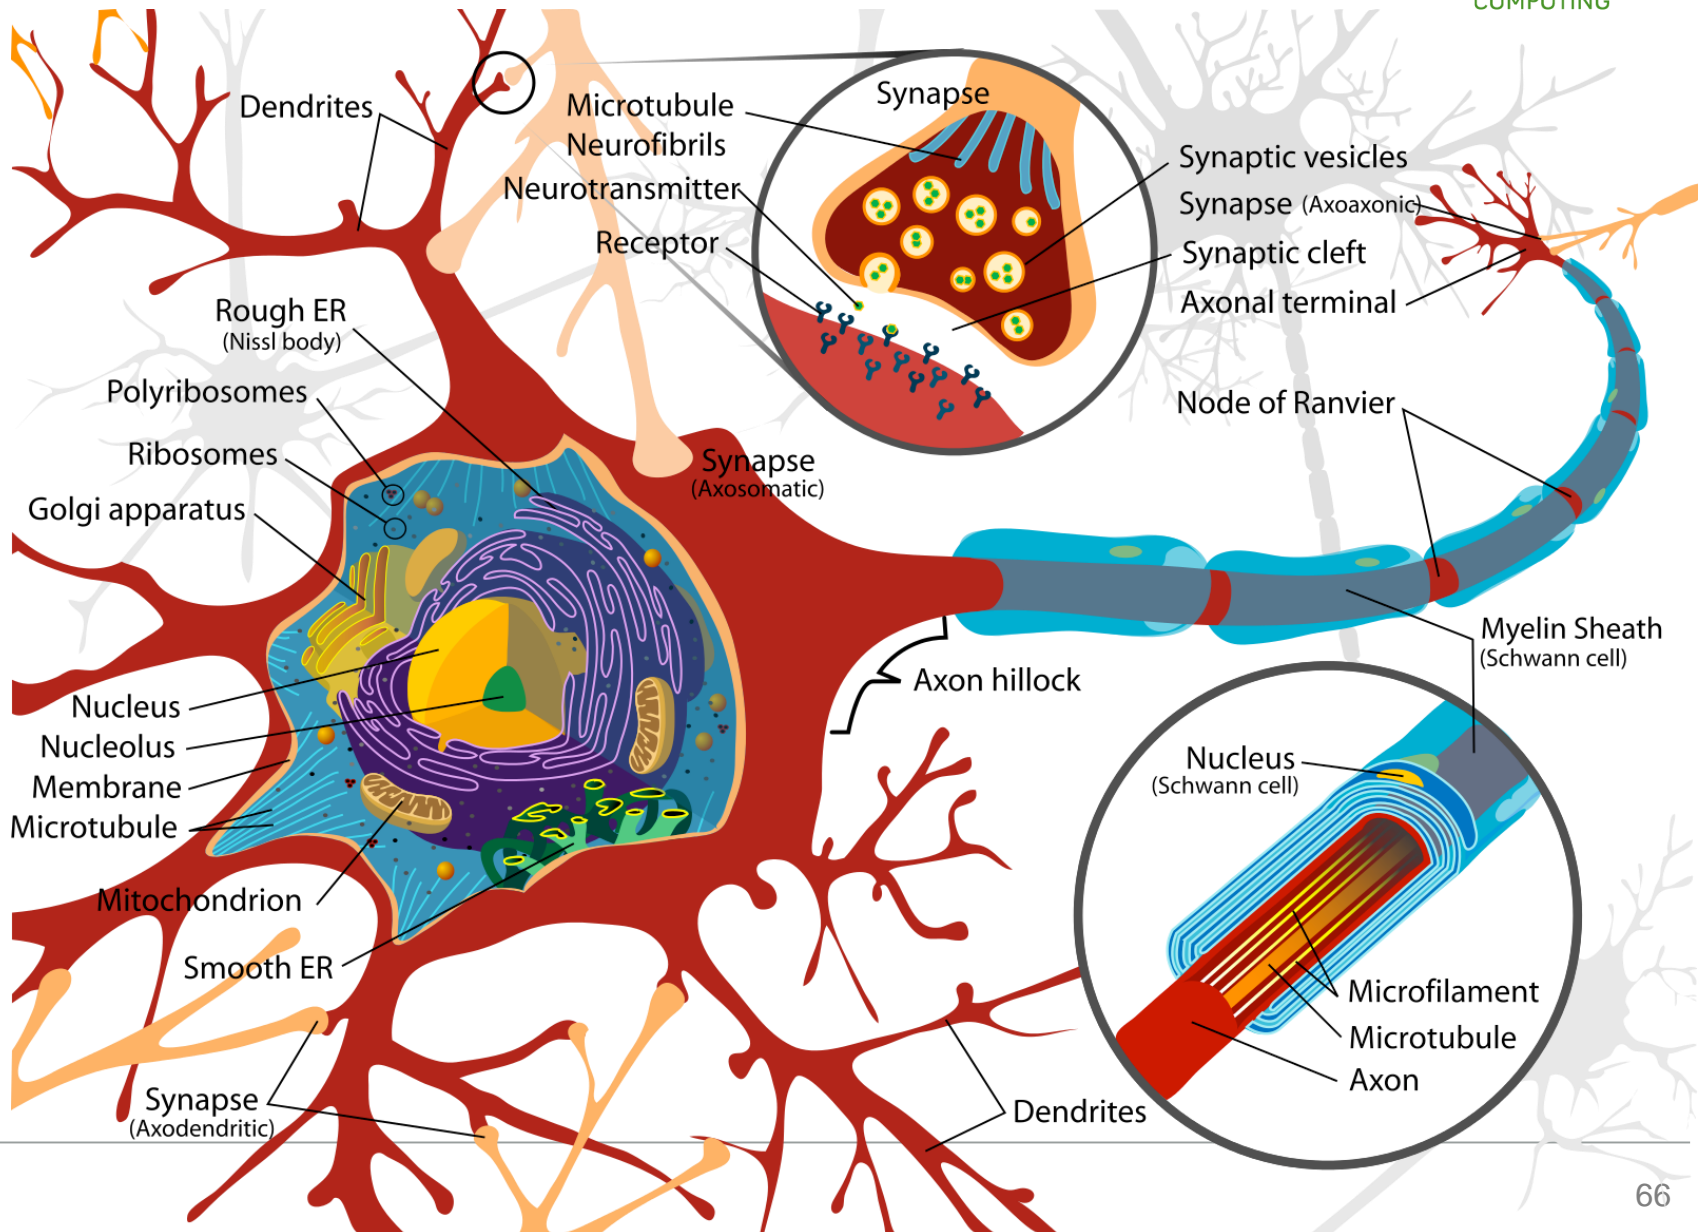
- Concepts for hardware solutions investigated in the context of OC
- Groups at Munich and Tübingen
- Result: Learning Classifier Table
 - Initial training and rule generation in software
 - Logic at runtime at FPGA
 - Scope in hardware: perform main loop for action selection and modify evaluation parameters

- Besides the hyper-rectangular condition structure, **alternative geometric shapes** have been proposed
 - Hyper-spherical or Hyper-ellipsoidal representation
 - Weaken the negative effect of high prediction errors in the corners of rectangles
- There are situations at which XCS is hardly able to learn
 - **Covering challenge**: Covering-Deletion Cycle
 - **Schema challenge**: from over-specification and over-generalisation to maximal general/specific and accurate subspaces
 - **Detrimental forgetting** of rarely sampled niches

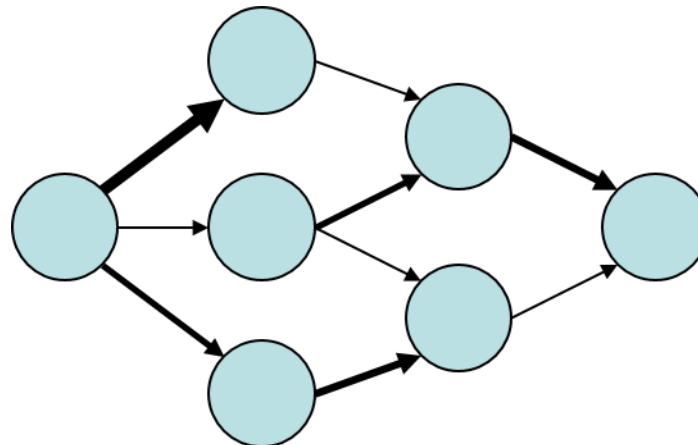
- Motivation
- Extended Classifier System
- XCS-O/C
- **Artificial Neural Networks**
- Conclusion and further readings



Neural networks in nature



- a set of **simple processing elements** (neurons)
- each neuron
 - may receive multiple **inputs**
 - provides an **output**
- input and output values are mediated by **weighted connections**
- the connected neurons form a directed graph, the **network**



What are ANNs used for?

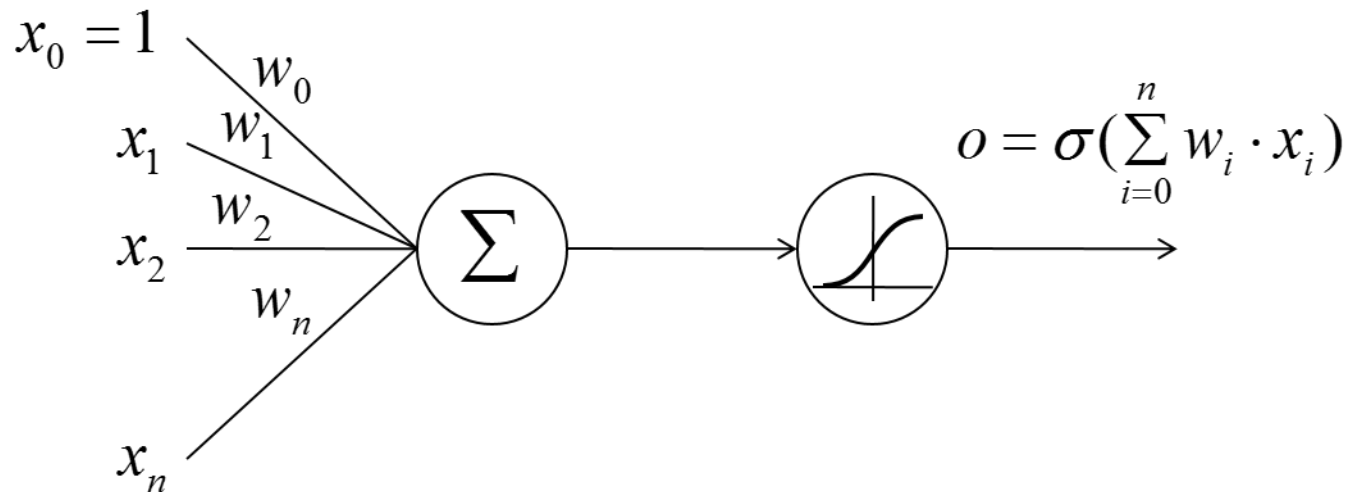
Function approximation!

(i.e. supervised learning)

How can functions be approximated using ANNs?

- altering the network's weights \Rightarrow calculations of the network are altered
- thus: adjust the weights so that the network maps certain inputs to certain outputs
- this process is referred to as training
- training relies on a set of examples, i.e. a set of inputs along with desired outputs

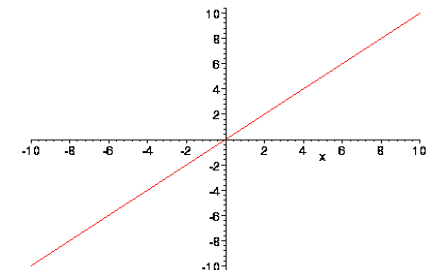
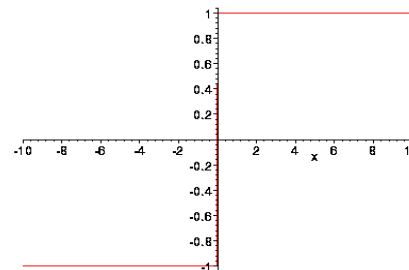
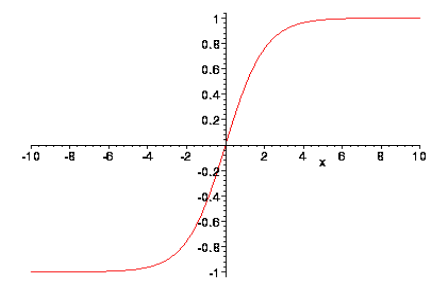
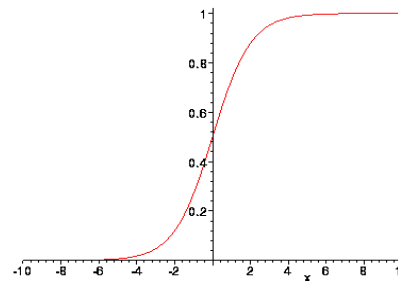
1. Calculate **weighted sum** of inputs.
2. Apply **activation** function σ .
3. Provide output.



- map net input value to neuron output
- many possibilities (linear, nonlinear, ...)
- choice can influence necessary training

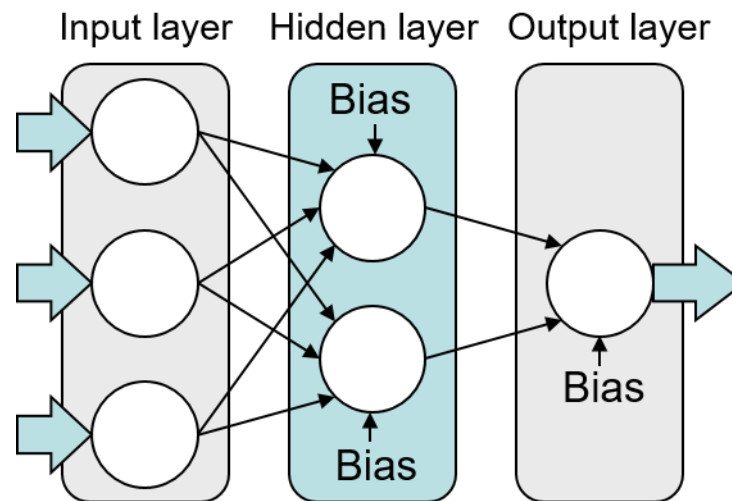
Typical activation functions

- logistic function (aka sigmoid)
- hyperbolic tangent function
- step function
- Identity function

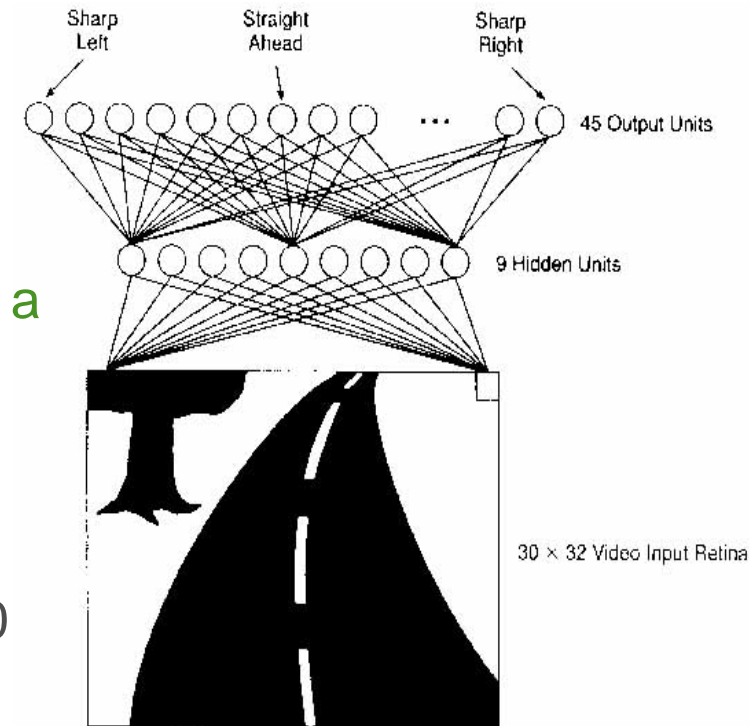


- feed-forward
- convolutional
- recurrent
- ...

- special case of feed-forward networks
- each layer **fully connected** to the next one
- no other connections



- ALVINN: **A**utonomous **L**and **V**ehicle **I**n a **N**eural **N**etwork
- learn to **steer a vehicle** driving at reasonable speed on a public highway
- steering decisions based on **images from a forward pointing camera**
- training by **watching a human driver**
- ALVINN is able to drive at speeds of around 100 km/h and for distances of 150 km after only 5 minutes of training



Pomerleau, D. A. (1993). Knowledge-based training of artificial neural networks for autonomous robot driving. In *Robot learning* (pp. 19-43). Springer US.

- Goal: Find weights such that input generates desired output!
- Idea: Minimize error of prediction!
 - ⇒ Optimization problem!
- Several options to solve this, e.g.
 - gradient descent (e.g. backpropagation)
 - evolutionary algorithms

e.g. mean squared error

$$\frac{1}{n} \sum_{i=1}^n (y_{true} - y)^2$$

- idea:
 - compute the network's prediction
 - calculate the prediction's error
 - adjust the network, layer by layer, starting from the output layer
- use **gradient of error measure** for weight updates
 - ⇒ error measure must be **differentiable**
- Note: Prediction is part of error measure ⇒ Activation function must be differentiable as well!

$$\Delta w_{ij} = -\eta \frac{\partial E}{\partial w_{ij}}$$

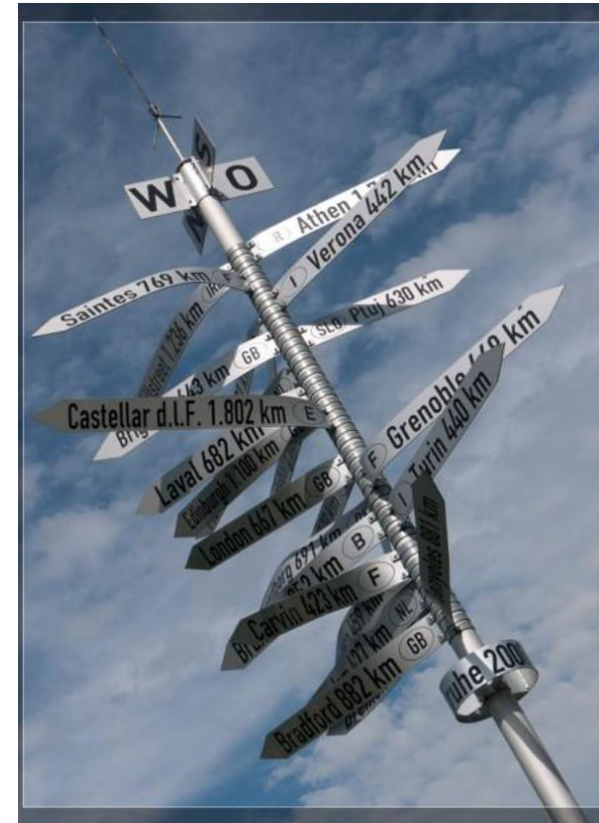
Drop the derivatives! Neuroevolution!

- Why not use evolutionary optimization instead of gradient descent for **weight update**?
- Why not evolve the network's **structure** as well?

- represent network weights as an individual of a population
- **fitness measure** for an individual: How well does a network with the individual's weights perform on the data?
- perform the usual GA operations on these individuals, e.g.
 - definitely mutation and selection
 - perhaps crossover
- result: network weights
- This works surprisingly well!
 - often less computationally intensive on (really) large networks
 - often the same or better performance as other optimization methods

- represent network structure as an individual of a population
- fitness measure for an individual: How well does a network with the individuals's structure perform on the data?
 - to calculate fitness: Have to train the network (using some optimization like backpropagation or another GA)
 - usually **computationally intensive**
- perform the usual GA operations on these individuals (mutation, crossover, selection, ...)
- result: network structure

- Motivation
- Extended Classifier System
- XCS-O/C
- Artificial Neural Networks
- **Conclusion and further readings**



This chapter:

- Introduced to the concepts of machine learning.
- Discussed the challenges of learning control strategies in organic systems.
- Presented the Extended Classifier System and its modifications for usage in organic systems.
- Compared concepts such as strength vs. accuracy, exploration vs. exploitation, etc.
- Highlighted how credit assignment is realised and summarised the corresponding challenges for real-world problems.

By now, students should be able to:

- Explain how learning is done in organic systems.
- Summarise the ideas of machine learning in technical applications.
- Outline the structure and process of XCS and its variants.
- Highlight the tasks of the major components and their impact on the learning process.
- Explain how what credit assignment is and how it is realised in OC.
- Discuss the major concepts and customisations used in organic systems.

- [Wilson1995]: Wilson, S. W.: Classifier Fitness Based on Accuracy. In: *Evolutionary Computation 3 (1995), no. 2, pp. 149-175*
- [Wilson1998]: Wilson, S. W.: Generalisation in the XCS Classifier System. *Morgan Kaufmann. Genetic Programming 1998: Proceedings of the Third Annual Conference, 1998, pp. 665-674*
- [Wilson2000]: Wilson, S.: Get Real! XCS with Continuous-Valued Inputs. In: Lanzi, P.; Stolzmann, W. & Wilson, S. (Eds.): *Springer Berlin Heidelberg. LNCS 1813, Learning Classifier Systems, 2000, pp. 209-219*
- [Stone2003]: Stone, C. & Bull, L.: For Real! XCS with Continuous-Valued Inputs. In: *Evolutionary Computation 11 (2003), no. 3, pp. 298-336*
- [Butz2002]: Butz, M. & Wilson, S. W.: An Algorithmic Description of XCS. In: *Soft Comput. 6 (2002), no. 3-4, pp. 144-153*

Questions ...?